

Azure Architect Mastery

Concepts · Labs · Scenarios · Whiteboards · Quizzes · Flashcards · AZ-305

Contents

Azure Architect Mastery — Enterprise APIs & AZ-305 Study Guide	1
1. Azure API Management (APIM)	1
2. Azure Service Bus	5
3. Azure RBAC & Governance	7
4. Cloud Networking	10
5. Azure Container Apps (ACA)	12
6. Authentication & Authorization Patterns (Entra ID)	14
7. Enterprise API Architecture & Design Patterns	16
8. Data Storage Design	19
9. Business Continuity & Disaster Recovery	21
10. Monitoring & Observability	23
11. Compute Design (Beyond Containers)	25
12. Data Integration, Analytics & Migration	26
13. Platform Extras: Key Vault, Hybrid, IaC	28
14. Concepts Explained Simply (Layman's Corner)	30
15. Real-World Architecture Walkthroughs	33
16. Hands-On Labs (CLI + Bicep, with Validation & Troubleshooting)	36
17. One-Page Cheat Sheets	41
18. Practice Questions (Q&A)	43
19. Flashcards	92

Azure Architect Mastery — Enterprise APIs & AZ-305 Study Guide

A comprehensive guide covering Azure API Management, Service Bus, RBAC, Cloud Networking, Container Apps, authentication/authorization patterns, and enterprise API architecture — aligned to the AZ-305 (Azure Solutions Architect Expert) exam.

1. Azure API Management (APIM)

1.1 What APIM Is

Azure API Management is a hybrid, multicloud platform for publishing, securing, transforming, and monitoring APIs. It sits between API consumers and backend services as a **facade**, decoupling clients from backend implementation.

Three core components:

- **API Gateway (data plane):** Accepts API calls, routes to backends, verifies keys/JWTs/certificates, enforces quotas and rate limits, transforms requests/responses, caches, and emits logs/metrics/traces.

- **Management plane:** Azure portal / ARM / Bicep / Terraform / CLI surface for provisioning, defining APIs, packaging into products, setting policies, and analytics.
- **Developer portal:** Auto-generated, customizable website where consumers discover APIs, read interactive docs, try APIs in a console, and get subscription keys.

1.2 Tiers (know these for the exam)

Tier	Key traits	Use case
Consumption	Serverless, per-call billing, scale-to-zero, no VNet injection, ~sub-second activation	Lightweight/serverless APIs, spiky traffic
Developer	Full features, no SLA, single unit	Dev/test, evaluation
Basic	99.95% SLA, small scale	Entry production
Standard	99.95% SLA, more scale units	Medium production
Premium	Multi-region deployment, VNet injection, availability zones, self-hosted gateway, workspaces, higher scale	Enterprise, hybrid, mission-critical
v2 tiers (Basic v2, Standard v2, Premium v2)	Faster provisioning/scaling, VNet integration (outbound) in Standard v2; Premium v2 adds VNet injection	Modern deployments needing agility

Exam signal: “multi-region,” “VNet injection,” “availability zones,” or “self-hosted gateway” → **Premium** (or Premium v2 for injection).

1.3 Core Concepts

- **APIs and operations:** An API maps to a backend; operations map to endpoints. Import from OpenAPI, WSDL, gRPC (self-hosted), Azure Functions, Logic Apps, Container Apps, AKS.
- **Products:** Bundles of one or more APIs with a title, terms, visibility, and policies. Can be *open* (no subscription needed) or *protected* (subscription key required).
- **Subscriptions and keys:** A subscription grants access to a product (or all APIs, or a single API) via primary/secondary keys sent in the `Ocp-Apim-Subscription-Key` header.
- **Groups:** Administrators, Developers, Guests — control product visibility in the developer portal.
- **Named values:** Key-value store for policy configuration; can reference **Key Vault secrets**.
- **Backends:** Reusable backend entities with credentials, mTLS client certificates, and **circuit breaker** rules + **load-balanced pools**.
- **Versions vs. revisions:**
 - **Versions** = breaking changes, exposed to consumers (path /v2/, query string, or header versioning schemes).
 - **Revisions** = non-breaking iterations of the same version; test a revision via `;rev=2` URL then make it current. Safe rollout mechanism.
- **Workspaces (Premium):** Decentralized API teams manage their own APIs/products/subscriptions with isolated runtime on workspace gateways — supports federated API management at enterprise scale.

1.4 Policies — the heart of APIM

Policies are XML statements executed in the gateway pipeline in four sections:

inbound → backend → outbound → on-error

Policy scopes and evaluation order: **Global → Workspace → Product → API → Operation**, composed via the `<base />` element (position of `<base/>` controls whether parent policies run before or after yours).

Must-know policies:

Policy	Purpose
rate-limit / rate-limit-by-key	Throttle burst traffic (sliding window, 429 on breach)
quota / quota-by-key	Long-term call/bandwidth caps (e.g., per month)
validate-jwt	Validate Entra ID / OAuth2 JWTs (issuer, audience, claims, signing keys via OpenID config)
validate-azure-ad-token	Simplified Entra-specific token validation
check-header, ip-filter	Basic gating
cors	Cross-origin support
set-header, set-query-parameter, rewrite-uri	Request transformation
set-backend-service	Dynamic routing to different backends
cache-lookup / cache-store	Response caching (built-in or external Redis)
send-request	Call out to another service mid-policy (e.g., token introspection)
mock-response	Return mock for API-first development
retry, forward-request	Backend resiliency behavior
authentication-managed-identity	Gateway acquires a managed identity token to call the backend
llm-token-limit / llm-semantic-cache (AI gateway)	Token quotas and semantic caching for LLM backends

Example — JWT validation at the gateway:

```
<inbound>
  <base />
  <validate-jwt header-name="Authorization" failed-validation-httpcode="401">
    <openid-config url="https://login.microsoftonline.com/{tenant}/v2.0/.well-known/openid-configuration">
      <audiences>
        <audience>api://my-api-app-id</audience>
      </audiences>
      <required-claims>
        <claim name="roles" match="any">
          <value>Orders.Read</value>
        </claim>
      </required-claims>
    </validate-jwt>
  <rate-limit calls="100" renewal-period="60" />
</inbound>
```

1.5 Networking Options

- **External VNet injection (Premium):** Gateway in your VNet, publicly reachable; can reach private backends.
- **Internal VNet injection (Premium):** Gateway only reachable inside the VNet (or via peering/VPN/ER). The standard enterprise pattern: **Application Gateway (WAF) or Front Door in front → internal APIM → private backends**.
- **Private endpoint:** Private inbound access to APIM (no VNet injection needed; inbound only).

- **VNet integration (Standard v2):** Outbound-only reach into a VNet to call private backends.
- **Self-hosted gateway:** Containerized gateway you run on-premises or in other clouds; management plane stays in Azure. Enables hybrid/multicloud API federation.

1.6 Reliability & Scale

- **Multi-region (Premium):** Gateway replicated across regions; primary region hosts the management plane. Combine with routing policies for regional backend affinity.
- **Availability zones (Premium):** Zone-redundant scale units.
- **Autoscale** rules on capacity metric; **caching** to shave backend load.
- **Observability:** Application Insights integration, Azure Monitor metrics/logs, built-in analytics, OpenTelemetry via self-hosted gateway.

1.7 Enterprise API Deployment Checklist (APIM)

1. Premium/internal VNet (or Standard v2 + private endpoints) for network isolation.
 2. Front Door or App Gateway + WAF for edge protection and global routing.
 3. validate-jwt with Entra ID at gateway; subscription keys only as an *identification* mechanism, never sole *authentication* for sensitive APIs.
 4. Products for consumer segmentation; quotas + rate limits per product/key.
 5. Named values backed by Key Vault; managed identity to backends.
 6. Versioning strategy declared up front; revisions for safe changes.
 7. CI/CD: APIOps (extract/publish via Git), Bicep/Terraform infrastructure.
 8. Diagnostics to Log Analytics; alerts on capacity, 4xx/5xx rates, latency.
-

Quick Check — API Management

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — APIM

W1. A bank wants to expose 40 internal APIs to fintech partners with per-partner quotas, OAuth security, and zero public exposure of backends. Sketch the architecture and name the APIM features you'd use.

Model answer: Front Door Premium (WAF, global TLS) → APIM Premium in **internal VNet mode** → private backends. Partners are onboarded via the **developer portal**; each gets a **subscription** under a partner-tier **product** carrying quota and rate-limit-by-key policies. validate-jwt against Entra ID (client credentials for partner apps) at global scope; authentication-managed-identity toward backends; Key Vault-backed **named values** for any residual secrets. Per-partner analytics come free from subscription attribution. Justify Premium: internal VNet + AZ redundancy; everything else would work on Standard v2 with VNet integration if injection weren't needed.

W2. During Black Friday, one consumer's buggy retry loop takes down the product API for everyone. Whiteboard the layered protections you should have had.

Model answer: Edge: WAF rate limiting + bot rules. Gateway: rate-limit-by-key (keyed on subscription/JWT sub) so one consumer exhausts only their bucket — the **throttling** pattern; quota for longer-horizon caps; **circuit breaker** on the backend entity so APIM stops hammering a degraded backend; retry policy with backoff only for idempotent GETs. Backend: **bulkhead** — separate ACA revisions/pools per consumer tier; **queue-based load leveling** for writes via Service Bus. Monitoring: alert on 429 ratio per subscription to catch the loop early. Key exam point: 429 the abuser, don't scale to absorb them.

Key Resources — APIM

- [APIM documentation home](#)
- [APIM policy reference \(every policy\)](#)
- [APIM tiers & feature comparison](#)
- [APIM VNet integration options](#)
- [APIM landing zone accelerator \(Architecture Center\)](#)
- [APIOps guidance](#)

2. Azure Service Bus

2.1 What It Is

Fully managed enterprise message broker supporting **queues** (point-to-point) and **topics/subscriptions** (publish-subscribe). Decouples producers from consumers, providing load-leveling, ordering, and reliable delivery over **AMQP 1.0**.

2.2 Queues vs. Topics

- **Queue:** One logical consumer group; messages pulled by **competing consumers**; each message processed by one receiver. Provides temporal decoupling and load leveling.
- **Topic + subscriptions:** Publisher sends once; each subscription gets an independent copy. Subscriptions support **filters** (SQL filters, correlation filters, boolean filters) and **actions** (modify properties), enabling content-based routing.

2.3 Tiers

Tier	Highlights
Basic	Queues only, 256 KB messages, no topics
Standard	Topics, sessions, transactions, dedup; shared capacity; 256 KB messages
Premium	Dedicated Messaging Units, predictable latency, 100 MB messages (large message support), VNet/private endpoints, availability zones, Geo-DR, CMK encryption, JMS 2.0

Exam signal: “predictable performance,” “private endpoints,” “geo-disaster recovery,” or “>1 MB messages” → **Premium**.

2.4 Key Messaging Features

- **Receive modes:**
 - **Peek-lock (at-least-once):** Message locked (default 60s, renewable); consumer must Complete/Abandon/Dead-letter/Defer. Failure → lock expiry → redelivery.
 - **Receive-and-delete (at-most-once):** Fastest, but message lost if consumer crashes.
- **Dead-letter queue (DLQ):** Automatic sub-queue per entity for poison messages (MaxDeliveryCount exceeded, TTL expired, filter evaluation errors). Path: queue/\$deadletterqueue.
- **Sessions:** Guarantee **FIFO ordering** and stateful processing for related messages sharing a SessionId. Required for strict ordering in Service Bus. Session state can persist workflow progress.
- **Duplicate detection:** Dedupes by MessageId within a configurable time window (idempotent enqueue).
- **Scheduled delivery:** Enqueue now, deliver at a future time.
- **Deferral:** Set aside a message to retrieve later by sequence number.

- **Transactions:** Atomically group operations against a single messaging entity (and cross-entity with transfer/send-via).
- **Auto-forwarding:** Chain queues/subscriptions (fan-in patterns, decoupling).
- **Prefetch + batching:** Throughput optimization.
- **Message TTL, auto-delete on idle, max delivery count.**

2.5 Reliability

- **Availability zones:** Automatic in regions that support them (Premium replicates across zones).
- **Geo-Disaster Recovery (Premium):** Namespace-level metadata pairing + alias; **metadata only** — messages are NOT replicated (know this!). Failover promotes secondary.
- **Geo-replication (Premium):** Newer full-data replication feature for supported regions.
- Client-side resilience: retry policies, idempotent handlers, DLQ monitoring.

2.6 Choosing the Right Messaging Service (classic exam question)

Service	Model	Use when
Service Bus	Broker: queues + pub/sub, transactions, ordering, DLQ	Enterprise messaging, commands, financial transactions, ordered workflows
Event Grid	Reactive discrete events , push delivery, serverless	React to state changes (“blob created”), event-driven automation, fan-out to handlers
Event Hubs	Event streaming , partitioned log, millions of events/sec	Telemetry, clickstreams, big-data ingestion, Kafka workloads
Storage Queues	Simple, cheap, massive queue (>80 GB)	Basic queueing, no ordering/dedup/sessions needed, audit log of all transactions via storage logs

Rule of thumb: **events** (fact happened, lightweight) → Event Grid/Event Hubs; **messages** (payload with contract, sender expects processing) → Service Bus.

2.7 Messaging Patterns

- **Competing consumers:** Multiple workers pull from one queue → horizontal scale + load leveling.
- **Publish-subscribe:** Topic fan-out to independent subscribers.
- **Claim-check:** Store large payload in Blob Storage; send message with reference (avoids size limits).
- **Saga / process manager:** Long-running distributed transactions coordinated via messages with compensating actions.
- **Queue-based load leveling:** Buffer between spiky producers and steady consumers.
- **Sequential convoy:** Sessions to process related message groups in order.
- **Dead-letter handling:** Automated DLQ processor with alerting; never let DLQs grow silently.

2.8 Security

- **Entra ID + RBAC (preferred):** Azure Service Bus Data Owner / Sender / Receiver roles, scoped to namespace/queue/topic. Use **managed identities** from compute.

- **SAS (legacy/compat):** Namespace or entity-scoped policies (Send/Listen/Manage). Rotate keys; prefer Entra ID.
 - **Network:** Private endpoints (Premium), service tags, IP firewall, disable public network access.
-

Quick Check — Service Bus

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Service Bus

W1. An airline's booking system must process seat reservations strictly in order per flight, survive consumer crashes without losing bookings, and never double-charge on retries. Whiteboard the messaging design.

Model answer: Service Bus **queue with sessions**, SessionId = flightNumber → FIFO per flight while different flights process in parallel (sequential convoy). **Peek-lock** receive with explicit Complete → at-least-once; consumers renew locks for long work. Producer sets deterministic MessageId = bookingId + **duplicate detection** window → idempotent enqueue; consumer-side idempotency (upsert by bookingId) guards double-charging on redelivery. MaxDeliveryCount + **DLQ** with an alerted DLQ processor for poison bookings. Premium tier if private endpoints/zone redundancy are required.

W2. Design the messaging backbone for an e-commerce order flow: order placed → inventory, payments, and email must all react; payment failures must be compensated; invoices are 20 MB PDFs.

Model answer: Order service publishes to a **topic** orders with three filtered **subscriptions** (inventory, payments, notifications) — pub/sub fan-out, correlation filters on event type. Payment failure triggers a **saga**: compensating messages (release inventory, cancel order) rather than distributed transactions. Invoices use **claim-check**: PDF to Blob Storage, message carries the URI. Email service scales with **competing consumers**; the whole chain is buffered (queue-based load leveling) so Black Friday bursts don't topple payments. Everything authenticated with managed identities + Service Bus data roles.

Key Resources — Service Bus

- [Service Bus documentation home](#)
- [Queues, topics & subscriptions concepts](#)
- [Compare messaging services \(SB vs Event Grid vs Event Hubs\)](#)
- [Sessions & FIFO](#)
- [Dead-letter queues](#)
- [Asynchronous messaging patterns \(Architecture Center\)](#)

3. Azure RBAC & Governance

3.1 RBAC Fundamentals

Azure RBAC = **authorization system** on Azure Resource Manager. An assignment is: **security principal + role definition + scope**.

- **Security principals:** user, group, service principal, managed identity.
- **Scopes (hierarchy):** Management group → Subscription → Resource group → Resource. Assignments **inherit downward**.
- **Role definition:** JSON with Actions, NotActions, DataActions, NotDataActions, Assignable-Scopes.
 - Actions = control-plane operations (ARM).

- DataActions = data-plane operations (e.g., read blob content, receive Service Bus message).
- NotActions = subtracted from Actions (NOT a deny — just not granted).
- **Deny assignments:** Explicit deny (created by Azure via managed apps/Blueprints/deployment stacks, not directly user-creatable). Deny wins over allow.
- **Effective permissions** = union of all role assignments minus deny assignments.

3.2 Key Built-in Roles

Role	Grants
Owner	Everything + assign roles
Contributor	Everything except role assignment / sharing
Reader	View only
User Access Administrator	Manage role assignments only
Role Based Access Control Administrator	Assign roles (more constrained than UAA)
Data roles (examples)	Storage Blob Data Reader/Contributor, Key Vault Secrets User, Service Bus Data Sender/Receiver, AcrPull/AcrPush

Least privilege: Prefer narrow built-in data roles over Owner/Contributor; scope as low as possible; assign to **groups**, not users; use **PIM** for eligibility instead of standing access.

3.3 Custom Roles & ABAC

- **Custom roles:** When no built-in role fits. Define JSON; AssignableScopes limits where it can be assigned. Limit: 5,000 custom roles per tenant.
- **ABAC (attribute-based access control):** Role assignment **conditions** that filter based on resource attributes (e.g., blob tags), request attributes, or principal attributes. Currently focused on Storage data actions. Adds fine-grained control on top of RBAC.

3.4 Entra ID Roles vs. Azure RBAC (classic confusion)

- **Entra ID roles** (Global Administrator, User Administrator, Application Administrator) govern the **directory/identity plane**.
- **Azure RBAC roles** govern **Azure resources**.
- They are separate systems. A Global Administrator has NO Azure resource access by default but can **elevate access** (gain User Access Administrator at root /) in emergencies.

3.5 Managed Identities

- **System-assigned:** Lifecycle tied to the resource; 1:1; deleted with the resource.
- **User-assigned:** Standalone resource; shareable across many resources; survives resource deletion; pre-provision permissions before compute exists (better for fleets and blue-green).
- Use for: Key Vault access, Service Bus send/receive, ACR pulls, SQL auth, APIM→backend auth, ACA→anything. **No secrets to manage or rotate.**

3.6 Privileged Identity Management (PIM)

- **Just-in-time** role activation with approval workflows, MFA, justification, time-bounded assignments, access reviews, and audit history.
- Applies to both Entra ID roles and Azure RBAC roles.
- Exam signal: “reduce standing privileged access,” “time-limited access with approval” → PIM.

3.7 Governance Toolbox (AZ-305 core)

Tool	Purpose
Management groups	Hierarchy above subscriptions for policy/RBAC inheritance
Azure Policy	Enforce/audit resource compliance (deny, audit, append, modify, deployIfNotExists); initiatives = policy sets
RBAC	Who can do what
Resource locks	CanNotDelete / ReadOnly, protect critical resources
Tags	Metadata for cost/ownership/environment; enforce via Policy
Microsoft Defender for Cloud	CSPM: secure score, regulatory compliance, workload protection
Azure landing zones (CAF)	Prescriptive MG hierarchy, platform/workload subscriptions, policy-driven guardrails
Deployment stacks	Manage resource collections as a unit with deny settings (successor to Blueprints, which are deprecated)

RBAC vs. Policy: RBAC controls *who* can act; Policy controls *what/how* resources may be configured regardless of who. They complement, not overlap.

Quick Check — RBAC & Governance

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — RBAC & Governance

W1. A 200-subscription enterprise has admins with permanent Owner everywhere and no configuration guardrails. Whiteboard the target governance model and the migration steps.

Model answer: Draw the **landing zone MG hierarchy**: Root → Platform (identity/connectivity/management) + Landing Zones (corp/online) + Sandbox + Decommissioned. Attach **Policy initiatives** at MG level (allowed regions, required tags, deny public storage, Defender on). Replace standing Owner with **PIM eligible** assignments (approval + MFA + 8h max), Reader as standing access, quarterly **access reviews**; assign everything to **Entra groups**. Break-glass accounts excluded from CA. Migration: inventory assignments → map to groups/roles → move subscriptions under MGs (policies in audit mode first) → flip to enforce → remove direct assignments.

W2. An auditor asks: “Prove that a compromised web app in subscription X could not have read the HR database in subscription Y.” Whiteboard your answer using the RBAC model.

Model answer: Walk the evaluation: the app’s **managed identity** has role assignments only at X’s resource group scope (least privilege, DataActions on its own storage only). No assignment exists at any scope covering Y — RBAC is **deny-by-default**; effective permissions = union of assignments, and none exist. Add layered evidence: Y’s SQL has **private endpoint** + public access disabled (network), requires Entra auth (no SQL logins), and **Policy** denies public exposure. Show the audit trail: az role assignment list at each scope + Entra sign-in logs proving no token was ever issued for Y’s resources.

Key Resources — RBAC & Governance

- [Azure RBAC documentation](#)
- [Built-in roles reference](#)
- [Azure Policy documentation](#)
- [PIM documentation](#)
- [Managed identities overview](#)
- [Azure landing zones \(CAF\)](#)

4. Cloud Networking

4.1 Virtual Networks

- **VNet:** Isolated private network in a region; address space in CIDR; divided into **subnets**.
- **NSG (Network Security Group):** Stateful L3/L4 allow/deny rules on subnets and/or NICs; processed by priority (100–4096); default rules allow VNet + LB inbound, deny other inbound. Use **service tags** (AzureCloud, Storage, Internet) and **ASGs** (application security groups — group VMs logically so rules reference workloads, not IPs).
- **UDR (route tables):** Override system routing, e.g., 0.0.0.0/0 → Azure Firewall (forced tunneling through an NVA).
- **Peering:** Connect VNets (same or cross-region, cross-subscription). Non-transitive — hub-spoke needs the hub firewall/gateway to route spoke-to-spoke.

4.2 Hybrid Connectivity

Option	Traits
VPN Gateway (S2S)	IPsec over internet, up to ~10 Gbps aggregate, cheaper, quick to set up
ExpressRoute	Private dedicated circuit via provider, up to 100 Gbps, SLA, predictable latency; use ER + VPN failover for resilience
P2S VPN	Individual clients (certs, Entra auth)
Virtual WAN	Microsoft-managed global hub-spoke fabric: branch (SD-WAN/VPN), ER, P2S, hub firewalls, any-to-any routing at scale

4.3 Hub-Spoke Topology (default enterprise answer)

- **Hub:** shared services — Azure Firewall/NVA, VPN/ER gateways, Bastion, DNS.
- **Spokes:** workloads, peered to hub. Spoke↔spoke via hub routing.
- Benefits: centralized security/egress control, cost sharing, isolation. Scales into **Azure landing zones**; at large scale consider **Virtual WAN**.

4.4 Private Connectivity to PaaS

- **Private endpoint (Private Link):** NIC with private IP in your subnet mapped to a *specific instance* of a PaaS resource (storage account, SQL, Key Vault, Service Bus Premium, APIM...). Traffic never leaves the Microsoft backbone; works cross-region and from on-prem; enables disabling all public access. Requires **private DNS zones** (e.g., privatelink.blob.core.windows.net) so FQDNs resolve to private IPs.
- **Service endpoint:** Subnet identity extended to a PaaS *service* (not instance); traffic stays on backbone but the service keeps a public IP; no on-prem reach. Simpler/cheaper, weaker isolation.
- **Exam rule:** “eliminate public exposure / on-prem access to PaaS / data-exfiltration control” → **Private endpoints + private DNS**.

4.5 Load Balancing & Global Delivery (decision tree)

Service	Layer	Scope	Use
Azure Load Balancer	L4	Regional	TCP/UDP, VM/VMSS backends, HA ports
Application Gateway	L7	Regional	HTTP(S), TLS termination, path/host routing, WAF , AKS ingress (AGIC)
Azure Front Door	L7	Global	Anycast edge, CDN, TLS offload, WAF, path routing, multi-region failover for HTTP apps
Traffic Manager	DNS	Global	DNS-based routing (priority/weighted/performance/geographic) any protocol

Combos: **Front Door** → **App Gateway** (global edge + regional WAF/ingress), **Front Door** → **internal APIM** → **backends**. Non-HTTP global → Traffic Manager (or Front Door TCP proxying where applicable).

4.6 Security Services

- **Azure Firewall:** Stateful managed firewall: L3-L7 rules, FQDN filtering, TLS inspection/IDPS (Premium SKU), threat intelligence. Central egress control in hub.
- **WAF (on App GW/Front Door):** OWASP core rule set, bot protection, custom rules.
- **DDoS Network/IP Protection:** Enhanced mitigation + cost protection on VNet public IPs.
- **Azure Bastion:** Browser-based RDP/SSH without public IPs on VMs.
- **NAT Gateway:** Predictable outbound SNAT at scale, avoids port exhaustion.

4.7 DNS

- **Azure DNS:** Public zones.
- **Private DNS zones:** Name resolution inside VNets (+ auto-registration of VM records; link zones to VNets).
- **Azure DNS Private Resolver:** Managed inbound/outbound endpoints for hybrid DNS (replace DNS-forwarder VMs; conditional forwarding on-prem ↔ Azure).

Quick Check — Networking

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Networking

W1. Whiteboard a hub-spoke network for a company with 2 regions, on-prem Express-Route, mandatory egress inspection, and PaaS services that must never be publicly reachable. Label every routing decision.

Model answer: Per region: hub VNet (Azure Firewall Premium zone-redundant, ER gateway, Bastion, DNS Private Resolver) + workload spokes peered to hub. **UDR 0.0.0.0/0** → **firewall** on every spoke subnet (egress inspection); spoke↔spoke also via firewall (peering non-transitive). ER private peering into both hubs, **VPN failover**. PaaS: **private endpoints** in spoke PE-subnets, public access disabled by Policy, **private DNS zones** linked to all VNets, on-prem resolution via Private Resolver inbound endpoint + conditional forwarders. Cross-region: hub-to-hub global peering. Address plan: carve a /16 per region up front.

W2. Your global HTTP app is slow for Asian users and went down during a regional outage last month. Whiteboard the traffic path from user to backend that fixes both, naming each load-balancing layer.

Model answer: User → **Front Door** (anycast edge nearest the user: TLS termination, caching for static assets, WAF, health-probed **priority/latency routing** across regional origins) → regional **Application Gateway** (WAF_v2, path-based routing, zone-redundant) → AKS/ACA backends across 3 zones (internal **Azure LB** under the ingress). Asia latency: Front Door edge + caching + (optionally) an Asian region origin. Outage: Front Door probes fail → automatic origin failover. Note the exam contrast: Traffic Manager would be DNS-only failover (slower, TTL-bound) and adds no WAF/caching.

Key Resources — Networking

- [Virtual network documentation](#)
- [Hub-spoke reference architecture](#)
- [Private Link & private endpoints](#)
- [Private endpoint DNS integration](#)
- [Load-balancing options decision guide](#)
- [Azure Firewall documentation](#)
- [ExpressRoute documentation](#)

5. Azure Container Apps (ACA)

5.1 What It Is

Serverless container platform built on **AKS + KEDA + Dapr + Envoy** (abstracted away). You bring containers; Azure runs, scales (including **to zero**), and upgrades the infrastructure. No K8s API access — that's the tradeoff vs. AKS.

5.2 Compute Selection (AZ-305 favorite)

Service	Choose when
Container Apps	Microservices, event-driven jobs, APIs; want K8s-grade capabilities (scale, revisions, service discovery, Dapr) without managing K8s
AKS	Full Kubernetes API control, custom operators/CRDs, service mesh choice, complex stateful workloads
App Service	Web apps: PaaS with slots, easy custom domains, mostly HTTP
ACI	Single containers, burst/batch, simple isolated jobs, virtual nodes for AKS burst
Functions	Event-driven code with bindings, per-execution model

5.3 Core Concepts

- **Environment:** Isolation + shared VNet and Log Analytics boundary for a set of container apps. Internal (VNet-only ingress) or external. **Workload profiles:** Consumption (serverless) + Dedicated (fixed vCPU/memory, GPU) in one environment.
- **Revisions:** Immutable snapshots of app config+image. **Traffic splitting** across revisions → blue-green and canary deployments. Single vs. multiple revision mode.
- **Ingress:** Managed HTTPS ingress (Envoy), automatic TLS, session affinity, IP restrictions, custom domains, TCP ingress support. Internal-only option for private microservices.
- **Scaling:** KEDA-based scale rules — HTTP concurrency, CPU/memory, or **event-driven scalers** (Service Bus queue length, Event Hubs, Kafka, cron...). Min replicas 0 (scale to zero) → cost-efficient; set min ≥ 1 to avoid cold starts.
- **Jobs:** Run-to-completion workloads — manual, scheduled (cron), or event-driven (KEDA-triggered) — alongside always-on apps.
- **Dapr integration:** Service invocation (mTLS, retries), state stores, pub/sub (e.g., Service Bus behind Dapr), bindings, observability — enabled per app.
- **Secrets & identity:** App-level secrets, Key Vault references via **managed identity**; system- or user-assigned MI for ACR pulls (no passwords) and Azure resource access.
- **Networking:** Bring your own VNet (infrastructure subnet), internal environments behind private DNS; egress control via UDR + Azure Firewall (workload profiles environments).

5.4 Typical Enterprise Pattern

Front Door (WAF) → APIM (internal) → **Container Apps environment (internal ingress)** running APIs → managed identities → Service Bus / SQL / Key Vault via private endpoints, Dapr pub/sub for async, ACR with private endpoint for images, GitHub Actions/Azure DevOps deploying by revision with canary traffic split.

Quick Check — Container Apps

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Container Apps

W1. A startup with 8 microservices and no Kubernetes skills asks you to design their platform: private APIs, async order processing, nightly reports, minimal idle cost. Whiteboard it on ACA.

Model answer: One **internal ACA environment** in their VNet (workload profiles). HTTP APIs: internal ingress, **HTTP scale rules**, minReplicas 1 for the customer-facing API (no cold start), 0 for admin APIs. Order worker: **KEDA Service Bus scale rule**, minReplicas 0 — scales with queue depth, free when idle. Nightly reports: **ACA Job** (cron). **Dapr** pub/sub over Service Bus for service-to-service async; managed identity for ACR pulls (AcrPull) and Key Vault references. Exposure: APIM Standard v2 with VNet integration in front. Cost story: scale-to-zero everywhere except one API.

W2. The team wants zero-downtime releases with instant rollback and gradual rollout, without building custom deployment tooling. Whiteboard the release flow on ACA.

Model answer: Set the app to **multiple revision mode**. Pipeline (GitHub Actions with OIDC federation — no secrets) builds → pushes to ACR → `az containerapp update` creates revision N+1 → **traffic split** 95/5 → watch App Insights failure rate/latency for the canary label → shift 50/50 → 0/100 → deactivate old revision. Rollback = set traffic 100% to revision N (instant, it's still warm). Draw the revisions as immutable snapshots behind Envoy ingress doing weighted routing. Mention session affinity caveat for stateful clients and health probes gating the new revision.

Key Resources — Container Apps

- [Container Apps documentation](#)
- [Revisions & traffic splitting](#)
- [Scaling \(KEDA rules\)](#)
- [Dapr integration](#)
- [Networking & internal environments](#)
- [Compute service decision tree](#)

6. Authentication & Authorization Patterns (Entra ID)

6.1 Protocol Foundations

- **OAuth 2.0** = authorization (delegated access via access tokens). **OpenID Connect (OIDC)** = authentication layer on top (ID tokens).
- **Tokens:** **ID token** (who the user is, for the client), **access token** (presented to APIs; validate audience/issuer/signature/expiry/claims), **refresh token** (get new tokens silently).
- **Microsoft identity platform** = Entra ID's OAuth2/OIDC implementation + MSAL libraries.

6.2 Choose the Right Flow (exam-critical)

Scenario	Flow
Web app / SPA / mobile signing in users Service-to-service (daemon, no user)	Authorization code + PKCE Client credentials (secret or better: certificate / managed identity / workload identity federation)
API calling a downstream API on the user's behalf Input-constrained device (TV, CLI) Legacy username/password	On-Behalf-Of (OBO) Device code ROPC — avoid (blocked by MFA/CA, no modern security)
Implicit flow	Deprecated — use auth code + PKCE

6.3 App Registrations, Scopes & App Roles

- **App registration** defines the identity of an app/API in Entra; **service principal** is its instance in a tenant; **enterprise application** = the SP view.
- **Expose an API:** define **scopes** (delegated permissions, e.g., `api://orders/Orders.Read`) consumed with user context, and **app roles** (application permissions for daemons, or role claims for users/groups).
- **Consent:** delegated permissions require user/admin consent; application permissions always require **admin consent**.
- **Validate in the API:** audience (aud), issuer (iss), signature (JWKS from OpenID metadata), expiry, then authorize on scp (delegated) or roles (app permissions). scp vs roles distinction is a classic exam trap.

6.4 Gateway-Centric AuthN/Z Pattern (enterprise APIs)

1. Client obtains access token from Entra ID (auth code + PKCE or client credentials).
2. Calls APIM with Authorization: Bearer <token> (+ subscription key for identification/analytics).
3. APIM `validate-jwt / validate-azure-ad-token` checks issuer, audience, and claims; rejects at the edge (defense in depth: backend validates again — zero trust).
4. APIM policy can do fine-grained authorization (claims → operations), token exchange, or send-request introspection.

5. APIM → backend using **managed identity** (authentication-managed-identity policy) or mTLS client certificate; backend locked to APIM via network (internal VNet/private endpoint) + identity.
6. Backend → data plane (SQL, Service Bus, Storage) via **managed identity + RBAC data roles**. No connection-string secrets anywhere; Key Vault for what remains.

6.5 Other Identity Patterns

- **Managed identities everywhere:** eliminate secrets for Azure-to-Azure calls.
 - **Workload identity federation:** external workloads (GitHub Actions, other clouds, AKS pods) exchange their native tokens for Entra tokens — no stored credentials in CI/CD.
 - **Microsoft Entra External ID (successor to Azure AD B2C):** CIAM — customer sign-up/sign-in, social identities, custom branding.
 - **B2B guests:** invite partner users into your tenant; govern with entitlement management + access reviews.
 - **Conditional Access:** policy engine (signals: user, device, location, risk → require MFA, compliant device, block). Pairs with **Identity Protection** risk detection.
 - **SAS vs. Entra:** For Service Bus/Storage prefer Entra ID + RBAC; SAS only for constrained delegation scenarios (time-boxed, scoped URLs) or legacy.
 - **mTLS:** client certificates for high-assurance service-to-service (APIM validates client certs; backends require APIM's cert).
 - **Zero trust:** verify explicitly (every hop authenticates), least privilege (narrow scopes/roles, PIM), assume breach (segmentation, private networking, logging).
-

Quick Check — Auth & Identity

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Auth & Identity

W1. Whiteboard the complete token journey when a doctor uses a hospital SPA to view lab results: SPA → API gateway → results API → downstream FHIR API. Name every flow and every validation.

Model answer: SPA signs the doctor in with **auth code + PKCE** → gets ID token (for the SPA) + access token (audience = results API, scope Results.Read). SPA calls APIM with the bearer token; APIM **validate-jwt** (signature via JWKS, iss, aud, exp, scp claim) rejects at the edge. Results API **re-validates** the token (zero trust), authorizes on scp, then uses **On-Behalf-Of** to exchange it for a FHIR-API token preserving the doctor's identity (audit trail shows the human, not a service account). APIM→backend locked with managed identity/mTLS + internal network. Conditional Access enforced MFA + compliant device at sign-in.

W2. A partner's nightly batch job and your own AKS pods both need to call your inventory API. Whiteboard the identity design with zero stored secrets anywhere.

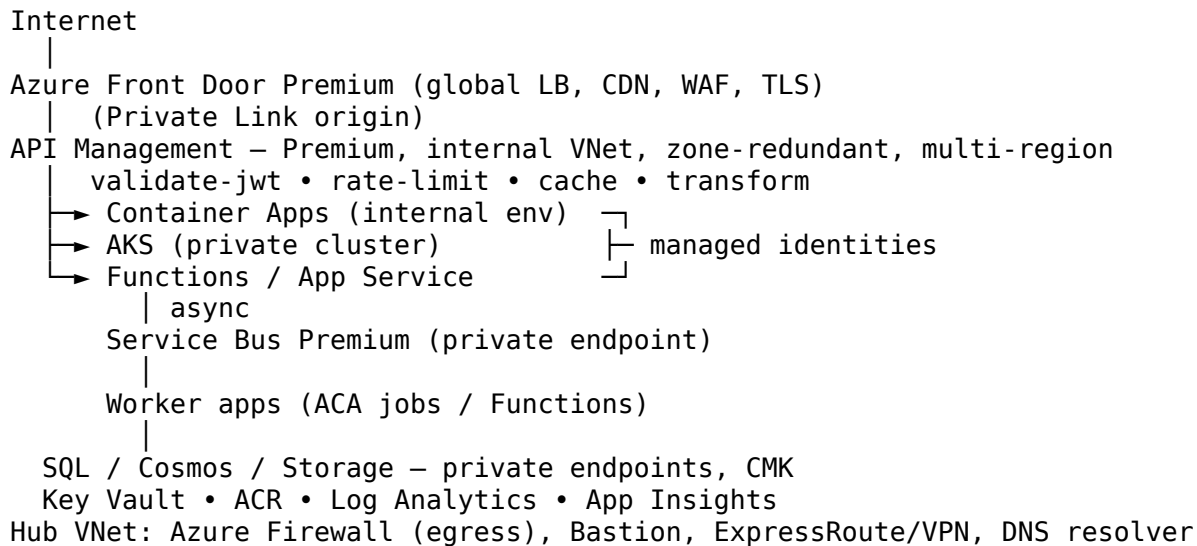
Model answer: Expose the API with **app roles** (e.g., Inventory.Sync) — application permissions with admin consent. Partner batch: their workload federates via **workload identity federation** (their cloud/CI token exchanged for your Entra tenant's token; multi-tenant app registration + admin consent) — no shared secret to rotate. AKS pods: **Entra Workload Identity** (federated service accounts) → tokens with roles: Inventory.Sync. API validates signature/iss/aud/exp then authorizes on roles (not scp — no user context). Rotate nothing; revoke by disabling the federated credential or SP. Monitor with sign-in logs per SP.

Key Resources — Auth & Identity

- [Microsoft identity platform documentation](#)
- [OAuth 2.0 / OIDC flows overview](#)
- [On-Behalf-Of flow](#)
- [Workload identity federation](#)
- [Conditional Access documentation](#)
- [Zero Trust guidance center](#)
- [APIM validate-jwt policy](#)

7. Enterprise API Architecture & Design Patterns

7.1 Reference Architecture (secure enterprise API platform)



7.2 Core Cloud Design Patterns (Azure Architecture Center)

Pattern	Problem it solves
Gateway (APIM)	Single entry: cross-cutting auth, throttling, transformation
Backends for Frontends (BFF)	Per-client-tailored APIs (mobile vs web)
Strangler Fig	Incrementally migrate a monolith by routing slices via the gateway
Gateway Offloading / Routing / Aggregation	TLS, routing, response composition at the gateway
Queue-Based Load Leveling	Buffer bursts with Service Bus between tiers
Competing Consumers	Scale-out message processing
Publisher-Subscriber	Decoupled event distribution
Claim-Check	Large payloads via Blob + message reference
Saga	Distributed transactions via compensating steps
CQRS + Event Sourcing	Separate read/write models; event log as source of truth
Retry + Circuit Breaker + Bulkhead	Transient fault handling, stop cascading failures, isolate pools
Cache-Aside	Load-on-miss caching (Redis)
Throttling	Protect services under load (APIM rate-limit/quota)

Pattern	Problem it solves
Health Endpoint Monitoring	Probes for LB/orchestrator decisions
Sidecar / Ambassador	Offload connectivity/observability (Dapr)
Anti-Corruption Layer	Isolate new systems from legacy semantics
Deployment Stamps	Repeatable scale units per tenant/region
Geode / active-active multi-region	Global low latency + resilience

7.3 Well-Architected Framework (WAF pillars — memorize)

1. **Reliability** — SLAs, redundancy (zones/regions), RTO/RPO, failure mode analysis, chaos testing.
2. **Security** — zero trust, identity as perimeter, encryption, network segmentation, Defender for Cloud.
3. **Cost Optimization** — right-size, reservations/savings plans, scale to zero, cost alerts.
4. **Operational Excellence** — IaC (Bicep/Terraform), CI/CD, observability, runbooks, safe deployment (rings, canary).
5. **Performance Efficiency** — scale out not up, caching, partitioning, async patterns, load testing.

Composite SLA math: serial components multiply ($99.95\% \times 99.9\% = 99.85\%$); redundant parallel components: $1 - (1 - A)^2$. Region pair + zone redundancy raises availability.

7.4 AZ-305 Exam Blueprint (updated April 17, 2026)

Domain	Weight
Design identity, governance, and monitoring solutions	25-30%
Design data storage solutions	25-30%
Design business continuity solutions	10-15%
Design infrastructure solutions	30-35%

- Passing score **700/1000**; ~50 questions incl. case studies; scenario/"you need to recommend" style.
- **Prerequisite for the Expert cert:** AZ-104 (Azure Administrator Associate).
- Question style: *requirements* → *best-fit service*. Learn the **decision trees** (compute, messaging, load balancing, storage, identity) and *cheapest-that-meets-requirements* logic.

7.5 Study & Practice Resources

Official (free):

- Exam page & registration: <https://learn.microsoft.com/credentials/certifications/exams/az-305/>
- Official study guide (skills measured): <https://learn.microsoft.com/credentials/certifications/resources/study-guides/az-305>
- **Free official Practice Assessment:** <https://learn.microsoft.com/credentials/certifications/exams/az-305/practice/assessment>
- MS Learn AZ-305 learning paths (Design identity/governance, storage, BC, infrastructure): <https://learn.microsoft.com/training/courses/az-305t00>
- Azure Architecture Center (patterns + reference architectures): <https://learn.microsoft.com/azure/architecture/>
- Well-Architected Framework: <https://learn.microsoft.com/azure/well-architected/>
- Cloud Adoption Framework: <https://learn.microsoft.com/azure/cloud-adoption-framework/>
- APIM docs: <https://learn.microsoft.com/azure/api-management/>
- Service Bus docs: <https://learn.microsoft.com/azure/service-bus-messaging/>
- Container Apps docs: <https://learn.microsoft.com/azure/container-apps/>
- Exam sandbox (question-format demo): <https://aka.ms/examdemo>

Practice tests & courses (paid/freemium):

- MeasureUp (Microsoft's official practice test partner): <https://www.measureup.com>
- Whizlabs AZ-305: <https://www.whizlabs.com/microsoft-azure-certification-az-305/>
- Tutorials Dojo AZ-305: <https://tutorialsdojo.com/courses/az-305-microsoft-azure-solutions-architect-practice-exams/>
- Udemy — John Savill / Scott Duffy / practice test sets: <https://www.udemy.com>
- John Savill's AZ-305 YouTube study playlist (highly recommended, free): <https://www.youtube.com/c/NTFAQO>
- ExamTopics community questions (verify answers yourself): <https://www.examttopics.com/exams/microsoft/305/>

Hands-on:

- Azure free account (\$200 credit): <https://azure.microsoft.com/free/>
- Microsoft Learn sandboxes (free, in-browser subscriptions inside modules)
- Build this guide's reference architecture yourself with Bicep — the single best prep activity.

7.6 8-Week Study Plan

Week	Focus
1	Identity & governance: Entra ID, RBAC, PIM, Policy, management groups, landing zones
2	Networking: VNets, hub-spoke, private endpoints, load-balancing decision tree, hybrid
3	Compute: ACA vs AKS vs App Service vs Functions; APIM deep dive
4	Data: SQL/Cosmos/Storage tiers, HA/DR options, analytics (Synapse/Fabric basics)
5	Messaging & integration: Service Bus, Event Grid, Event Hubs, Logic Apps; app architecture patterns
6	Business continuity: backup, ASR, RTO/RPO design, multi-region patterns
7	Monitoring + Well-Architected review; MS Learn practice assessment until consistently >85%
8	Case-study drills, MeasureUp/Tutorials Dojo full mocks, weak-area review, sit the exam

Quick Check — Architecture & Exam

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Architecture & Patterns

W1. Whiteboard the strangler-fig migration of a monolithic insurance system to microservices over 18 months, showing what exists at month 0, 6, and 18.

Model answer: Month 0: APIM (the facade) routes 100% to the monolith — deploying the gateway first is the key move. Month 6: /claims and /quotes carved out to Container Apps (highest-change, best-bounded contexts first); APIM routes those paths to new services, everything else to the monolith; an **anti-corruption layer** translates legacy models; shared data still in the legacy DB with a sync/change-feed bridge. Month 18: monolith reduced to a residual module or retired; each service owns its store (CQRS where read/write

shapes diverge); Service Bus events between services; the facade never changed from the consumer's view — that's the pattern's whole point.

W2. Your CEO asks: “Why does our 99.99% target cost 4× more than 99.9%?” Whiteboard the explanation an architect gives.

Model answer: Draw the serial chain: FD 99.99 × APIM 99.95 × app 99.95 × SQL 99.99 ≈ **99.88%** — a chain is weaker than its weakest link, so 99.9 is roughly the single-region ceiling. Getting to 99.99 means **parallel redundancy**: zone-redundant everything (small premium) then a second active region $(1-(1-A)^2)$ — doubling infrastructure, adding data replication (Cosmos multi-write or SQL FG), global routing, and the ops cost of failover testing. Show the availability ladder vs cost curve and the honest alternative: negotiate the SLO per workload tier — most workloads don't need 99.99.

Key Resources — Architecture & Patterns

- [Cloud design patterns catalog](#)
- [Azure Well-Architected Framework](#)
- [Azure Architecture Center \(reference architectures\)](#)
- [Microservices on Azure guide](#)
- [AZ-305 study guide \(official\)](#)
- [Free AZ-305 practice assessment](#)

8. Data Storage Design

8.1 Relational: Azure SQL Family

Option	Traits	Choose when
Azure SQL Database	Fully managed single database; serverless option; Hyperscale	New cloud apps, per-database scaling
SQL Managed Instance	Near-100% SQL Server compatibility (SQL Agent, cross-DB queries, CLR), VNet-native	Lift-and-shift with instance-level features
SQL Server on VMs	Full OS/engine control, bring-your-own license	Legacy dependencies, OS access, full control

Purchasing models:

- **DTU** (Basic/Standard/Premium): bundled compute+IO blend; simple, cheaper for small workloads; no serverless.
- **vCore** (General Purpose / Business Critical / Hyperscale): independent compute/storage sizing, **Azure Hybrid Benefit** (reuse SQL licenses), reserved capacity, serverless (auto-pause) in GP.

Service tiers to know:

- **General Purpose**: remote storage, 99.99%, budget default.
- **Business Critical**: local SSD + Always On replicas, lowest latency, **built-in readable secondary** (ApplicationIntent=ReadOnly), zone-redundant option.
- **Hyperscale**: up to 100+ TB, fast scale-out with named replicas, snapshot-based near-instant backups/restores.
- **Serverless (GP)**: auto-scale + **auto-pause** — pay storage only while paused; ideal for intermittent dev/test.
- **Elastic pools**: many variable-usage databases share resources — SaaS multi-tenant cost pattern.

8.2 NoSQL: Azure Cosmos DB

- **Multi-model:** NoSQL (Core/SQL API), MongoDB, Cassandra, Gremlin, Table; **PostgreSQL** flavor for distributed relational.
- **Request Units (RU/s):** normalized cost of operations. Provisioned throughput (standard or **autoscale**) or **serverless** (per-request, spiky/low traffic).
- **Partitioning:** logical partitions by **partition key** — choose high-cardinality keys matching dominant query filters; avoid hot partitions. 20 GB logical partition limit.
- **Consistency levels (strongest→weakest): Strong → Bounded Staleness → Session (default) → Consistent Prefix → Eventual.** Session = read-your-own-writes per client; the usual sweet spot. Strong across regions costs latency/availability.
- **Global distribution:** turnkey multi-region replication; **multi-region writes** for active-active write availability; automatic/manual failover.
- **SLAs:** up to 99.999% read/write with multi-region; <10 ms point reads/writes at P99.
- Extras: change feed (event sourcing/integration), TTL, unique keys, analytical store (Synapse Link / mirroring to Fabric).

8.3 Azure Storage Accounts

Services: Blob (objects), ADLS Gen2 (hierarchical namespace for analytics), Files (SMB/NFS shares), Queues, Tables.

Blob access tiers: **Hot** (frequent) → **Cool** (≥30 days, lower storage/higher access cost) → **Cold** (≥90 days) → **Archive** (≥180 days, offline, hours to rehydrate). **Lifecycle management policies** auto-tier/delete by age or last access.

Redundancy (memorize):

SKU	Copies	Survives
LRS	3 in one datacenter	Drive/rack failure
ZRS	3 across zones	Datacenter (zone) failure
GRS / RA-GRS	LRS ×2 regions (async)	Regional disaster (RA- adds read access to secondary)
GZRS / RA-GZRS	ZRS + LRS in pair region	Zone + regional failure — highest durability

Security/data protection: SSE by default (+ optional CMK, infrastructure/double encryption), immutable storage (WORM legal hold/time-based), soft delete (blob/container/share), versioning, point-in-time restore, private endpoints, SAS (prefer user-delegation), Entra RBAC data roles.

Azure Files extras: identity-based SMB auth (Entra Kerberos/AD DS), **Azure File Sync** = cloud tiering + multi-site cache of on-prem file servers.

8.4 Choosing a Data Store (decision logic)

- Relational OLTP, strong schema, joins → **Azure SQL** (MI if lift-and-shift needs instance features).
- Global low-latency, flexible schema, massive scale → **Cosmos DB**.
- Objects/files/media, data lake → **Blob / ADLS Gen2**.
- Lift-and-shift SMB shares → **Azure Files (+ File Sync)**.
- Caching/session state → **Azure Cache for Redis**.
- Warehousing/analytics → **Synapse / Fabric**; big-data engineering → **Databricks**.
- Time-series/telemetry at scale → **Data Explorer (ADX)** / Fabric Real-Time Intelligence.

Exam lens: match *consistency, latency, scale, query shape, and cost* requirements — the cheapest store that satisfies them wins.

Quick Check — Data Storage

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Data Storage

W1. A ticketing platform sells globally: 200k reads/sec on event catalogs, strict-consistency seat inventory per venue, and 7-year immutable purchase receipts. Whiteboard the polyglot data design.

Model answer: Three stores, three jobs. Catalog: **Cosmos DB**, partition key eventId, **Session** consistency, multi-region replicas near users, autoscale RU/s (+ Redis cache-aside for the hottest events). Seat inventory: needs serializable “one seat, one buyer” semantics — **Azure SQL** Business Critical zone-redundant with failover groups, or Cosmos with **Bounded Staleness + optimistic concurrency (ETags)** scoped per venue; justify whichever you pick by the consistency requirement, not fashion. Receipts: **Blob with time-based immutability (WORM)** for 7 years, lifecycle to Archive after 180 days, private endpoints throughout. Close with the exam lens: each requirement named the store.

W2. Your Cosmos bill exploded and p99 latency spiked. The container stores IoT readings with partition key = deviceType (4 values). Whiteboard the diagnosis and the fix, including how you migrate.

Model answer: Diagnose: 4 logical partitions → **hot partitions**; RU/s divides across physical partitions so the hot one throttles (429s → retries → more RUs), and 20 GB/logical-partition looms. Fix: repartition to high-cardinality key aligned to queries — deviceId (or hierarchical/synthetic key deviceId_day for time-series). Cosmos keys are immutable → **migrate via change feed**: create new container, dual-write or change-feed copy old→new, backfill, cutover reads, retire old. Add autoscale RU/s and TTL on raw readings; consider ADX for long-range analytics instead of hoarding in Cosmos.

Key Resources — Data Storage

- [Azure SQL documentation](#)
- [Cosmos DB documentation](#)
- [Cosmos partitioning guidance](#)
- [Cosmos consistency levels](#)
- [Storage account documentation](#)
- [Data store decision guide \(Architecture Center\)](#)
- [Storage redundancy explained](#)

9. Business Continuity & Disaster Recovery

9.1 Definitions

- **RTO** (Recovery Time Objective): max acceptable downtime.
- **RPO** (Recovery Point Objective): max acceptable data loss window.
- **HA** = survive component failures (zones, replicas); **DR** = survive regional failure (pair region, backups); **Backup** = point-in-time recovery from deletion/corruption/ransomware — backups are NOT DR replication and replication is NOT backup.

9.2 Azure Backup

- **Recovery Services vault** (VMs, SQL/SAP in VMs, Files, on-prem via MARS/MABS) and **Backup vault** (Blobs, Disks, PostgreSQL).
- Policies define schedule + retention (daily/weekly/monthly/yearly GFS).

- **Soft delete** (retains deleted backups 14+ days), **immutable vaults**, **multi-user authorization** — ransomware defenses.
- **Cross-region restore** with GRS vaults.
- VM backup = app-consistent snapshots; instant restore from local snapshots.

9.3 Azure Site Recovery (ASR)

- Continuous **replication** of VMs (Azure↔Azure regions, VMware/Hyper-V/physical→Azure) for regional DR.
- **Recovery plans**: ordered multi-tier failover with scripts/runbooks; **test failover** into isolated networks without impacting production (do this regularly!).
- RPO typically seconds-minutes; RTO minutes (vs hours for restore-from-backup).
- Exam split: **need low RTO/RPO regional DR → ASR; need point-in-time restore → Backup.**

9.4 Database HA/DR

Service	Mechanism
Azure SQL	Zone redundancy (BC/GP); auto-failover groups (RW + RO listener endpoints, group failover across regions); active geo-replication (per-DB readable secondaries, manual failover); PITR 1-35 days + LTR up to 10 years; geo-restore from geo-backups
Cosmos DB	Multi-region replication, automatic failover, multi-region writes (RPO≈0), continuous backup (PITR) or periodic
MySQL/PostgreSQL Flexible	Zone-redundant HA standby, read replicas (cross-region), geo-redundant backup
Storage	GRS/GZRS + customer-initiated account failover; blob object replication

9.5 Multi-Region Application Patterns

- **Active-passive (cold/warm/hot standby)**: cheaper, higher RTO; Front Door/Traffic Manager priority routing.
- **Active-active**: both regions serve traffic (weighted/performance routing); needs multi-master or partitioned data (Cosmos multi-write, or SQL failover group with read-locality).
- Sequence: define RTO/RPO per workload criticality tier → pick zone redundancy first (cheap) → add regional DR only where justified → automate failover → **test failover regularly.**
- **Chaos Studio** for resilience fault-injection experiments.

Quick Check — Business Continuity

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Business Continuity

W1. The board asks for a DR strategy for 3 workload tiers: payments (RTO 5 min/RPO 0), internal ERP (RTO 4 h/RPO 15 min), archives (RTO 72 h). Whiteboard one slide per tier with services and rough cost logic.

Model answer: Payments: **active-active** two regions — Cosmos multi-region writes (RPO≈0) or SQL failover groups + zone redundancy everywhere, Front Door health-probe routing; most expensive by far, justified only here. ERP: **warm standby** — ASR-replicated VMs (RPO minutes) + SQL failover group, scaled-down secondary, scripted recovery plan; test failover quarterly; fraction of the cost. Archives: **backup-only** — immutable, GRS-replicated vaults, restore-from-backup meets 72 h; cheapest. Draw the matrix RTO/RPO → pattern → cost; the exam (and the board) reward matching spend to tier, not gold-plating everything.

W2. Ransomware encrypted production AND the attacker had Owner rights for 6 hours. Whiteboard why your backup design still saves the company, control by control.

Model answer: Walk the attacker's attempts: delete backups → **soft delete** retains them 14+ days; purge the vault → **immutable vault policy** blocks it; disable protection/change policy → **multi-user authorization** demands a second identity via Resource Guard; encrypt data → replication copies it, but **point-in-time restores** predate infection; tamper region-wide → **GRS vault + cross-region restore**. Recovery: isolate, restore to clean point in an isolated subscription, rotate identities (PIM should have prevented 6-hour standing Owner — note it), forensics from immutable Log Analytics export. Lesson line: replication ≠ backup; immutability + MUA + JIT access is the trio.

Key Resources — Business Continuity

- [Azure Backup documentation](#)
- [Azure Site Recovery documentation](#)
- [SQL auto-failover groups](#)
- [Well-Architected reliability pillar](#)
- [Azure reliability documentation \(per-service guides\)](#)
- [Backup security features \(immutability, MUA\)](#)

10. Monitoring & Observability

10.1 Azure Monitor Stack

- **Metrics:** numeric time-series (near-real-time), metrics explorer, metric alerts.
- **Logs: Log Analytics workspace** — KQL query store for resource logs, activity logs, custom logs.
- **Application Insights:** APM — request/dependency telemetry, exceptions, **distributed tracing** (correlate across APIM→ACA→Service Bus), live metrics, **availability (web) tests**, Application Map, smart detection. Workspace-based (data lands in Log Analytics).
- **Alerts:** metric, log (KQL scheduled), activity-log, smart detection → **action groups** (email, SMS, push, webhook, Logic App, Azure Function, ITSM). **Alert processing rules** suppress/route at scale.
- **Visualize:** workbooks (interactive, parameterized), dashboards, Grafana (managed).
- **Diagnostic settings** per resource route platform logs/metrics to: Log Analytics, Storage (cheap archive), Event Hubs (SIEM/3rd-party export).

10.2 Workspace Design (exam scenario)

- **Central workspace** per environment/region = easiest cross-resource correlation + centralized RBAC. Default recommendation.
- Split workspaces for: data sovereignty per region, chargeback isolation, different retention needs.
- **Table-level RBAC / resource-context access** lets teams see only their resources' logs in a shared workspace.
- Control cost: retention settings (interactive vs long-term/archive tiers), daily cap, Basic/Auxiliary table plans, sampling in App Insights.

10.3 KQL Survival Kit

```
requests
| where timestamp > ago(1h)
| where success == false
| summarize failures = count() by name, resultCode, bin(timestamp, 5m)
| order by failures desc
```

Know: where, summarize ... by bin(), project, join, render. AZ-305 tests *choosing* the tool (metric alert vs log alert vs availability test) more than writing KQL.

10.4 Security & Health Monitoring

- **Microsoft Defender for Cloud:** CSPM (secure score, regulatory compliance) + workload protection plans (servers, storage, SQL, containers, Key Vault, APIs).
 - **Microsoft Sentinel:** cloud-native **SIEM/SOAR** on Log Analytics — connectors, analytics rules, hunting, automation playbooks. Exam signal: “correlate security events across sources / SOC” → Sentinel.
 - **Service Health** (Azure platform incidents, planned maintenance — alertable) vs **Resource Health** (your resource’s availability).
 - **Azure Advisor:** WAF-aligned recommendations (cost, security, reliability, performance, opex).
 - **Network Watcher:** NSG flow logs/VNet flow logs, connection monitor, packet capture, IP flow verify.
-

Quick Check — Monitoring

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Monitoring

W1. Whiteboard the observability design for the enterprise API platform (Front Door → APIM → ACA → Service Bus → SQL): what’s collected where, and the 5 alerts you’d ship on day one.

Model answer: One central **Log Analytics workspace**; diagnostic settings on every layer; **App Insights** (workspace-based) on APIM + ACA with **distributed tracing** so one operation ID crosses the whole chain; Service Bus/SQL metrics + logs in. Day-one alerts: (1) availability web test fails from 3 regions; (2) 5xx ratio > 2% over 15 min (log alert); (3) p95 latency breach per API (metric); (4) **Service Bus DLQ depth > 0** — silent DLQs kill order flows; (5) SQL failover/Service Health event. All fire one **action group** (on-call + Teams webhook + auto-runbook where safe). Workbook for the golden signals; cost guarded by table plans + sampling.

W2. A tenant of your multi-tenant SaaS demands access to “their logs only,” while your SOC needs everything and Finance wants per-tenant cost attribution. Whiteboard the workspace design.

Model answer: Keep **one central workspace** (SOC correlation beats fragmentation) — Sentinel on top for the SOC. Tenant access: **resource-context RBAC** — tenants get Reader on their own spoke’s resources, so queries auto-scope to their logs; or expose curated **workbooks**/exported views instead of raw workspace access. Finance: enforce a tenantId tag/custom dimension everywhere; usage queries split ingestion by resource/tag for chargeback; per-table retention keeps costs honest. Only split into separate workspaces if a regulator demands physical separation or sovereignty — name that trade-off explicitly.

Key Resources — Monitoring

- [Azure Monitor documentation](#)
- [Application Insights overview](#)
- [KQL tutorial](#)
- [Log Analytics workspace design](#)
- [Alerts & action groups](#)
- [Microsoft Sentinel documentation](#)
- [Defender for Cloud documentation](#)

11. Compute Design (Beyond Containers)

11.1 Full Compute Decision Tree

Requirement	Service
Full OS control / legacy apps	VMs (+ VMSS for scale)
Web apps, slots, minimal ops	App Service
Event-driven code, per-execution billing	Functions
Microservices without K8s ops	Container Apps
Full Kubernetes control	AKS
Single simple container	ACI
Large-scale parallel/HPC jobs	Azure Batch
VDI / desktop streaming	Azure Virtual Desktop

11.2 VMs & Scale Sets

- **Series:** B (burstable), D (general), E (memory), F (compute), L (storage), N (GPU), M (huge memory).
- **Availability:** single premium-SSD VM 99.9% → availability set (fault/update domains, one DC) 99.95% → **availability zones 99.99%**.
- **VMSS:** autoscale on metrics/schedule; **flexible orchestration** mixes VM types and spreads zones; supports Spot mix.
- **Spot VMs:** up to ~90% discount, evictable — batch/CI/stateless only.
- Disks: Ultra/Premium SSD v2/Premium/Standard SSD/HDD; host caching; ephemeral OS disks.
- Cost levers: reservations (1/3-yr), **savings plans** (flexible compute commit), Azure Hybrid Benefit (Windows/SQL licenses), right-sizing via Advisor.

11.3 Azure Functions Hosting

Plan	Traits
Consumption	Scale to zero, per-execution, 5-10 min timeout, cold starts
Flex Consumption	Faster scale, VNet integration, per-instance concurrency
Premium (Elastic)	Pre-warmed instances (no cold start), VNet, unlimited duration
Dedicated (App Service plan)	Predictable cost, reuse existing plan
Durable Functions	Stateful orchestrations: function chaining, fan-out/fan-in, async HTTP, monitor, human interaction, saga orchestration

Triggers/bindings: HTTP, Timer, Service Bus, Event Grid, Event Hubs, Cosmos change feed, Blob, Queue — the glue of event-driven architecture.

11.4 AKS Architect View

- **Control plane free/managed;** you pay for nodes. Standard/Premium tiers add uptime SLA.
 - Node pools (system/user, Spot, GPU), **cluster autoscaler** + HPA/KEDA, **private cluster** (private API server), Azure CNI vs kubenet/CNI overlay.
 - Identity: **Entra workload identity** for pods, managed identity for kubelet, Entra-integrated RBAC + Kubernetes RBAC.
 - Ingress: App Gateway for Containers / AGIC, NGINX, or service mesh (Istio add-on).
 - Ops: Fleet Manager (multi-cluster), node auto-upgrade channels, Defender for Containers, GitOps via Flux.
-

Quick Check — Compute

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Compute

W1. Map these five workloads to compute services on a whiteboard, defending each in one sentence: legacy COBOL bridge (Windows, registry hacks), customer web portal, video transcoding batch, 60-microservice platform (K8s team), monthly invoice generation triggered by queue.

Model answer: COBOL bridge → **VM** (registry/OS control; nothing else allows it) with zones or availability set. Portal → **App Service** (slots, certs, no container requirement). Transcoding → **Batch with Spot VMs** (parallel, interruptible = 90% discount; checkpoint jobs). Microservices with a real K8s team → **AKS** (they'll want CRDs/mesh; without the team it'd be ACA). Invoices → **Functions with Service Bus trigger** (per-execution billing, scale to zero) or ACA Job. Close with the decision tree: control ↔ ops burden, and 'cheapest that meets requirements' as the tiebreak.

W2. Whiteboard the cost optimization review for 300 VMs running 24/7: what do you look at, in what order, and what savings bands do you quote?

Model answer: Order: (1) **delete zombies** (Advisor idle/unattached disks/IPs — 100% saving); (2) **right-size** overprovisioned SKUs (Advisor CPU/mem data — 30-50%); (3) **schedule** dev/test off-hours (auto-shutdown — ~65% for 8x5 workloads); (4) commit: **reservations** for steady prod (~40-60% at 3-yr) vs **savings plan** where flexibility matters (~30-50%); (5) **Azure Hybrid Benefit** on Windows/SQL (up to ~40% more); (6) refactor candidates → PaaS/containers/Spot for interruptible tiers. Draw the compounding: right-size *then* reserve — reserving an oversized VM locks in waste. Governance: budgets + cost alerts + tags for ownership.

Key Resources — Compute

- [Compute decision tree \(Architecture Center\)](#)
- [Virtual machines documentation](#)
- [VMSS documentation](#)
- [Functions hosting options](#)
- [AKS baseline architecture](#)
- [Azure reservations & savings plans](#)

12. Data Integration, Analytics & Migration

12.1 Integration & Analytics Services

Service	Role	Exam signal
Azure Data Factory (ADF)	Managed ETL/ELT orchestration; 90+ connectors; mapping data flows; self-hosted integration runtime for on-prem sources	“orchestrate/copy data pipelines, hybrid sources”
Microsoft Fabric	Unified SaaS analytics: OneLake, Lakehouse/Warehouse, Data Engineering, Real-Time Intelligence, Power BI	“unified analytics platform,” successor direction to Synapse
Azure Synapse Analytics	Dedicated/serverless SQL pools, Spark, pipelines	Existing warehouse workloads (roadmap → Fabric)
Azure Databricks	First-class Spark platform: data engineering, ML, Delta Lake	“Spark/ML/lakehouse engineering”
Azure Stream Analytics	SQL-on-streams from Event Hubs/IoT Hub	“real-time SQL queries on event streams”
Azure Data Explorer / ADX	Telemetry/time-series interactive analytics (KQL)	“sub-second queries over billions of log rows”
Logic Apps	Low-code workflow integration, 1400+ connectors, B2B	“workflow across SaaS + on-prem, low-code”
Event Grid / Service Bus / Event Hubs	Messaging backbone (see §2.6)	—

Logic Apps vs Functions: declarative connector-driven workflows vs code; combine freely (Durable Functions when code-first orchestration).

12.2 Migration Framework

The 5 Rs: **Rehost** (lift-and-shift IaaS), **Refactor/Repackage** (minor changes → PaaS, e.g., SQL MI, App Service), **Rearchitect** (microservices/containers), **Rebuild** (cloud-native rewrite), **Replace** (SaaS).

Tooling:

- **Azure Migrate:** discovery + assessment (dependency mapping, right-size + cost estimates) and migration of servers, databases, web apps, VDI.
- **Database Migration Service (DMS):** online/offline moves to Azure SQL family; pair with **Data Migration Assistant** (compat assessment) and SQL best-practice assessments.
- **Storage Migration Service / AzCopy / Azure File Sync:** file server data.
- **Azure Data Box:** offline bulk transfer (TB-PB) when network transfer is impractical.

Process (CAF Migrate): assess → replicate/test → migrate (cutover) → optimize (right-size, reservations, PaaS modernization).

Quick Check — Migration & Integration

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Migration & Integration

W1. A retailer must exit its datacenter in 9 months: 400 VMs, 30 SQL Servers, a 500 TB file archive, and one core app the business wants “modernized.” Whiteboard the migration waves.

Model answer: Frame with CAF + the **5 Rs**. Wave 0 (month 1-2): **Azure Migrate** discovery + dependency mapping; build the **landing zone** first. Wave 1: low-risk **rehost** of dev/test + independent apps via Azure Migrate (momentum + learning). Wave 2: SQL via **DMA assess** → **DMS online** to SQL MI (refactor — instance features preserved, minimal downtime cutovers). Archive: **Data Box** (500 TB over WAN is months; trucks win) → Blob Cool/Archive with lifecycle. Wave 3: core app **rearchitected** to ACA + Service Bus behind APIM using strangler fig — in parallel, not blocking the exit. Wave 4: optimize (right-size, reservations, decommission). Timeline bar with the datacenter exit dependent only on waves 1-2.

W2. Whiteboard the analytics pipeline for a retailer that needs: nightly ETL from 12 on-prem systems, real-time dashboard of store sales, and data scientists exploring 5 years of history.

Model answer: Ingest: **ADF with self-hosted IR** for the 12 on-prem sources (orchestrated nightly ELT) → **ADLS Gen2** raw/curated zones (medallion). Real-time: store POS events → **Event Hubs** → **Stream Analytics** (windowed aggregates) → Power BI real-time / **Fabric Real-Time Intelligence**. Exploration: **Fabric Lakehouse/OneLake** (or Databricks for heavy Spark/ML) over the same lake — one copy of data, multiple engines. Serving: Fabric Warehouse/semantic model for BI. Governance: Purview lineage, private endpoints, managed identities end-to-end. Name why each service and its exam signal.

Key Resources — Migration & Integration

- [Cloud Adoption Framework — Migrate](#)
- [Azure Migrate documentation](#)
- [Database Migration Service](#)
- [Data Factory documentation](#)
- [Microsoft Fabric documentation](#)
- [Azure Data Box](#)

13. Platform Extras: Key Vault, Hybrid, IaC

13.1 Key Vault Deep Dive

- Objects: **secrets, keys, certificates**. Standard vs **Premium (HSM-backed keys)**; **Managed HSM** = dedicated FIPS 140-2 Level 3 pool for strict compliance.
- **Soft delete (on by default) + purge protection**: recoverable deletions; purge protection blocks permanent delete until retention lapses — required for CMK scenarios.
- Access models: **Azure RBAC (recommended)** vs legacy access policies. Data roles: Key Vault Secrets User/Officer, Crypto User, Certificates Officer, Administrator.
- Per-vault-per-app-per-environment isolation limits blast radius; private endpoints + disable public access; firewall with trusted-services bypass.
- Certificate integration: auto-renewal with integrated CAs, App Service/App GW/APIM pull certs via managed identity.
- Monitor with diagnostic logs (AuditEvent) — alert on anomalous access.
- **App Configuration** complements Key Vault: feature flags + non-secret settings with Key Vault references.

13.2 Hybrid & Multicloud: Azure Arc

- **Arc-enabled servers:** project on-prem/other-cloud machines into ARM — Policy, RBAC, tags, Defender, Monitor agents, extensions.
- **Arc-enabled Kubernetes:** GitOps config, Policy for any CNCF cluster.
- **Arc-enabled data services:** SQL MI / PostgreSQL running on your infrastructure with cloud management.
- **Azure Local (HCI):** Azure-managed hyperconverged on-prem infrastructure.
- Exam signal: “manage/govern on-prem + AWS VMs with Azure Policy like Azure resources” → **Arc**.

13.3 Infrastructure as Code & Safe Deployment

- **Bicep:** Azure-native DSL over ARM — day-0 resource support, modules, what-if preview. **Terraform:** multicloud, state-file model, huge ecosystem. Either is exam-acceptable; both beat portal clicking.
- **Deployment stacks:** manage a Bicep/ARM deployment as a unit — deny settings (block out-of-band edits) and cleanup of removed resources (Blueprints successor).
- **Template specs:** versioned, RBAC-shared templates in Azure.
- **Azure Verified Modules / landing zone accelerators:** Microsoft-maintained IaC building blocks.
- Safe deployment practice: environments promoted via pipelines (GitHub Actions/Azure DevOps with OIDC federation — no secrets), what-if/plan gates, canary rings, feature flags (App Configuration), automated rollback.
- **Azure Load Testing** (Locust/JMeter managed) + **Chaos Studio** = performance and resilience validation in CI/CD.

Quick Check — Key Vault, Hybrid & IaC

Interactive quick check available in the web app. Full Q&A with explanations: see the Practice Questions appendix.

Whiteboard Challenges — Key Vault, Hybrid & IaC

W1. Whiteboard the secrets architecture for 40 microservices across dev/test/prod: who can read what, how rotation works, and what an auditor sees.

Model answer: Vault-per-app-per-environment (blast radius) — or at minimum per-team-per-env; all with RBAC model, private endpoints, purge protection. Apps read via **managed identity + Key Vault Secrets User** scoped to their own vault; humans get PIM-eligible Officer roles in prod, standing access only in dev. Most ‘secrets’ eliminated entirely: managed identities/federation for Azure-to-Azure, so vaults hold only third-party keys + certs. Rotation: Key Vault rotation policies + Event Grid SecretNearExpiry → Function rotates; apps fetch latest version at startup/interval. Config vs secrets: **App Configuration** with KV references. Auditor sees: AuditEvent diagnostic logs per access, RBAC assignments, rotation timestamps.

W2. The platform team deploys by portal-clicking and prod drifts weekly. Whiteboard the target IaC operating model and how you stop the drift — technically, not by policy memo.

Model answer: Single Git repo of **Bicep modules** (or Terraform) = source of truth; environments as parameter files; PR review + what-if output posted to the PR; pipeline (GitHub Actions, **OIDC federation**, no secrets) promotes dev→test→prod with approvals. Drift is stopped mechanically: deploy via **deployment stacks with denyWriteAndDelete** — portal edits are rejected by the platform, not by memo; developers keep Reader in prod, PIM for break-glass (which the stack still constrains). Detection for the rest: what-if drift

checks on schedule + Policy audit. Rollback = redeploy previous commit. Load Testing + Chaos Studio as pipeline gates for the critical paths.

Key Resources — Key Vault, Hybrid & IaC

- [Key Vault documentation](#)
- [Key Vault best practices](#)
- [Azure Arc documentation](#)
- [Bicep documentation](#)
- [Deployment stacks](#)
- [App Configuration documentation](#)

14. Concepts Explained Simply (Layman's Corner)

Every core concept, explained with an everyday analogy first, then the technical takeaway. If you can explain it this simply to someone else, you understand it.

14.1 The API & Messaging World

API Management = a hotel front desk. Guests (clients) never wander into the kitchen or laundry room (your backends). They ask the front desk, which checks their room key (subscription key/JWT), enforces house rules (rate limits), answers common questions from memory (caching), and forwards requests to the right department (routing). If the kitchen moves to a new floor, guests never notice — the desk just forwards differently. *Technical takeaway: APIM decouples consumers from backend implementation and centralizes cross-cutting concerns.*

Rate limit vs quota = highway speed limit vs monthly fuel allowance. The speed limit (rate-limit) stops you going too fast right now; the fuel allowance (quota) caps total distance this month. You can obey one and still violate the other. *Both throttle, but on different time scales.*

Versions vs revisions = new edition of a book vs fixing typos. A second edition (version) changes the story — readers choose to buy it. A reprint with typo fixes (revision) silently replaces stock — nobody's reading experience breaks.

A queue = the deli counter ticket machine. Customers (messages) take a ticket and wait; any free clerk (competing consumer) serves the next ticket. Nobody is served twice, nobody is skipped, and a lunch rush just makes the line longer instead of crashing the shop (load leveling).

A topic = a magazine subscription list. The publisher prints once; every subscriber gets their own copy, and some subscribe only to the sports edition (filters).

Dead-letter queue = the post office's undeliverable-mail shelf. After several failed delivery attempts, the letter goes on the shelf with a note explaining why — a human investigates instead of the mail carrier retrying forever.

Sessions = a checkout lane dedicated to one family. All items for order #123 go through the same lane in order, even if other lanes are free. That's FIFO per session key.

Claim-check = a coat check. You don't pass a winter coat (10 MB payload) hand-to-hand through the party; you check it and pass a small ticket (blob reference).

Event vs message = a doorbell vs a courier package. An event ("someone's at the door") is a lightweight fact; whoever cares reacts. A message (the package) is a payload the sender expects someone to process — signature required.

14.2 Identity & Access

Authentication vs authorization = airport ID check vs boarding pass. ID proves who you are (authN); the boarding pass proves what you may do — board this flight, this seat, this class (authZ). Both checks happen, in that order.

Azure RBAC = office keycards. The card (role assignment) encodes who you are (principal), which doors it opens (role definition), and which building/floor it works in (scope). Facilities gives out cards per team (groups), and the master-key cabinet is behind a sign-out sheet (PIM).

PIM = the master-key sign-out sheet. Nobody carries the master key home. You sign it out with a reason, a manager approves, it auto-returns in 4 hours, and the logbook shows every borrow. *Just-in-time privileged access.*

Managed identity = an employee badge issued by the building itself. The app never carries a password in its wallet; the platform vouches for it. Nothing to steal, rotate, or accidentally commit to GitHub.

OAuth access token = a valet key. It starts your car and opens the door but not the trunk or glovebox (limited scopes), and it expires. You never hand over your real key (password).

On-Behalf-Of = a valet asking the concierge for another valet key. The restaurant valet (API A) can't use your car key at the parking garage across town (API B); he exchanges it, with proof he acts for you, at the key desk (Entra).

Zero trust = a nightclub where every door has a bouncer. Getting past the front door doesn't get you backstage. Every door re-checks ID (verify explicitly), your wristband only opens what you paid for (least privilege), and cameras run everywhere assuming someone snuck in (assume breach).

Key Vault = a hotel safe with a logbook. Valuables (secrets/keys/certs) live in the safe, not under mattresses (config files). Every opening is logged, the safe can't be thrown away while things are inside (purge protection), and staff access it with their badge (managed identity), not a shared combination.

14.3 Networking

VNet/subnets = an office floor plan. The floor (VNet) has rooms (subnets); NSGs are the door badge-readers deciding who may enter each room. ASGs are job-title stickers — “all printers,” “all accountants” — so rules say “accountants may reach printers” instead of listing room numbers.

Hub-spoke = an airport hub. Every regional flight (spoke VNet) connects through the hub, where security screening (firewall), customs (gateways), and the control tower (DNS, monitoring) are centralized. Two spokes talking = fly via the hub. Peering being non-transitive = there are no direct flights between small towns.

Private endpoint = a private elevator installed directly into your office. The public lobby entrance (public endpoint) is bricked over; the service physically appears inside your floor plan with its own room number (private IP). Even people who know the street address (FQDN) get routed to your elevator (private DNS).

Service endpoint = a staff-only side door — but the building still has a public lobby. Your subnet gets a trusted path, yet the service keeps its public address. That's why exams prefer private endpoints for “no public exposure.”

ExpressRoute vs VPN = a private rail line vs driving on the public highway in an armored truck. Both get cargo there safely (encryption), but the rail line (ER) never touches public roads, has guaranteed schedules (SLA), and carries far more freight.

Front Door vs Application Gateway vs Load Balancer vs Traffic Manager = global concierge vs building receptionist vs elevator dispatcher vs phone directory. Front Door greets the world at every city (global edge, HTTP). App Gateway manages one building's visitors in detail (regional L7, WAF). Load Balancer just sends people to the next open elevator (L4). Traffic Manager is the directory that tells you which office to call (DNS) — it never sees you walk in.

WAF = the metal detector at the entrance. The receptionist (gateway) checks appointments; the metal detector checks for weapons (SQL injection, XSS) regardless of who carries them.

14.4 Compute & Containers

VM vs App Service vs Container Apps vs AKS vs Functions = owning a house vs renting an apartment vs a serviced co-working office vs managing an office tower vs paying per meeting room hour. More control to the left, less maintenance to the right. The exam gives you a family (requirements) and asks which home fits.

Scale to zero = motion-sensor lights. The room is dark (zero cost) until someone walks in; there's a half-second flicker (cold start) as they turn on.

Revisions with traffic splitting = a restaurant testing a new recipe on 10% of tables. Complaints? Old recipe returns instantly. Praise? Roll it to every table. That's canary deployment.

KEDA scaling on queue length = a supermarket opening checkouts when lines grow. Nobody opens lane 7 because the clock says 5 pm (schedule guessing); they open it because six people are waiting (queue depth — the honest signal).

Availability zones = keeping spare car keys in different buildings. A fire in one building (datacenter) can't take out every key. Region pairs = a spare set in another city entirely (regional disaster).

14.5 Data & Recovery

Storage redundancy = photo backups. LRS: three copies in one shoebox. ZRS: copies in three rooms of the house. GRS: a copy mailed to grandma in another city (but she gets it a few minutes late — async). GZRS: three rooms AND grandma. Choose by asking "what disaster am I paying to survive?"

Blob tiers = closet / attic / storage unit / bank vault. Hot: daily clothes in the closet. Cool: winter coats in the attic. Cold: the storage unit across town. Archive: the bank vault — cheap, but retrieving takes hours and an appointment (rehydration).

Cosmos consistency levels = group-chat message delivery. Strong: nobody sees a message until everyone can (slow, perfectly synced). Session: *you* always see your own messages instantly (the default sweet spot). Eventual: everyone gets every message eventually, possibly out of order (fastest).

Partition key = choosing how to file cabinets. File customer orders by customer ID and lookups are instant; file everything under "2026" and one drawer jams while others sit empty (hot partition).

RTO vs RPO = "how long until the shop reopens?" vs "how many sales receipts did we lose?" Two different fears, two different price tags. Every DR design starts by putting numbers on both.

Backup vs replication = photo album vs a mirror. The mirror (replication) instantly shows everything — including the ketchup you just spilled on your shirt (corruption/ransomware). The album (backup) lets you go back to before the spill. You need both.

Site Recovery = a fully furnished second apartment kept in sync. Disaster strikes; you drive over and the fridge is already stocked (minutes of RTO). Backup alone = rebuilding the apartment from moving boxes (hours/days).

Composite SLA = a chain of old Christmas lights. Serial: every extra bulb is one more thing that can kill the whole string (multiply availabilities — they only go down). Parallel: two strings side by side — both must fail for darkness (availability shoots up).

14.6 Monitoring & Governance

Metrics vs logs = the car dashboard vs the trip journal. The speedometer (metrics) is instant and cheap — great for alarms. The journal (logs) records everything for later questions — "why did we detour on Tuesday?" (KQL).

Application Insights distributed tracing = a barcode that follows one package end-to-end. When a customer says “my order vanished,” you scan one ID and see every warehouse, truck, and doorstep it touched — across APIM, Container Apps, and Service Bus.

Azure Policy vs RBAC = building codes vs door keys. Keys (RBAC) decide who may build; building codes (Policy) decide what anyone may build — no wooden shacks (non-compliant SKUs) even if you own the land.

Management groups = a family tree for rules. House rules set by grandparents (root MG) automatically apply to every child and grandchild subscription; teenagers (sandbox subscriptions) get a looser curfew.

Azure Advisor = a home inspector who visits monthly. “You’re heating an empty room (idle VM), your locks are outdated (security), and your gutters need cleaning before the storm (reliability).”

15. Real-World Architecture Walkthroughs

Five end-to-end enterprise scenarios. For each: the business context, the design, why each component was chosen, how it fails safely, and where the money is saved.

15.1 FinSecure Bank — Hybrid Identity & Governance (10,000 users)

Context: A financial services firm with on-prem Active Directory migrates 10,000 users to the cloud. Regulators demand MFA, auditable privileged access, and least privilege. Budget favors minimal new infrastructure.

Design:

On-prem AD DS —(Entra Connect: password hash sync)→ Microsoft Entra ID

Conditional Access policies —————
(MFA, compliant device, block legacy) |



Management group hierarchy: Root — Platform — Landing Zones — Sandbox

| Azure Policy initiatives (allowed regions, required tags, Defender on)

| RBAC via Entra groups; PIM for Owner/Contributor; access reviews quarterly

Key decisions & why:

- **Password hash sync over ADFS/PTA:** cheapest, most resilient (auth works even if on-prem is down), no federation farm to patch. ADFS chosen only when regulators literally forbid hashes in cloud — they rarely do (it syncs a hash *of the hash*).
- **Conditional Access baseline:** require MFA for all users, block legacy authentication (the #1 password-spray vector), require compliant devices for admins.
- **PIM with approval + 4-hour activations** for Owner/Contributor/User Access Administrator; standing access is Reader-only. Access reviews auto-remove stale grants.
- **Break-glass:** two cloud-only emergency accounts, excluded from CA, monitored with alerts on every sign-in.
- **Governance:** management groups carry policy initiatives so every new subscription lands compliant on day one (landing zone pattern).

Failure modes: On-prem outage → cloud sign-in unaffected (PHS). Entra Connect server dies → sync pauses, auth continues; rebuild from config export. Admin credential theft → CA blocks unknown device, PIM means no standing privilege to steal, audit logs show the attempt.

Cost notes: PHS is free; PIM/CA need Entra ID P2 licenses for admins (scope licenses to who needs them, not all 10,000).

15.2 ShopSphere — Global E-Commerce Data Platform (GDPR)

Context: E-commerce across EU + US, 40 ms page-load target, GDPR requires EU customer data to stay in the EU. Catalog reads dominate; orders must never be lost.

Design:

Front Door Premium (WAF, caching)

└─ EU region: App tier – Cosmos DB (catalog, session) – Azure SQL (orders)
└─ US region: App tier – Cosmos DB (replica) – Azure SQL (orders)

Cosmos: multi-region write, Session consistency

SQL: separate EU/US servers (data residency), zone-redundant, failover groups intra-geo

Blob + ADLS Gen2 (EU): clickstream → lifecycle to Cool/Archive

Service Bus: order events → fulfillment workers

Key decisions & why:

- **Cosmos DB for catalog/session:** global low-latency reads with multi-region writes; **Session consistency** — a shopper always sees their own cart instantly; nobody needs Strong for a product page.
- **Azure SQL for orders:** money needs ACID transactions and joins. **Residency by architecture:** EU orders live on an EU server, US on US — GDPR compliance by design, not by policy document. Failover groups pair within the same geography (EU→EU) so DR never violates residency.
- **Partition keys:** catalog by productId, cart by sessionId — high cardinality, matches queries.
- **Claim-check + Service Bus:** order-placed messages are small; invoice PDFs go to blob. DLQ monitored with alerts — a silent DLQ is where orders go to die.
- **Lifecycle policies:** clickstream hot 7 days → Cool 30 → Archive 365 → delete at 730 (GDPR minimization).

Failure modes: Region loss → Front Door reroutes; Cosmos serves writes from surviving region (RPO≈0); SQL failover group promotes secondary within the geo (seconds of data at risk — measured, accepted RPO). WAF absorbs OWASP attacks; rate limiting at APIM protects checkout APIs. Data corruption → SQL PITR (35 days) + Cosmos continuous backup.

Cost notes: Cosmos autoscale RU/s (spiky retail traffic), Front Door caching slashes origin egress, Archive tier for clickstream, reserved capacity on the steady-state SQL tiers.

15.3 ForgeWorks — Mission-Critical HA/DR (RTO < 15 min, RPO < 1 h)

Context: Manufacturing execution system; every hour of downtime stops the production line. Two regions, strict-but-not-zero recovery targets, limited ops team.

Design:

Region A (primary, zone-redundant)

App Gateway (WAF, zonal) ← Front Door

VMSS across 3 zones (app)

Azure SQL Business Critical, ZR

ASR replication: legacy VMs

Recovery Services vault: daily backups, GRS, immutable, cross-region restore

Region B (warm standby)

App Gateway ← Front Door priority routing → App Gateway

VMSS min capacity (scaled-down)

Failover group secondary

Replica VMs (test failover quarterly)

Key decisions & why:

- **Zones first, region second:** zone redundancy handles the common failures (single DC) at near-zero design cost; the second region exists only for true regional disaster.
- **Warm standby sizing:** RTO of 15 min forbids cold rebuild-from-IaC (too slow) but doesn't require full active-active (too expensive). Minimal VMSS capacity in B scales out during failover.
- **ASR for the legacy VMs** (continuous replication, scripted recovery plans, RPO minutes); **failover groups** for SQL (listener endpoints — the app config never changes). Both comfortably beat the 1-hour RPO.

- **Backup ≠ DR:** immutable, GRS-replicated backups defend against ransomware — replication alone would faithfully replicate the encryption of your files by an attacker.
- **Quarterly test failovers** into isolated networks; an untested DR plan is a rumor, not a plan.
- **Health model:** availability tests + Service Health alerts trigger the documented failover runbook (automated via Automation runbooks; humans approve).

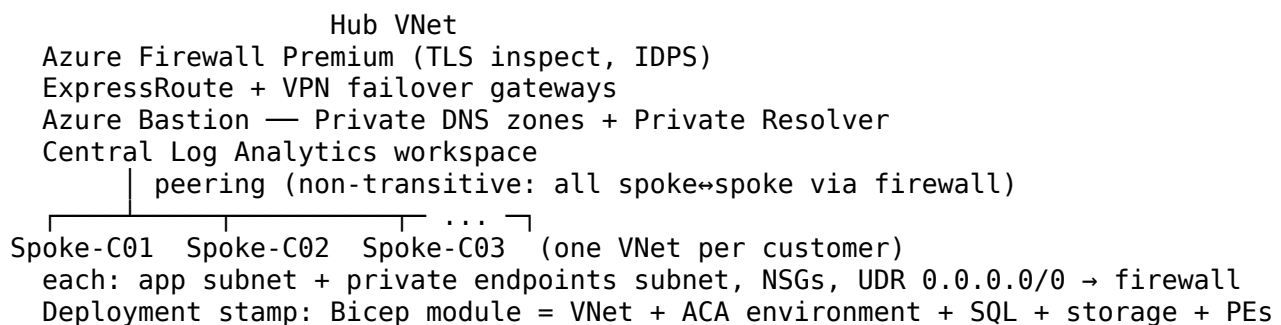
Failure modes: Zone loss → transparent (zonal VMSS + ZR SQL). Region loss → Front Door priority flips, VMSS-B scales, SQL FG fails over, ASR recovery plan boots legacy tier in order: measured RTO ~12 min. Ransomware → immutable backups + MUA restore to clean point.

Cost notes: Standby VMSS at minimal instance count; reservations on Region A steady compute; Spot for the batch analytics tier; B-series for the rarely-used jump boxes.

15.4 CloudGate SaaS — Secure Hub-Spoke for 50 Customer Environments

Context: A B2B SaaS provider hosts 50 isolated customer environments. Customers demand tenant isolation evidence; the provider wants centralized egress control, DNS, and logging without 50 copies of everything.

Design:



Key decisions & why:

- **Spoke-per-customer = deployment stamp:** hard network isolation satisfies auditors better than logical row-level isolation; one Bicep module stamps out customer N+1 in minutes.
- **All traffic through the hub firewall:** UDRs force spoke egress and spoke↔spoke through Firewall Premium — one place for TLS inspection, IDPS, FQDN allow-lists, and one set of logs.
- **Private endpoints everywhere:** each customer's SQL/storage resolves via central private DNS zones; public access disabled on every PaaS resource. Azure Policy denies public-endpoint creation — governance backs up architecture.
- **Central Log Analytics with resource-context RBAC:** provider SOC sees everything; a customer-facing support role sees only that customer's resources. Table-level retention keeps firewall logs 90 days, app logs 30.
- **Scale watch-outs:** peering and UDR limits, firewall SNAT ports, IP address plan (each spoke gets a /24 from a reserved supernet — plan the whole /16 on day one).


Failure modes: Firewall is the choke point → deployed zone-redundant; its failure mode is the top availability risk and is monitored accordingly. Compromised customer app → NSGs + firewall rules prevent lateral movement to other spokes; per-spoke managed identities scope blast radius. DDoS → Network Protection plan on hub public IPs.

Cost notes: One firewall/gateway/Bastion set shared across 50 customers is the entire point of hub-spoke: isolation without 50× the platform bill. Firewall Premium justified by compliance; logs tiered aggressively.

15.5 RetailRun — Kubernetes at Scale with GitOps

Context: A retailer runs 60 microservices on AKS. Requirements: no secrets in pipelines or pods, image provenance, zero-downtime deploys, cluster recreatable in hours.

Design:



Key decisions & why:

- **GitOps (Flux) over push deploys:** the cluster pulls its desired state from Git — the cluster is cattle; rebuilding = point Flux at the repo. Drift is auto-corrected, and the Git history is the change log auditors ask for.
- **OIDC federation for CI and Workload Identity for pods:** zero long-lived credentials anywhere in the chain. The pipeline exchanges GitHub's token for Entra tokens; pods exchange service-account tokens the same way.
- **Private everything:** private API server, ACR behind private endpoint, approved-registry policy — supply chain locked to images you built and scanned.
- **Zero-downtime:** rolling updates + PodDisruptionBudgets; blue-green at ingress for risky releases; KEDA scales order processors on Service Bus queue depth (scale on the honest signal, not CPU).
- **Spot node pool** for stateless batch (price-eviction tolerant), tainted so only batch tolerates it.

Failure modes: Zone loss → zonal node pools + PDBs keep quorum. Bad deploy → Git revert = rollback. Leaked pipeline? Nothing to leak — federation tokens are minutes-lived. Node compromise → workload identity scopes per-pod, Defender for Containers alerts, network policies limit east-west.

Cost notes: Spot pool (~70-90% off) for batch, autoscaler floor low at night, reservations for the steady system pool, ACR Premium only if geo-replication is needed.

16. Hands-On Labs (CLI + Bicep, with Validation & Troubleshooting)

Do these in an Azure free account or MS Learn sandbox. Each lab: goal → commands → validation checklist → a deliberately broken configuration to diagnose. Clean up with `az group delete` after each lab. Replace \$RG, \$LOC (e.g., westeurope) throughout.

Lab 1 — APIM: Publish, Protect, and Throttle an API

Goal: Front a public demo API with APIM, add a rate limit and JWT validation.

```
az group create -n $RG -l $LOC
az apim create -n apim-lab-$RANDOM -g $RG --publisher-email you@example.com \
  --publisher-name "Lab" --sku-name Consumption # Consumption = fast + cheap for labs

# Import a demo OpenAPI backend
az apim api import -g $RG --service-name <apim-name> --path conference \
  --specification-url https://conferenceapi.azurewebsites.net?format=json \
  --specification-format OpenApiJson --api-id conf
```

Portal → API → All operations → Inbound policy:

```
<inbound>
<base />
```

```
<rate-limit-by-key calls="5" renewal-period="30"
  counter-key="@context.Request.IpAddress)" />
</inbound>
```

Validate: ☐ Calling the API 6x in 30 s from one IP returns **429** on the 6th. ☐ Trace (Ocp-Apim-Trace) shows the policy firing. ☐ A subscription key is required (or API set to open intentionally).

Troubleshoot this: A teammate reports every caller shares one rate-limit bucket. *Find 3 issues:* (1) rate-limit used instead of rate-limit-by-key; (2) counter-key set to a constant string; (3) policy placed at product scope after <base/> of an API that already short-circuits with return-response.

Lab 2 — Service Bus: Queues, DLQ, and Sessions

```
az servicebus namespace create -n sb-lab-$RANDOM -g $RG -l $LOC --sku Standard
az servicebus queue create -n orders -g $RG --namespace-name <ns> \
  --max-delivery-count 3 --enable-dead-lettering-on-message-expiration true \
  --default-message-time-to-live PT5M
az servicebus queue create -n orders-fifo -g $RG --namespace-name <ns> \
  --enable-session true
# Send/receive without keys: assign yourself data roles
az role assignment create --assignee <your-upn> \
  --role "Azure Service Bus Data Owner" --scope <namespace-resource-id>
```

Send test messages with the portal's Service Bus Explorer: send 1 message, receive in **peek-lock** and abandon it 3 times.

Validate: ☐ After 3 abandons the message appears in the **DLQ** with reason MaxDeliveryCountExceeded. ☐ Messages sent to orders-fifo without a SessionId are rejected. ☐ Two messages with SessionId A are received in order.

Troubleshoot this: Consumers process duplicates and messages vanish under load. *Find 3 issues:* (1) receive-and-delete mode used, so crashes lose messages; (2) lock duration 30 s but processing takes 90 s with no lock renewal → redelivery duplicates; (3) TTL 5 min silently expires backlog — no one monitors the DLQ.

Lab 3 — RBAC & Managed Identity: Zero-Secret Storage Access

```
az group create -n $RG -l $LOC
az storage account create -n stlab$RANDOM -g $RG --sku Standard_LRS --min-tls-version TLS1_2
az vm create -n vm-lab -g $RG --image Ubuntu2204 --size Standard_B1s \
  --assign-identity --generate-ssh-keys # system-assigned MI
# Grant ONLY data-plane read at container scope (least privilege)
az role assignment create \
  --assignee <vm-principal-id> \
  --role "Storage Blob Data Reader" \
  --scope "<storage-id>/blobServices/default/containers/docs"
```

On the VM: get a token from IMDS and read a blob — no keys anywhere:

```
curl -s 'http://169.254.169.254/metadata/identity/oauth2/token?api-version=2018-02-01&resource='
```

Validate: ☐ Blob GET with the token succeeds. ☐ Blob PUT fails (Reader ≠ Contributor). ☐ az role assignment list --scope <storage-id> shows no broader grants. ☐ Storage account keys were never used.

Troubleshoot this: The app on the VM gets 403s. *Find 3 issues:* (1) role assigned at the wrong scope (another container); (2) role is Reader (control plane) not Storage Blob **Data** Reader; (3) token requested for the wrong resource/audience (<https://management.azure.com/>).

Lab 4 — Hub-Spoke with Private Endpoint & Private DNS

```
az network vnet create -n vnet-hub -g $RG --address-prefix 10.0.0.0/16 -l $LOC \
--subnet-name AzureFirewallSubnet --subnet-prefix 10.0.1.0/26
az network vnet create -n vnet-spoke -g $RG --address-prefix 10.1.0.0/16 -l $LOC \
--subnet-name app --subnet-prefix 10.1.1.0/24
az network vnet peering create -n hub-to-spoke -g $RG --vnet-name vnet-hub \
--remote-vnet vnet-spoke --allow-vnet-access
az network vnet peering create -n spoke-to-hub -g $RG --vnet-name vnet-spoke \
--remote-vnet vnet-hub --allow-vnet-access

# Private endpoint for a storage account into the spoke
az storage account create -n stpe$RANDOM -g $RG --public-network-access Disabled
az network private-endpoint create -n pe-blob -g $RG --vnet-name vnet-spoke \
--subnet app --connection-name blobconn \
--private-connection-resource-id <storage-id> --group-id blob
az network private-dns zone create -g $RG -n privatelink.blob.core.windows.net
az network private-dns link vnet create -g $RG -n link-spoke \
--zone-name privatelink.blob.core.windows.net --virtual-network vnet-spoke \
--registration-enabled false
az network private-endpoint dns-zone-group create -g $RG \
--endpoint-name pe-blob -n zg --private-dns-zone privatelink.blob.core.windows.net \
--zone-name blob
```

Validate: ☐ From a VM in the spoke, nslookup <account>.blob.core.windows.net returns a **10.1.1.x** address. ☐ Blob access works from the VNet, fails from the internet. ☐ Peering state shows **Connected** both ways.

Troubleshoot this: VM resolves the storage FQDN to a public IP. *Find 3 issues:* (1) private DNS zone not linked to the VM's VNet; (2) DNS zone group missing so no A record was created; (3) VM uses custom DNS servers that don't forward to Azure DNS (168.63.129.16).

Lab 5 — Container Apps: Revisions, Traffic Split, KEDA Scale on Service Bus

```
az containerapp env create -n aca-env -g $RG -l $LOC
az containerapp create -n web -g $RG --environment aca-env \
--image mcr.microsoft.com/k8se/quickstart:latest \
--ingress external --target-port 80 \
--min-replicas 0 --max-replicas 5 --revisions-mode multiple

# New revision + canary
az containerapp update -n web -g $RG --set-env-vars VERSION=v2
az containerapp ingress traffic set -n web -g $RG \
--revision-weight <rev1>=90 <rev2>=10
```

Scale a worker on queue length

```
az containerapp create -n worker -g $RG --environment aca-env \
--image <your-worker-image> --min-replicas 0 --max-replicas 10 \
--scale-rule-name sbqueue --scale-rule-type azure-servicebus \
--scale-rule-metadata queueName=orders namespace=<ns> messageCount=5 \
--scale-rule-auth connection=sb-conn --secrets sb-conn=<connection-string>
```

Validate: ☐ az containerapp revision list shows two revisions with 90/10 weights; ~10% of curls return v2. ☐ With 0 messages the worker shows **0 replicas**; enqueue 50 messages → replicas grow; drain → back to 0. ☐ App scales to zero after idle (check replica count, note the cold-start on next request).

Troubleshoot this: The canary never receives traffic and the worker never scales. *Find 3 issues:* (1) app left in **single** revision mode so weights are ignored; (2) scale rule references the wrong queue

name; (3) min-replicas set to 1 masks scale-to-zero claims while the KEDA auth secret is invalid (check az containerapp logs).

Lab 6 — Entra ID: Protect an API with App Roles (validate-jwt end-to-end)

```
# API app registration exposing a role
az ad app create --display-name orders-api
az ad app update --id <api-app-id> --identifier-uri api://orders-api
# add appRole "Orders.Read" via portal (Manifest) or Graph
# Client credentials grant for a daemon client:
az ad app create --display-name orders-client
az ad app credential reset --id <client-app-id> # secret for lab only
# Admin-consent the client's application permission to Orders.Read
```

Get a token and inspect it:

```
curl -X POST https://login.microsoftonline.com/<tenant>/oauth2/v2.0/token \
  -d "client_id=<client-app-id>&client_secret=<secret>" \
  -d "grant_type=client_credentials&scope=api://orders-api/.default"
# paste access_token into jwt.ms → check aud + roles claims
```

Then attach the validate-jwt policy from §1.4 to the Lab 1 APIM API requiring roles contains Orders.Read.

Validate: ☐ Token shows aud: api://orders-api and roles: ["Orders.Read"]. ☐ APIM returns 401 with no/expired token, 401 with wrong audience, 200 with the right one. ☐ In real projects the client uses **managed identity/federation**, not the lab secret.

Troubleshoot this: Every token is rejected with 401. *Find 3 issues:* (1) policy audience checks the client app ID instead of the API's; (2) openid-config URL uses the wrong tenant; (3) client requested scope=api://orders-api/Orders.Read with client_credentials — app-only flows must use /.default.

Lab 7 — Backup & DR Drill: SQL PITR + Storage Recovery

```
az sql server create -n sql-lab-$RANDOM -g $RG -l $LOC -u labadmin -p '<strong-pw>'
az sql db create -n appdb -g $RG -s <server> --service-objective GP_S_Gen5_1 \
  --backup-storage-redundancy Zone
# create a table, insert rows, note the time, then DROP the table (the "disaster")
az sql db restore -g $RG -s <server> -n appdb --dest-name appdb-restored \
  --time "2026-07-17T10:30:00Z" # restore to before the drop
```

Storage protection drill:

```
az storage account blob-service-properties update -n <account> -g $RG \
  --enable-delete-retention true --delete-retention-days 7 \
  --enable-versioning true
# upload a blob, overwrite it, delete it – then recover both the version and the deleted blob
```

Validate: ☐ Restored DB contains the dropped table (restore is to a **new** database — plan the app cutover). ☐ Deleted blob is recoverable within retention; the pre-overwrite version restores. ☐ You know your measured RTO for this restore — write it down; that number *is* your DR documentation.

Troubleshoot this: A real incident recovers nothing. *Find 3 issues:* (1) soft delete enabled *after* the deletion happened; (2) restore attempted to the same DB name (unsupported) losing time; (3) retention 7 days but the corruption was noticed on day 9 — retention must exceed detection time.

Lab 8 — IaC: Bicep + Deployment Stack with Deny Settings

```
// main.bicep – a governed stamp
param location string = resourceGroup().location
```

```
resource sa 'Microsoft.Storage/storageAccounts@2023-05-01' = {
  name: 'stack${uniqueString(resourceGroup().id)}'
  location: location
  sku: { name: 'Standard_ZRS' }
  kind: 'StorageV2'
  properties: { minimumTlsVersion: 'TLS1_2', allowBlobPublicAccess: false }
}

az deployment group what-if -g $RG --template-file main.bicep # preview first
az stack group create -n lab-stack -g $RG --template-file main.bicep \
  --deny-settings-mode denyWriteAndDelete --action-on-unmanage deleteResources
# Now try to change the account in the portal → blocked by the stack's deny assignment
```

Validate: ☐ what-if showed the plan before deploy. ☐ Portal edit of the storage account is denied.
☐ Removing the resource from the template and updating the stack deletes it (managed lifecycle).
☐ az stack group show lists managed resources.

Troubleshoot this: Drift keeps appearing in prod. *Find 3 issues:* (1) stack created with --deny-settings-mode none; (2) an Owner is in the deny-settings excluded principals; (3) hotfixes applied in the portal instead of via the pipeline — the process, not the tool, is broken.

Lab 9 — SQL Failover Group + Storage Lifecycle

```
az sql server create -g rg-data -n sqllab9-pri -l westeurope -u azadmin -p '<pwd>'
az sql server create -g rg-data -n sqllab9-sec -l northeurope -u azadmin -p '<pwd>'
az sql db create -g rg-data -s sqllab9-pri -n appdb --service-objective S0
az sql failover-group create -g rg-data -s sqllab9-pri -n fg-lab9 \
  --partner-server sqllab9-sec --add-db appdb --failover-policy Automatic --grace-period 1

# Storage lifecycle: cool after 30d, archive 90d, delete 365d
az storage account management-policy create --account-name <acct> -g rg-data --policy '{
  "rules":[{"name":"tier","enabled":true,"type":"Lifecycle","definition":{"
    "filters":{"blobTypes":["blockBlob"]},
    "actions":{"baseBlob":{"
      "tierToCool":{"daysAfterModificationGreaterThan":30},
      "tierToArchive":{"daysAfterModificationGreaterThan":90},
      "delete":{"daysAfterModificationGreaterThan":365}}}}}]}'
```

Validate: ☐ Connect via fg-lab9.database.windows.net (the listener, not the server name). ☐
 az sql failover-group set-primary on the secondary completes and the same connection string still works. ☐ Lifecycle policy appears in az storage account management-policy show.

Troubleshoot this: (a) the app breaks after failover; (b) blobs never move to Cool; (c) reading an archived blob fails. *Answers:* (a) the app connects to sqllab9-pri... directly instead of the failover-group listener; (b) last-modified dates too recent / the policy runs ~daily — wait a cycle, or filters exclude the container prefix; (c) archived blobs must be rehydrated before read — that's by design.

Lab 10 — Monitoring & Alerting

```
az monitor log-analytics workspace create -g rg-mon -n law-lab10
az monitor diag-settings create --resource <apimOrAppResourceId> -n ds-lab10 \
  --workspace law-lab10 --logs '[{"categoryGroup":"allLogs","enabled":true}]' \
  --metrics '[{"category":"AllMetrics","enabled":true}]'
az monitor action-group create -g rg-mon -n ag-oncall --short-name oncall \
  --action email admin.you@example.com
az monitor scheduled-query create -g rg-mon -n alert-5xx \
  --scopes <workspaceId> --condition "count > 5" \
  --condition-query "requests | where success == false | where timestamp > ago(15m)" \
  --evaluation-frequency 5m --window-size 15m --action-groups ag-oncall
```


Validate: ☐ KQL requests | summarize count() by resultCode returns data. ☐ Forced failures trigger the alert and the email arrives. ☐ App Insights Application Map shows the dependency chain.

Troubleshoot this: (a) no data in the workspace; (b) the alert never fires though failures occur; (c) workspace cost spikes. *Answers:* (a) diagnostic settings missing/pointed elsewhere, or 5-10 min ingestion latency; (b) query window/frequency mismatch or threshold too high — test the KQL manually first; (c) verbose categories (allLogs on chatty resources) — switch tables to Basic plan, add sampling, set a daily cap.

Capstone: rebuild the §7.1 reference architecture end-to-end with Bicep + GitHub Actions OIDC. If you can explain every resource's purpose to a colleague, you're exam-ready for the design questions.

Lab habit for the exam: after every lab ask the AZ-305 question — “*which requirement (cost, RTO, isolation, least privilege) did each flag serve?*” The exam tests the why, not the syntax.

17. One-Page Cheat Sheets

17.1 Identity & Security

Need	Answer
User sign-in flow (web/SPA/mobile)	Auth code + PKCE
Daemon	Client credentials (cert/MI/federation > secret)
API → downstream API as user	On-Behalf-Of
No secrets Azure↔Azure	Managed identity + RBAC data roles
No secrets CI/CD	OIDC workload identity federation
JIT privileged access	PIM (eligible, approval, MFA, reviews)
Enforce MFA/device/location	Conditional Access (+ Identity Protection risk)
Consumers (CIAM)	Entra External ID · Partners → B2B guests
Delegated vs app permissions	scp claim vs roles claim; app perms need admin consent
Secrets/keys/certs	Key Vault (RBAC model, soft delete + purge protection, PE); FIPS L3 → Managed HSM

17.2 Storage & Data

Need	Answer
Redundancy ladder	LRS → ZRS → GRS/RA-GRS → GZRS/RA-GZRS
Tiers	Hot / Cool(30d) / Cold(90d) / Archive(180d, offline) + lifecycle policies
WORM compliance	Immutable storage (time-based / legal hold)
Data lake	ADLS Gen2 (hierarchical namespace)
Lift-and-shift SQL w/ Agent, cross-DB	SQL Managed Instance
100+ TB, instant restore	Hyperscale
Intermittent dev/test DB	Serverless (auto-pause)
Multi-tenant SaaS DBs	Elastic pools
Global NoSQL, RPO≈0 writes	Cosmos DB multi-region writes; Session consistency default
Cache	Azure Cache for Redis (cache-aside)
Telemetry analytics (KQL)	Azure Data Explorer / Fabric RTI

17.3 Compute & Containers

Need	Answer
Decision ladder	VM → App Service → Functions → ACA → AKS → ACI → Batch
VM SLA ladder	99.9% single → 99.95% avail. set → 99.99% zones
Cheap interruptible	Spot · Steady 24/7 → reservations · Flexible commit → savings plan
Functions no cold start + VNet	Premium plan
Stateful orchestration	Durable Functions (fan-out/fan-in, saga)
Microservices, no K8s ops	Container Apps (KEDA scale-to-zero, revisions, Dapr)
Full K8s control	AKS (private cluster, workload identity, GitOps)
APIM tiers	Consumption (serverless) · Premium (VNet, multi-region, zones, self-hosted GW) · v2 = faster + VNet integration

17.4 Networking

Need	Answer
LB decision	L4 regional: LB · L7 regional+WAF: App GW · L7 global HTTP: Front Door · DNS any-protocol: Traffic Manager
Private PaaS	Private endpoint + private DNS zone (service endpoint = weaker, no on-prem)
Topology	Hub-spoke (firewall, Bastion, gateways, DNS in hub); big scale → Virtual WAN
Hybrid	VPN (internet, cheap) · ExpressRoute (private, SLA) · ER+VPN failover
Egress control	UDR 0.0.0.0/0 → Azure Firewall (Premium = TLS inspect/IDPS)
Peering	Non-transitive; global peering crosses regions/subs
Admin access	Bastion (no public IPs) · SNAT scale → NAT Gateway
Hybrid DNS	DNS Private Resolver (conditional forwarding both ways)

17.5 BC/DR & Monitoring

Need	Answer
Definitions	RTO = downtime cap · RPO = data-loss cap
Ladder	Zones (cheap, 99.99%) → multi-region (priority/active-active via FD/TM)
VM regional DR, low RTO/RPO	ASR + recovery plans + test failovers
Point-in-time recovery	Azure Backup (soft delete, immutable vault, MUA vs ransomware)
SQL cross-region	Auto-failover groups (listener = no conn-string change); LTR up to 10 yrs
Storage DR	(RA-)G(Z)RS + customer-initiated failover (Last Sync Time loss)

Need	Answer
Composite SLA Logs vs metrics	Serial: multiply · Parallel: $1 - (1 - A)^2$ Log Analytics (KQL, log alerts) vs metrics (fast alerts)
APM/tracing Export	App Insights (availability tests, App Map) Diagnostic settings → LA / Storage (archive) / Event Hubs (SIEM)
SIEM/SOAR	Sentinel · Posture: Defender for Cloud · Recs: Advisor

18. Practice Questions (Q&A)

293 exam-style questions. Answers immediately follow each question.

API Management

Q1. Your company needs APIM deployed into a VNet so the gateway is reachable only from internal networks. Which tier supports internal VNet injection?

- A. Consumption
- B. Standard
- C. Premium
- D. Basic v2

Answer: C. VNet injection (external or internal mode) requires Premium (or Premium v2). Standard v2 only supports outbound VNet integration.

Q2. Which APIM component do API consumers use to discover APIs, read docs, test calls, and obtain subscription keys?

- A. Management plane
- B. API gateway
- C. Developer portal
- D. Azure portal

Answer: C. The developer portal is the auto-generated, customizable site for API discovery, interactive docs, and key management.

Q3. You must introduce a breaking change to an API while existing consumers continue using the current contract. What should you use?

- A. A revision
- B. A version
- C. A new product
- D. A named value

Answer: B. Versions are for breaking changes and are explicitly selected by consumers. Revisions are for non-breaking changes.

Q4. You want to test a non-breaking change to a production API before making it live. What is the safest APIM-native approach?

- A. Deploy a second APIM instance
- B. Create a revision and test via ;rev=2 URL, then set it current
- C. Create a new version v2
- D. Edit the live API directly

Answer: B. Revisions allow safe iteration; access a non-current revision via ;rev=N and promote it when ready.

Q5. In what order are APIM policy sections executed for a successful request?

- A. backend, inbound, outbound
- B. inbound, backend, outbound
- C. outbound, inbound, backend
- D. inbound, outbound, backend

Answer: B. Pipeline: inbound → backend → outbound, with on-error running if any section throws.

Q6. Which policy enforces a monthly call volume cap per subscription?

- A. rate-limit
- B. quota
- C. ip-filter
- D. cache-lookup

Answer: B. quota enforces long-term limits (calls or bandwidth per renewal period); rate-limit handles short-term burst throttling.

Q7. What does the element do in a policy definition?

- A. Sets the backend base URL
- B. Inherits and positions policies from the enclosing (broader) scope
- C. Resets all policies
- D. Defines the default error response

Answer: B. injects the parent scope's policies at that position, controlling execution order across Global→Product→API→Operation scopes.

Q8. Which policy validates an Entra ID-issued bearer token, checking issuer, audience, and claims via OpenID configuration metadata?

- A. check-header
- B. validate-jwt
- C. authentication-basic
- D. validate-parameters

Answer: B. validate-jwt validates JWTs using the OpenID config endpoint for signing keys; validate-azure-ad-token is the Entra-specific variant.

Q9. APIM must call a backend App Service without any stored credentials. What is the best approach?

- A. Basic auth with Key Vault password
- B. authentication-managed-identity policy with the APIM managed identity
- C. Subscription key forwarding
- D. Self-signed client certificate

Answer: B. The authentication-managed-identity policy acquires a token for the backend using APIM's managed identity — no secrets.

Q10. Which tier supports deploying APIM gateways to multiple Azure regions from one instance?

- A. Standard
- B. Premium
- C. Basic
- D. Consumption

Answer: B. Multi-region deployment is a Premium feature; the primary region hosts the management plane.

Q11. You need the APIM gateway to run on-premises next to legacy backends while keeping management in Azure. What do you deploy?

- A. Azure Arc VM
- B. Self-hosted gateway container
- C. ExpressRoute
- D. Consumption tier

Answer: B. The self-hosted gateway is a containerized gateway for on-prem/other clouds, federated to the Azure management plane.

Q12. What are subscription keys best treated as in a secure enterprise API design?

- A. Sole authentication mechanism
- B. Consumer identification/analytics mechanism, combined with OAuth 2.0
- C. Encryption keys
- D. Backend credentials

Answer: B. Keys identify consumers for quotas/analytics; real authentication should use OAuth 2.0/JWT validation.

Q13. Which construct bundles APIs with terms of use and visibility, and can require subscription approval?

- A. Backend
- B. Product
- C. Named value
- D. Workspace

Answer: B. Products package APIs for consumers with visibility, terms, approval workflow, and product-scope policies.

Q14. Where should APIM policy secrets like API keys for backends be stored?

- A. Inline in policy XML
- B. Named values referencing Key Vault secrets
- C. Developer portal
- D. Subscription keys

Answer: B. Named values can be Key Vault-backed, keeping secrets out of policy definitions with rotation support.

Q15. A retail platform wants product teams to independently manage their own APIs and subscriptions in one Premium APIM instance with runtime isolation. What feature fits?

- A. Products
- B. Workspaces
- C. Revisions
- D. Groups

Answer: B. Workspaces give teams isolated management and dedicated workspace gateways for federated API management.

Q16. Which policy would you use to return a stubbed response so client teams can develop before the backend exists?

- A. mock-response
- B. send-request
- C. return-response only on error
- D. rewrite-uri

Answer: A. mock-response returns sample/schema-based responses, enabling API-first development.

Q17. What is the correct enterprise pattern for exposing an internal-VNet APIM globally with WAF protection?

- A. Public IP directly on APIM
- B. Front Door (or App Gateway) with WAF in front of internal APIM
- C. Peering to every consumer VNet
- D. Service endpoints

Answer: B. Internal APIM stays private; Front Door/App GW terminates public traffic, applies WAF, and forwards privately.

Q18. Which APIM feature helps reduce backend load for repeated identical GET requests?

- A. quota
- B. cache-lookup/cache-store policies
- C. validate-jwt
- D. set-variable

Answer: B. Response caching (built-in or external Redis) serves repeats from cache.

Q19. Standard v2 differs from classic Standard primarily by offering what networking capability?

- A. Inbound private endpoint only
- B. Outbound VNet integration to reach private backends
- C. Full VNet injection
- D. ExpressRoute termination

Answer: B. Standard v2 supports VNet integration (outbound) so the gateway can call private backends; injection needs Premium/Premium v2.

Q20. Which policy would enforce per-consumer throttling keyed on the caller's IP or JWT subject rather than subscription?

- A. rate-limit
- B. rate-limit-by-key
- C. quota
- D. ip-filter

Answer: B. rate-limit-by-key lets you define the counter key via expression (IP, claim, header, etc.).

Q21. What happens when rate-limit is exceeded?

- A. Requests queue at the gateway
- B. 429 Too Many Requests returned
- C. Request forwarded with warning header
- D. Subscription suspended

Answer: B. The gateway rejects excess calls with 429 and a Retry-After header.

Q22. For AI/LLM backends behind APIM, which policies manage token consumption and reuse semantically similar completions?

- A. quota and cache-lookup
- B. llm-token-limit and llm-semantic-cache
- C. rate-limit and cors
- D. validate-jwt and set-header

Answer: B. The AI gateway capabilities include llm-token-limit (token quotas) and llm-semantic-cache.

Q23. Which deployment practice manages APIM configuration as code across dev/test/prod instances?

- A. Manual portal edits

- B. APIOps with Git + extractor/publisher pipelines
- C. Copying policies via clipboard
- D. Developer portal export

Answer: B. APIOps applies GitOps to APIM: extract configuration, review via PRs, publish to environments.

Service Bus

Q24. Which Service Bus entity delivers each message to exactly one of several competing receivers?

- A. Topic
- B. Queue
- C. Subscription
- D. Relay

Answer: B. Queues are point-to-point: competing consumers each lock and process distinct messages.

Q25. Publishers send one message; three independent systems must each receive a copy with different filtering. What do you use?

- A. Three queues and client-side fan-out
- B. A topic with three filtered subscriptions
- C. Event Hubs partitions
- D. A single queue with sessions

Answer: B. Topics fan out to subscriptions; SQL/correlation filters give each subscriber only relevant messages.

Q26. Which receive mode guarantees at-least-once delivery?

- A. Receive-and-delete
- B. Peek-lock
- C. Prefetch
- D. Auto-complete off only

Answer: B. Peek-lock locks the message until the consumer completes it; failures cause redelivery — at-least-once.

Q27. Messages that exceed MaxDeliveryCount are automatically moved to...

- A. An Azure Storage blob
- B. The dead-letter queue
- C. A retry topic
- D. Deleted permanently

Answer: B. The DLQ captures poison messages for inspection/reprocessing; monitor it actively.

Q28. You need strict FIFO processing of messages belonging to the same order ID. What feature do you enable?

- A. Duplicate detection
- B. Sessions with SessionId = order ID
- C. Partitioning
- D. Scheduled delivery

Answer: B. Sessions guarantee ordered, single-consumer handling per session key.

Q29. How do you make enqueueing idempotent when producers might retry sends?

- A. Sessions
- B. Duplicate detection with a deterministic MessageId

- C. Prefetch
- D. TTL

Answer: B. Duplicate detection drops repeats of the same MessageId within the configured window.

Q30. Which tier is required for private endpoints, availability zone replication, Geo-DR, and 100 MB messages?

- A. Basic
- B. Standard
- C. Premium
- D. Any tier

Answer: C. Premium provides dedicated messaging units, network isolation, zone redundancy, Geo-DR, and large messages.

Q31. What does Service Bus Geo-Disaster Recovery replicate?

- A. Messages and metadata
- B. Metadata only (entities/config), not messages
- C. Messages only
- D. Nothing until failover

Answer: B. Classic Geo-DR pairs namespaces and replicates metadata only — in-flight messages are not copied. A key exam fact.

Q32. A producer bursts to 10x normal load while consumers process steadily. Which pattern does a queue between them implement?

- A. Circuit breaker
- B. Queue-based load leveling
- C. Cache-aside
- D. Sidecar

Answer: B. The queue buffers bursts so consumers process at sustainable rates — load leveling.

Q33. You must send a 250 MB payload through Service Bus Standard. Best practice?

- A. Enable large message support
- B. Claim-check: store payload in Blob Storage, send a reference message
- C. Split into 1000 messages
- D. Switch to Event Grid

Answer: B. Claim-check keeps messages small; consumers fetch the payload from storage via the reference.

Q34. Which is the preferred authentication for applications sending to Service Bus?

- A. SAS keys in app config
- B. Entra ID with managed identity and Azure Service Bus Data Sender role
- C. Connection string in code
- D. Anonymous with IP filter

Answer: B. Managed identity + narrowly scoped data-plane RBAC roles avoids secret management.

Q35. An IoT platform ingests millions of telemetry events per second for stream analytics. Which service fits better than Service Bus?

- A. Storage Queues
- B. Event Hubs
- C. Event Grid
- D. SignalR

Answer: B. Event Hubs is the partitioned event-streaming service for high-throughput telemetry; Service Bus is for enterprise messaging.

Q36. You need serverless functions to react when blobs are created in a storage account. Which service delivers these discrete events push-style?

- A. Service Bus topic
- B. Event Grid
- C. Event Hubs
- D. Storage Queue polling

Answer: B. Event Grid is built for reactive discrete-event delivery with push subscriptions to handlers.

Q37. Your queue storage requirement exceeds 80 GB and you need only simple queueing semantics at lowest cost. Choose:

- A. Service Bus Premium
- B. Azure Storage Queues
- C. Event Hubs Dedicated
- D. Service Bus Standard

Answer: B. Storage Queues support massive storage cheaply; Service Bus queues cap out far lower but add rich features.

Q38. What is the effect of setting a message's Scheduled_enqueue_time?

- A. Message expires at that time
- B. Message becomes visible for delivery at that time
- C. Consumer must poll at that time
- D. Message is deferred

Answer: B. Scheduled messages are enqueued immediately but not available until the scheduled time.

Q39. A consumer needs to postpone processing a specific message and retrieve it later by sequence number. Which feature?

- A. Abandon
- B. Deferral
- C. Dead-lettering
- D. Peek

Answer: B. Defer sets the message aside; it is only retrievable explicitly by sequence number.

Q40. Which filter type on a subscription offers the best throughput for exact-match routing on a property?

- A. SQL filter
- B. Correlation filter
- C. Boolean filter
- D. No filter

Answer: B. Correlation filters are cheap exact-match comparisons; SQL filters are more flexible but costlier.

Q41. What does auto-forwarding enable?

- A. Automatic retries
- B. Chaining a queue/subscription to another entity for fan-in or decoupling
- C. Sending to multiple namespaces
- D. Priority queues

Answer: B. Auto-forwarding moves messages automatically to a destination entity within the namespace.

Q42. During a regional outage with classic Geo-DR configured, what does initiating failover do?

- A. Moves all messages to secondary
- B. Points the alias at the secondary namespace so clients reconnect; unprocessed primary messages remain behind
- C. Restores from backup
- D. Nothing without Microsoft support

Answer: B. Failover re-points the alias; because only metadata replicates, in-flight primary messages are not transferred.

Q43. Which pattern coordinates a long-running distributed transaction across services using messages and compensating actions?

- A. CQRS
- B. Saga
- C. Bulkhead
- D. Retry

Answer: B. Sagas replace ACID distributed transactions with a sequence of local transactions plus compensations.

Q44. A message lock expires while a consumer is still processing. What happens?

- A. Processing continues safely
- B. The message becomes available again and may be processed twice — handlers must be idempotent
- C. Message dead-letters immediately
- D. Namespace throttles

Answer: B. Lock expiry causes redelivery; at-least-once semantics demand idempotent consumers or lock renewal.

Q45. Which protocol does Service Bus primarily use for messaging?

- A. MQTT
- B. AMQP 1.0
- C. gRPC
- D. WebSub

Answer: B. Service Bus speaks AMQP 1.0 (with HTTPS/REST fallback for some operations).

RBAC & Governance

Q46. What three elements make up an Azure role assignment?

- A. User, password, scope
- B. Security principal, role definition, scope
- C. Policy, initiative, assignment
- D. Tenant, subscription, resource

Answer: B. Assignment = who (principal) + what (role definition) + where (scope).

Q47. At which scopes can Azure RBAC roles be assigned?

- A. Subscription only
- B. Management group, subscription, resource group, resource
- C. Tenant and resource only
- D. Resource group only

Answer: B. All four levels; assignments inherit down the hierarchy.

Q48. Which property in a role definition grants data-plane operations like reading blob contents?

- A. Actions
- B. DataActions
- C. NotActions
- D. AssignableScopes

Answer: B. DataActions cover data-plane operations; Actions cover control-plane (ARM) operations.

Q49. What does NotActions do in a role definition?

- A. Explicitly denies operations
- B. Subtracts operations from Actions — it is not a deny
- C. Blocks inheritance
- D. Disables the role

Answer: B. NotActions narrows the granted set; another role can still grant those operations. Deny requires deny assignments.

Q50. Which built-in role can do everything except manage role assignments and shares?

- A. Owner
- B. Contributor
- C. User Access Administrator
- D. Reader

Answer: B. Contributor manages resources but cannot grant access; Owner adds role assignment rights.

Q51. A user has Reader on a subscription via group A and Contributor on one RG via group B. Effective rights on that RG?

- A. Reader only
- B. Contributor (union of assignments)
- C. No access — conflict
- D. Owner

Answer: B. RBAC is additive: effective permissions are the union of all assignments at or above the scope.

Q52. By default, what Azure resource permissions does an Entra ID Global Administrator have?

- A. Owner on all subscriptions
- B. None — but they can elevate access to User Access Administrator at root scope
- C. Contributor everywhere
- D. Reader everywhere

Answer: B. Directory roles and Azure RBAC are separate; elevation at '/' is an audited break-glass capability.

Q53. Your security team requires admins to activate the Owner role only when needed, with approval and MFA. What do you implement?

- A. Custom role
- B. Privileged Identity Management (PIM) eligible assignments
- C. Conditional Access
- D. Resource locks

Answer: B. PIM provides just-in-time activation, approvals, MFA, time limits, and audit for privileged roles.

Q54. Which managed identity type survives deletion of the compute resource and can be shared across many resources?

- A. System-assigned
- B. User-assigned
- C. Service principal with secret
- D. Device identity

Answer: B. User-assigned MIs are standalone resources with independent lifecycle — ideal for fleets and pre-provisioned permissions.

Q55. An app on Azure VMs must read Key Vault secrets with no credentials in config. Recommend:

- A. Store the vault key in appsettings
- B. Managed identity + Key Vault Secrets User role
- C. Shared access signature
- D. Certificate on disk

Answer: B. Managed identity plus the data-plane Key Vault Secrets User role (RBAC model) eliminates stored secrets.

Q56. You need to allow access only to blobs tagged project=alpha for certain principals. Which capability adds this filter to a role assignment?

- A. Custom role NotActions
- B. ABAC conditions on the role assignment
- C. Deny assignment
- D. Azure Policy

Answer: B. ABAC conditions refine storage data-action assignments based on attributes like blob index tags.

Q57. Who can create deny assignments directly?

- A. Owners
- B. Global admins
- C. No one — Azure creates them via deployment stacks/managed apps
- D. User Access Administrators

Answer: C. Deny assignments are system-created (e.g., deployment stacks deny settings, managed apps), not directly user-creatable.

Q58. Enforcing that all resources deploy only to approved regions across 60 subscriptions is best done with...

- A. Custom RBAC roles
- B. Azure Policy assigned at a management group
- C. Resource locks
- D. Tags

Answer: B. Policy governs resource configuration; assigning at MG scope covers all child subscriptions. RBAC can't express this.

Q59. What distinguishes Azure Policy from RBAC?

- A. Nothing, they overlap
- B. RBAC controls who can act; Policy controls what configurations are allowed regardless of who
- C. Policy replaces RBAC
- D. RBAC applies only to users

Answer: B. They are complementary: identity-based access vs. resource compliance guardrails.

Q60. Which policy effect automatically deploys a Log Analytics agent when a VM is created without one?

- A. Deny
- B. Audit
- C. DeployIfNotExists
- D. Append

Answer: C. DeployIfNotExists remediates by deploying required configuration after resource creation.

Q61. To prevent accidental deletion of a production resource group while allowing normal edits, apply:

- A. ReadOnly lock
- B. CanNotDelete lock
- C. Deny assignment
- D. Policy deny

Answer: B. CanNotDelete allows modify but blocks delete; ReadOnly would block updates too.

Q62. What is the recommended structure for applying baseline policies and RBAC across many subscriptions in an enterprise?

- A. Per-subscription manual config
- B. Management group hierarchy (landing zones) with inherited policy and RBAC
- C. Single giant subscription
- D. Tags

Answer: B. Landing zone MG hierarchies (platform/workloads/sandbox) inherit governance downward — the CAF approach.

Q63. A custom role must be assignable only within two subscriptions. Which property enforces this?

- A. Actions
- B. AssignableScopes
- C. DataActions
- D. Condition

Answer: B. AssignableScopes limits where the custom role definition may be assigned.

Q64. Best practice for granting the finance team access to their resources:

- A. Assign roles to each user individually
- B. Assign roles to an Entra ID group at the narrowest sufficient scope
- C. Give Owner at subscription
- D. Share one service account

Answer: B. Group-based, least-privilege, narrow-scope assignments simplify lifecycle management.

Q65. Which role should a CI/CD pipeline get to push images to Azure Container Registry, following least privilege?

- A. Owner on the registry
- B. AcrPush
- C. Contributor on the RG
- D. AcrPull

Answer: B. AcrPush grants push+pull only — least privilege for build pipelines.

Q66. What limits standing access risk from permanent role assignments during audits?

- A. More custom roles

- B. PIM access reviews and eligible (not active) assignments
- C. Longer key rotation
- D. Resource locks

Answer: B. Access reviews + JIT eligibility ensure privileges exist only when justified.

Q67. An external GitHub Actions workflow must deploy to Azure without a stored secret. Recommend:

- A. Service principal with client secret in GitHub
- B. Workload identity federation (OIDC) to a service principal/managed identity
- C. Personal access token
- D. SAS token

Answer: B. Federated credentials let GitHub's OIDC token be exchanged for Entra tokens — no long-lived secrets.

Networking

Q68. NSG rules are evaluated by...

- A. Name alphabetically
- B. Priority number, lowest first, first match wins
- C. Creation order
- D. Random

Answer: B. Rules process from priority 100 upward; the first match applies and processing stops.

Q69. You need web-tier VMs referenced in NSG rules by workload rather than IP address. Use:

- A. Service tags
- B. Application security groups (ASGs)
- C. UDRs
- D. Private endpoints

Answer: B. ASGs group NICs logically so rules say 'web-asg → db-asg' instead of IP ranges.

Q70. Is VNet peering transitive (spoke A ↔ hub ↔ spoke B means A reaches B)?

- A. Yes, automatically
- B. No — spoke-to-spoke needs routing through an NVA/firewall or gateway in the hub (or direct peering)
- C. Only cross-region
- D. Only with NSGs

Answer: B. Peering is non-transitive; hub-spoke designs route spoke↔spoke via the hub firewall with UDRs.

Q71. To force all outbound internet traffic from spokes through Azure Firewall, configure:

- A. NSG outbound rules
- B. UDR with 0.0.0.0/0 next hop = firewall private IP on spoke subnets
- C. Peering settings
- D. DDoS protection

Answer: B. A default route via UDR steers egress to the firewall for centralized inspection.

Q72. Which hybrid option provides private connectivity that does NOT traverse the public internet with an SLA?

- A. S2S VPN
- B. ExpressRoute
- C. P2S VPN

- D. Peering

Answer: B. ExpressRoute is a dedicated private circuit through a connectivity provider; VPNs ride the internet.

Q73. Recommended resilience design for a critical ExpressRoute connection:

- A. Second ER circuit only in same peering location
- B. Site-to-site VPN as failover path (or ER in a second peering location)
- C. Nothing needed, ER has SLA
- D. Public endpoints

Answer: B. ER + VPN failover (or geo-redundant circuits) is the standard guidance for hybrid resilience.

Q74. What distinguishes a private endpoint from a service endpoint?

- A. Nothing
- B. Private endpoint gives a private IP mapped to a specific resource instance, reachable from on-prem, allowing public access to be disabled; service endpoint keeps the service public
- C. Service endpoints are newer
- D. Private endpoints are free

Answer: B. Instance-level private IP + on-prem reachability + public-access lockdown = Private Link. Service endpoints only optimize the path from a subnet.

Q75. After creating a private endpoint for a storage account, name resolution must return the private IP. What do you configure?

- A. Hosts file on every VM
- B. Private DNS zone privatelink.blob.core.windows.net linked to the VNets
- C. Public DNS CNAME
- D. NSG rule

Answer: B. Private DNS zones for the privatelink subdomain make the FQDN resolve privately inside linked VNets.

Q76. On-premises clients must resolve Azure private endpoint names. Which managed service handles hybrid DNS forwarding?

- A. Azure DNS public zones
- B. Azure DNS Private Resolver (inbound endpoint)
- C. Traffic Manager
- D. App Gateway

Answer: B. On-prem DNS conditionally forwards to the Private Resolver inbound endpoint, which resolves against private zones.

Q77. A global HTTP application needs edge acceleration, WAF, and automatic failover between two regions. Best fit:

- A. Azure Load Balancer
- B. Azure Front Door
- C. Traffic Manager
- D. App Gateway alone

Answer: B. Front Door is the global L7 anycast entry with caching, WAF, and origin failover.

Q78. Global routing for a non-HTTP (TCP port 1433-style) service across regions is best served by:

- A. Front Door Standard
- B. Traffic Manager (DNS-based)
- C. App Gateway
- D. Bastion

Answer: B. Traffic Manager is protocol-agnostic DNS routing; Front Door focuses on HTTP(S).

Q79. Which service provides regional L7 load balancing with TLS termination, URL path routing, and WAF?

- A. Azure LB
- B. Application Gateway
- C. NAT Gateway
- D. Front Door

Answer: B. App Gateway is the regional L7 (HTTP) load balancer with WAF_v2 SKU.

Q80. Which load balancer would you choose for high-performance L4 TCP traffic to a VM scale set within one region?

- A. Application Gateway
- B. Azure Load Balancer (Standard)
- C. Front Door
- D. Traffic Manager

Answer: B. Standard LB is the regional L4 option with HA ports and zone redundancy.

Q81. Admins need RDP/SSH to VMs that must not have public IPs. Recommend:

- A. Public LB NAT rules
- B. Azure Bastion
- C. P2S VPN for every admin laptop
- D. Corporate proxy

Answer: B. Bastion provides browser-based RDP/SSH over TLS directly in the portal without exposing VM public IPs.

Q82. Many VMs share outbound internet access and you must avoid SNAT port exhaustion with predictable egress IPs. Use:

- A. Default outbound access
- B. NAT Gateway on the subnet
- C. Instance-level public IPs
- D. Service endpoints

Answer: B. NAT Gateway provides scalable, predictable SNAT for outbound flows.

Q83. Which Azure Firewall SKU adds TLS inspection, IDPS, and URL filtering?

- A. Basic
- B. Standard
- C. Premium
- D. WAF

Answer: C. Firewall Premium adds TLS inspection, IDPS, URL filtering, and web categories.

Q84. Connecting 40 branches, multiple regions, ER circuits, and P2S users with managed any-to-any routing suggests:

- A. Manual hub-spoke with many UDRs
- B. Azure Virtual WAN
- C. One big flat VNet
- D. Peering mesh

Answer: B. Virtual WAN is the managed global transit backbone for large-scale branch/hybrid connectivity.

Q85. What do NSG service tags like 'Storage' or 'AzureCloud' represent?

- A. Custom IP lists you maintain

- B. Microsoft-managed groups of IP prefixes for Azure services
- C. DNS names
- D. Subnets

Answer: B. Service tags are auto-updated prefix groups, avoiding manual IP maintenance in rules.

Q86. Where should DDoS Network Protection be enabled to protect public IPs of a production app?

- A. On each VM
- B. On the virtual network (plan covering its public IPs)
- C. On the subscription only
- D. It is always on with no option

Answer: B. DDoS Network Protection is enabled on VNets, covering associated public IP resources (IP Protection covers single IPs). Basic infrastructure protection is always on, but the enhanced plan is opt-in.

Q87. Two VNets in different regions and subscriptions must communicate privately. Simplest option:

- A. S2S VPN between them
- B. Global VNet peering
- C. ExpressRoute for both
- D. Public IPs with NSGs

Answer: B. Global peering connects VNets across regions/subscriptions over the Microsoft backbone.

Q88. A workload requires that traffic to Azure SQL never traverse public endpoints and be reachable from on-premises via ExpressRoute. Solution:

- A. Service endpoint on the subnet
- B. Private endpoint for the SQL server + private DNS + ER private peering
- C. SQL firewall IP rules
- D. Traffic Manager

Answer: B. Only private endpoints give a routable private IP reachable over ER/VPN with public access disabled.

Q89. Which combination fronts a multi-region AKS/API platform with global entry and regional WAF ingress?

- A. Traffic Manager → Bastion
- B. Front Door → regional Application Gateways → backends
- C. LB → Front Door
- D. NAT GW → App GW

Answer: B. Global Front Door for edge/failover, regional App GW (WAF) as ingress is the canonical layered pattern.

Container Apps

Q90. Azure Container Apps is built on which open-source technologies (abstracted from the user)?

- A. Docker Swarm + Consul
- B. Kubernetes + KEDA + Dapr + Envoy
- C. Nomad + Istio
- D. Mesos + Linkerd

Answer: B. ACA runs on AKS infrastructure with KEDA scaling, Dapr building blocks, and Envoy ingress — without exposing the K8s API.

Q91. A team wants microservices with autoscaling and service discovery but has no Kubernetes expertise and doesn't need the K8s API. Recommend:

- A. AKS
- B. Container Apps
- C. ACI
- D. VMSS

Answer: B. ACA gives K8s-grade capabilities as a serverless platform; AKS is for teams needing full API control.

Q92. When must you choose AKS over Container Apps?

- A. Any HTTP API
- B. You need custom CRDs/operators, node-level control, or a specific service mesh
- C. Scale to zero
- D. Dapr usage

Answer: B. ACA hides the K8s API — custom operators, DaemonSets, node customization require AKS.

Q93. What is a Container Apps environment?

- A. A resource group
- B. An isolation boundary sharing VNet and Log Analytics for a set of apps
- C. A Docker network
- D. A subscription

Answer: B. Apps in one environment share networking (optionally your VNet) and logging; internal environments keep ingress private.

Q94. How do you run a gradual rollout sending 10% of traffic to a new app version in ACA?

- A. Deploy a second app and DNS round-robin
- B. Multiple revision mode with traffic splitting 90/10 across revisions
- C. Feature flags only
- D. Not possible

Answer: B. Revisions are immutable versions; traffic weights across revisions enable canary/blue-green.

Q95. Which technology powers ACA scale rules, including scaling on Service Bus queue length?

- A. HPA only
- B. KEDA
- C. Cluster autoscaler
- D. Cron

Answer: B. KEDA scalers drive event-driven scaling (Service Bus, Event Hubs, Kafka, cron, HTTP concurrency, CPU/memory).

Q96. What is the cost benefit of setting minReplicas to 0?

- A. None
- B. Scale to zero — no compute cost when idle (with cold-start tradeoff)
- C. Free egress
- D. Reserved pricing

Answer: B. Consumption workloads can scale to zero; latency-sensitive apps often keep $\text{min} \geq 1$ to avoid cold starts.

Q97. A nightly batch container must run to completion on a schedule inside the same environment as your apps. Use:

- A. Always-on app with sleep loop

- B. ACA Jobs (scheduled trigger)
- C. Logic Apps only
- D. AKS CronJob only

Answer: B. ACA Jobs support manual, scheduled (cron), and event-driven run-to-completion workloads.

Q98. Which Dapr building block would ACA microservices use for pub/sub over Service Bus without SDK-level coupling?

- A. Dapr state store
- B. Dapr pub/sub component backed by Service Bus
- C. Dapr secrets
- D. Envoy filter

Answer: B. Dapr pub/sub abstracts the broker; the component config points at Service Bus, keeping app code portable.

Q99. Best-practice authentication for ACA pulling images from Azure Container Registry:

- A. Admin username/password
- B. Managed identity with AcrPull role
- C. Anonymous pull
- D. Docker Hub mirror

Answer: B. Managed identity + AcrPull avoids registry credentials entirely.

Q100. How should an internal-only microservices platform on ACA be exposed to other private workloads?

- A. External environment with IP restrictions
- B. Internal environment (VNet) with internal ingress and private DNS
- C. Public ingress + WAF
- D. Port forwarding

Answer: B. Internal environments place ingress on a private IP in your VNet; consumers resolve via private DNS.

Q101. Which ACA feature lets one environment host both serverless apps and dedicated-compute (or GPU) apps?

- A. Revisions
- B. Workload profiles
- C. Dapr
- D. Jobs

Answer: B. Workload profiles combine Consumption and Dedicated (D/E-series, GPU) profiles in a single environment.

Q102. To control ACA egress traffic through Azure Firewall you should...

- A. Use consumption-only environment defaults
- B. Deploy a workload profiles environment in your VNet and apply UDRs to the firewall
- C. Enable NAT on the app
- D. Impossible

Answer: B. Workload-profile VNet environments support UDRs for firewall-controlled egress.

Q103. A single short-lived container needs to run for an isolated build task, no orchestration. Cheapest, simplest service:

- A. AKS
- B. Azure Container Instances
- C. App Service

- D. Service Fabric

Answer: B. ACI runs single containers on demand per-second billing — no orchestration overhead.

Q104. Where do ACA apps get secrets at runtime following best practice?

- A. Baked into the image
- B. Key Vault references resolved via managed identity (or app secrets)
- C. Plain env vars in source control
- D. ConfigMaps

Answer: B. Key Vault-backed secret references with managed identity keep secrets centralized and rotated.

Q105. What does ACA session affinity do?

- A. Encrypts sessions
- B. Routes a client's requests to the same replica (sticky sessions)
- C. Splits traffic by revision
- D. Caches responses

Answer: B. Sticky sessions help stateful web workloads; stateless design remains preferable.

Q106. An HTTP API on ACA should scale based on concurrent requests per replica. Which scale rule?

- A. CPU rule only
- B. HTTP scale rule (concurrent requests)
- C. Cron rule
- D. Storage queue rule

Answer: B. The HTTP scaler targets concurrent requests per replica — the natural rule for synchronous APIs.

Q107. Which choice is best for a team needing web app hosting with deployment slots, easy certs/domains, and no containers required?

- A. Container Apps
- B. App Service
- C. AKS
- D. ACI

Answer: B. App Service is the PaaS web host with slots and integrated domain/cert management.

Auth & Identity

Q108. Which flow should a single-page application use to sign in users?

- A. Implicit flow
- B. Authorization code flow with PKCE
- C. Client credentials
- D. ROPC

Answer: B. Auth code + PKCE replaced implicit flow for SPAs — tokens via secure code exchange.

Q109. A background daemon with no user context needs to call an API. Which flow?

- A. Device code
- B. Client credentials
- C. Auth code + PKCE
- D. OBO

Answer: B. Client credentials issues app-only tokens; prefer certificates, managed identity, or federation over secrets.

Q110. API A receives a user's token and must call downstream API B as that user. Which flow?

- A. Client credentials
- B. On-Behalf-Of (OBO)
- C. Device code
- D. Refresh token replay

Answer: B. OBO exchanges the incoming user assertion for a token to the downstream API preserving user identity.

Q111. Which token should a client present to a protected API?

- A. ID token
- B. Access token
- C. Refresh token
- D. SAML assertion always

Answer: B. Access tokens are for APIs; ID tokens are for the client's own authentication of the user.

Q112. An API validating a JWT must check which claims at minimum?

- A. Only expiry
- B. Signature, issuer (iss), audience (aud), expiry (exp), then scopes/roles
- C. Only audience
- D. Header only

Answer: B. Full validation: signature against JWKS, iss, aud, exp/nbf, then authorize via scp/roles claims.

Q113. In Entra tokens, what indicates delegated permissions vs application permissions?

- A. aud vs iss
- B. scp claim (delegated) vs roles claim (application)
- C. tid vs oid
- D. email claim

Answer: B. scp carries delegated scopes with user context; roles carries app roles for app-only tokens (or user role assignments).

Q114. Application (app-only) permissions on Microsoft Graph require what before use?

- A. User consent
- B. Admin consent
- C. Nothing
- D. MFA

Answer: B. Application permissions always need admin consent; delegated ones may too depending on scope.

Q115. A smart-TV app needs user sign-in but has no browser/keyboard. Which flow?

- A. Client credentials
- B. Device code flow
- C. Implicit
- D. ROPC

Answer: B. Device code shows a code the user enters on another device to complete sign-in.

Q116. Why is ROPC (resource owner password credentials) discouraged?

- A. Too slow
- B. App handles raw passwords, breaks MFA/Conditional Access, no federation
- C. Requires certificates
- D. Deprecated TLS

Answer: B. ROPC is legacy-only: incompatible with modern security controls and passwordless.

Q117. What is the difference between an app registration and an enterprise application (service principal)?

- A. Identical
- B. Registration is the global app definition; the service principal is its instance/identity in a tenant
- C. SP is only for SAML
- D. Registration is per-user

Answer: B. One registration (in the home tenant) can have SPs in many tenants (multi-tenant apps).

Q118. Your customer-facing product needs social logins, custom-branded sign-up, and consumer identity at scale. Use:

- A. Entra ID employee tenant
- B. Microsoft Entra External ID (CIAM, successor to B2C)
- C. B2B guest invites
- D. ADFS

Answer: B. External ID is Microsoft's CIAM offering for consumer sign-up/sign-in and social identities.

Q119. Partner-company employees need access to your internal app using their own credentials. Use:

- A. Create local accounts
- B. Entra B2B collaboration (guest users)
- C. External ID consumer tenant
- D. Shared password

Answer: B. B2B guests federate partners into your tenant, governed by access reviews and entitlement management.

Q120. Which Entra feature enforces 'require MFA when sign-in risk is high, block legacy auth'?

- A. PIM
- B. Conditional Access (with Identity Protection signals)
- C. RBAC
- D. Access reviews

Answer: B. Conditional Access evaluates signals and applies controls; Identity Protection supplies risk scores.

Q121. Best way for APIM to enforce that only tokens containing the Orders.Read scope reach the orders API:

- A. Check subscription key
- B. validate-jwt with required-claims on scp/scope
- C. IP filtering
- D. Backend checks only

Answer: B. Gateway-level claim checks reject unauthorized calls early; backend still re-validates (defense in depth).

Q122. Should the backend API skip token validation because APIM already validated the JWT?

- A. Yes, saves latency
- B. No — zero trust: every hop validates; network position is not identity
- C. Only in dev
- D. Only for GETs

Answer: B. Assume breach: backends must independently validate tokens even behind a gateway.

Q123. Which mechanism gives an AKS pod or external workload an Entra identity without storing any secret?

- A. Client secret in a K8s secret
- B. Workload identity federation
- C. SAS tokens
- D. Cert on node disk

Answer: B. Federation trusts the platform's OIDC token and exchanges it for Entra tokens — no stored credentials.

Q124. For time-boxed, anonymous-ish download access to a single blob for an external user, which is appropriate?

- A. Entra ID account creation
- B. User-delegation SAS with short expiry
- C. Public container
- D. Owner role

Answer: B. SAS remains right for constrained delegation; user-delegation SAS is signed with Entra credentials — safer than account keys.

Q125. High-assurance service-to-service calls between APIM and a payments backend should add which transport control?

- A. Plain HTTP inside VNet
- B. Mutual TLS (client certificates) plus network isolation
- C. Longer JWTs
- D. Basic auth

Answer: B. mTLS authenticates both ends at transport layer, layered with token auth and private networking.

Q126. What are the three guiding principles of Zero Trust?

- A. Encrypt, replicate, monitor
- B. Verify explicitly; use least-privilege access; assume breach
- C. Patch, scan, alert
- D. Isolate, air-gap, backup

Answer: B. These principles drive per-request verification, JIT/JEA access, and segmentation/telemetry.

Q127. A multi-tenant SaaS API must accept tokens from many customer tenants. What must its token validation handle?

- A. Single fixed issuer
- B. Issuer validation per allowed tenant (tid) with common metadata endpoint
- C. No issuer check
- D. Only audience

Answer: B. Multi-tenant apps validate the issuer pattern/tenant allowlist since iss varies per tenant.

Q128. Where should refresh tokens for a confidential web app be kept?

- A. Browser localStorage

- B. Server-side, protected (e.g., encrypted session/store) — never exposed to JS
- C. URL parameters
- D. Cookies without flags

Answer: B. Confidential clients keep tokens server-side; SPAs rely on MSAL cache with short-lived tokens instead.

Q129. Which Key Vault capability supports automatic rotation stories and eliminates app restarts for secret updates?

- A. Hard-coded secrets
- B. Versioned secrets with rotation policies and apps referencing latest via MI
- C. One static secret
- D. Export to env file at build

Answer: B. Rotation policies plus managed-identity retrieval of current versions enables zero-touch rotation.

Design Patterns

Q130. Incrementally replacing a legacy monolith by routing selected paths to new microservices through a gateway is which pattern?

- A. Sidecar
- B. Strangler Fig
- C. Bulkhead
- D. Geode

Answer: B. Strangler Fig migrates functionality slice by slice behind a routing facade until the monolith is retired.

Q131. Mobile and web clients need differently shaped APIs. Which pattern avoids one bloated general-purpose API?

- A. CQRS
- B. Backends for Frontends (BFF)
- C. Retry
- D. Cache-aside

Answer: B. BFF gives each client type its own tailored backend, often implemented at/behind the gateway.

Q132. A downstream dependency is failing; you want to stop hammering it and fail fast until it recovers. Pattern?

- A. Retry with backoff
- B. Circuit breaker
- C. Throttling
- D. Saga

Answer: B. Circuit breakers trip after failures, short-circuit calls, and probe for recovery — preventing cascade.

Q133. Which pattern pairs naturally WITH circuit breaker for transient faults?

- A. Retry with exponential backoff and jitter
- B. Strangler
- C. Deployment stamps
- D. BFF

Answer: A. Retry handles transient blips; the breaker stops retries when failure is persistent.

Q134. Isolating resource pools so one tenant's overload cannot exhaust capacity for others is which pattern?

- A. Bulkhead
- B. Claim-check
- C. Ambassador
- D. Gateway aggregation

Answer: A. Bulkheads partition connections/compute per consumer/workload, containing failures.

Q135. Separating the read model from the write model, often with different stores, is...

- A. Saga
- B. CQRS
- C. Event sourcing only
- D. Sharding

Answer: B. CQRS splits command and query paths; frequently combined with event sourcing.

Q136. Storing state changes as an append-only sequence of events replayed to build state is...

- A. Cache-aside
- B. Event sourcing
- C. Materialized view only
- D. Snapshotting

Answer: B. Event sourcing keeps the event log as the source of truth; projections build read models.

Q137. An app loads data into Redis on cache miss and invalidates on update. Which pattern?

- A. Write-through
- B. Cache-aside
- C. CDN
- D. Static content hosting

Answer: B. Cache-aside (lazy loading): check cache → miss → load from store → populate cache.

Q138. Combining multiple backend calls into one client round trip at the gateway is...

- A. Gateway routing
- B. Gateway aggregation
- C. Gateway offloading
- D. Ambassador

Answer: B. Aggregation reduces chattiness for clients; routing directs requests; offloading moves cross-cutting concerns.

Q139. Terminating TLS and handling auth at APIM instead of in each microservice is which gateway pattern?

- A. Aggregation
- B. Gateway offloading
- C. Strangler
- D. BFF

Answer: B. Offloading centralizes cross-cutting functionality (TLS, auth, throttling) at the gateway.

Q140. Deploying identical isolated stamps of the full stack per tenant group or region is...

- A. Geode
- B. Deployment Stamps
- C. Bulkhead
- D. Blue-green

Answer: B. Stamps are repeatable scale/isolation units; geodes are active-active regional nodes behind global routing.

Q141. A translation layer preventing a new system from being polluted by legacy models/protocols is...

- A. Adapter sidecar
- B. Anti-Corruption Layer
- C. Facade cache
- D. Bridge table

Answer: B. ACL isolates bounded contexts, translating between new and legacy semantics.

Q142. Dapr running alongside each container app providing mTLS, retries, and pub/sub is an example of which pattern?

- A. Ambassador only
- B. Sidecar
- C. Gateway
- D. Broker

Answer: B. Dapr is a sidecar: co-located helper extending the app without code coupling.

Q143. Health probes exposed by services so load balancers and orchestrators can gate traffic is which pattern?

- A. Watchdog
- B. Health Endpoint Monitoring
- C. Heartbeat table
- D. Ping-pong

Answer: B. Health endpoints let infrastructure verify functional status, not just process liveness.

AZ-305 Exam

Q144. Two services in series have 99.95% and 99.9% SLAs. Approximate composite availability?

- A. 99.95%
- B. 99.85%
- C. 99.99%
- D. 99.5%

Answer: B. Serial dependencies multiply: $0.9995 \times 0.999 \approx 0.99850 \rightarrow 99.85\%$.

Q145. Deploying an instance in two independent regions each with 99.9% availability behind global failover yields roughly...

- A. 99.8%
- B. 99.9999% theoretical ($1 - 0.001^2$)
- C. 99.9%
- D. 98%

Answer: B. Parallel redundancy: $1 - (1 - 0.999)^2 = 99.9999\%$ theoretical ceiling (real-world lower due to failover mechanics).

Q146. RPO of 15 minutes means...

- A. Recovery must finish in 15 minutes
- B. Up to 15 minutes of data loss is acceptable
- C. Backups run every 15 minutes always
- D. SLA is 15 min/month downtime

Answer: B. RPO = max tolerable data loss window; RTO = max tolerable downtime.

Q147. Which Well-Architected pillar does 'use IaC, safe deployment practices, and observability' map to?

- A. Reliability
- B. Operational Excellence
- C. Cost Optimization
- D. Security

Answer: B. Operational Excellence covers DevOps practices, automation, and monitoring of operations.

Q148. A solution must survive a datacenter failure within a region at minimal redesign. Recommend:

- A. Region pair failover only
- B. Availability zone-redundant deployment
- C. Single VM with premium SSD
- D. Bigger VM SKU

Answer: B. Zone redundancy protects against datacenter-level failure within the region.

Q149. Which AZ-305 domain carries the highest weight (April 2026 blueprint)?

- A. Identity/governance/monitoring
- B. Data storage
- C. Business continuity
- D. Infrastructure (30-35%)

Answer: D. Design infrastructure solutions is 30-35%; identity/governance and data storage are 25-30% each; BC is 10-15%.

Q150. What is the prerequisite certification for the Azure Solutions Architect Expert credential?

- A. AZ-900
- B. AZ-104 (Azure Administrator Associate)
- C. AZ-204
- D. SC-100

Answer: B. Passing AZ-305 plus holding AZ-104 earns the Expert certification.

Q151. Company policy: minimize cost; workload is a dev/test API with sporadic traffic tolerating cold starts. Best APIM tier?

- A. Premium
- B. Consumption
- C. Standard
- D. Developer with SLA

Answer: B. Consumption is serverless pay-per-call with scale-to-zero — cheapest for sporadic traffic; exam logic = cheapest meeting requirements.

Q152. 'You need to recommend a messaging solution that guarantees FIFO and supports transactions.' Recommend:

- A. Event Grid
- B. Service Bus with sessions
- C. Storage Queues
- D. Event Hubs

Answer: B. Only Service Bus offers sessions (FIFO) plus transactions among these options.

Q153. A case study requires ‘no secrets in code or config anywhere.’ Which combination satisfies it end to end?

- A. Key Vault with access keys
- B. Managed identities + Entra RBAC data roles + Key Vault (MI-accessed) for third-party secrets
- C. Encrypted appsettings
- D. SAS tokens rotated monthly

Answer: B. MI removes Azure-to-Azure secrets; Key Vault (accessed by MI) holds unavoidable external secrets.

Q154. Requirement: global users, HTTP app, sub-50ms static content latency, WAF, cheapest single service. Recommend:

- A. Traffic Manager + App GW per region
- B. Azure Front Door (Standard/Premium)
- C. Azure CDN classic + LB
- D. App Gateway

Answer: B. Front Door combines global anycast entry, caching/CDN, WAF, and origin failover in one service.

Q155. Monitoring design: you need KQL queries across app traces and platform logs with alerting. Recommend:

- A. Storage account diagnostics only
- B. Azure Monitor with Log Analytics workspace + App Insights, alert rules & action groups
- C. Event Hubs export only
- D. Metrics explorer only

Answer: B. Log Analytics is the KQL store; App Insights feeds APM data; alert rules + action groups automate response.

Q156. Which tool estimates monthly cost of a proposed architecture before deployment?

- A. Cost analysis blade
- B. Azure Pricing Calculator
- C. Advisor
- D. TCO calculator only

Answer: B. Pricing Calculator models planned architectures; TCO compares on-prem migration economics; Cost analysis tracks actual spend.

Q157. Steady-state 24/7 production VMs for 3 years, minimizing cost, no flexibility needed. Recommend:

- A. Pay-as-you-go
- B. Reserved instances (3-year) or savings plan
- C. Spot VMs
- D. Dev/test pricing

Answer: B. Reservations give the deepest discount for predictable long-running compute; Spot is for interruptible work.

Q158. Which service provides recommendations across cost, security, reliability, performance, and operational excellence?

- A. Defender for Cloud only

- B. Azure Advisor
- C. Policy compliance
- D. Service Health

Answer: B. Advisor surfaces WAF-aligned recommendations; Defender focuses on security posture.

Q159. A regulated workload needs customer-managed keys, double encryption, and immutable audit logs. Which services combine?

- A. Default SSE only
- B. Key Vault (or Managed HSM) CMK + infrastructure encryption + immutable storage/Log Analytics export
- C. BitLocker on VMs
- D. Client-side encryption only

Answer: B. CMK via Key Vault/HSM, infrastructure (double) encryption, and immutability policies satisfy common compliance asks.

Q160. Design for 'API platform must remain available during regional outage; writes must continue.' Data tier suggestion:

- A. Single-region SQL with backups
- B. Cosmos DB multi-region writes (or SQL failover groups for RW failover)
- C. Read replicas only
- D. Local Redis

Answer: B. Multi-region write (Cosmos) or auto-failover groups keep write availability across regional failure.

Q161. Hybrid identity: users must sign in with the same password on-prem and cloud, minimal infrastructure, cloud auth. Recommend:

- A. ADFS federation
- B. Entra password hash synchronization
- C. Pass-through authentication with agents
- D. Local accounts

Answer: B. PHS is the simplest, most resilient hybrid auth; PTA/ADFS add infrastructure and failure modes.

Q162. Which assessment should you take (free, official) to gauge AZ-305 readiness?

- A. Any braindump site
- B. Microsoft Learn Practice Assessment for AZ-305
- C. Random YouTube quiz
- D. AZ-104 practice

Answer: B. Microsoft's free official practice assessment mirrors exam style; braindumps violate NDA and teach wrong answers.

Q163. In case-study questions, the most reliable answering strategy is:

- A. Pick the most powerful service
- B. Map explicit requirements/constraints to the cheapest compliant service; ignore unstated needs
- C. Always choose Premium tiers
- D. Choose newest service

Answer: B. AZ-305 rewards requirement-driven minimal designs — never gold-plate beyond stated constraints.

Q164. 'Recommend a solution to enforce org-wide resource naming and tagging.' Best answer:

- A. RBAC custom roles
- B. Azure Policy (deny/append/modify) at management group scope
- C. Manual reviews
- D. Locks

Answer: B. Policy is the governance engine for configuration standards; RBAC cannot express naming rules.

Q165. Which migration framework phase sequence does Microsoft's Cloud Adoption Framework define?

- A. Build-run-retire
- B. Strategy → Plan → Ready → Adopt (Migrate/Innovate) → Govern + Manage
- C. Lift and shift only
- D. Assess-deploy

Answer: B. CAF: Strategy, Plan, Ready (landing zones), Adopt, with Govern and Manage as ongoing disciplines.

Q166. An architecture review finds a single APIM Developer-tier instance serving production traffic. Biggest issue?

- A. Too expensive
- B. No SLA on Developer tier — production requires Basic+ (and redundancy for critical apps)
- C. Wrong region
- D. Too many policies

Answer: B. Developer tier carries no SLA and is licensed for evaluation/dev/test — a classic exam trap.

Q167. To protect an APIM-fronted API platform against OWASP Top 10 web attacks, add:

- A. More rate limits
- B. WAF (Front Door Premium or App Gateway WAF_v2) in front of APIM
- C. NSGs only
- D. Larger tier

Answer: B. APIM is not a WAF; layer a WAF at the edge for OWASP rule-set protection.

Q168. Requirement: 'audit all privileged role activations and require approval for Owner.' Combine:

- A. RBAC + locks
- B. PIM (eligible assignments, approval, audit) + Entra audit logs
- C. Policy + tags
- D. Defender plans

Answer: B. PIM handles JIT activation with approvals; audit logs record activations for compliance.

Q169. Which two artifacts should anchor every AZ-305 design recommendation?

- A. Marketing sheets
- B. Well-Architected Framework pillars + Cloud Adoption Framework guidance
- C. Pricing page + roadmap
- D. Blog posts

Answer: B. Microsoft explicitly frames the architect role around WAF and CAF alignment.

Data Storage

Q170. A lift-and-shift SQL Server app needs SQL Agent jobs and cross-database queries with minimal changes. Recommend:

- A. Azure SQL Database

- B. SQL Managed Instance
- C. SQL on VM only
- D. Cosmos DB

Answer: B. Managed Instance offers near-100% SQL Server compatibility (Agent, cross-DB) as a managed service; SQL Database lacks these instance features.

Q171. Which purchasing model allows Azure Hybrid Benefit and independent compute/storage scaling for Azure SQL?

- A. DTU
- B. vCore
- C. Serverless only
- D. Elastic

Answer: B. vCore decouples compute/storage and supports Hybrid Benefit and reserved capacity; DTU bundles everything.

Q172. A dev database is used a few hours daily; cost must be minimal with auto-pause when idle. Recommend:

- A. Business Critical
- B. SQL serverless (General Purpose) with auto-pause
- C. Hyperscale
- D. Provisioned Premium DTU

Answer: B. Serverless auto-scales and auto-pauses — you pay only storage while paused.

Q173. A SaaS vendor hosts 400 tenant databases with unpredictable, spiky usage. Most cost-efficient:

- A. 400 provisioned single DBs
- B. Elastic pool sharing resources
- C. One giant database
- D. Cosmos DB serverless

Answer: B. Elastic pools share eDTU/vCores across databases with complementary usage peaks — the classic multi-tenant answer.

Q174. Requirement: 60 TB OLTP database, rapid scale-out of read replicas, near-instant restores. Which tier?

- A. General Purpose
- B. Business Critical
- C. Hyperscale
- D. Basic

Answer: C. Hyperscale supports 100+ TB, snapshot-based fast backup/restore, and quickly added named replicas.

Q175. Which SQL tier includes a free-to-use readable secondary replica via Application-Intent=ReadOnly?

- A. General Purpose
- B. Business Critical
- C. Serverless
- D. DTU Standard

Answer: B. Business Critical's Always On replicas include a built-in readable secondary for read offloading.

Q176. In Cosmos DB, what is the default and most commonly recommended consistency level?

- A. Strong

- B. Eventual
- C. Session
- D. Bounded staleness

Answer: C. Session guarantees read-your-own-writes within a client session — the latency/consistency sweet spot.

Q177. Order Cosmos DB consistency levels strongest to weakest:

- A. Strong, Session, Bounded, Prefix, Eventual
- B. Strong, Bounded Staleness, Session, Consistent Prefix, Eventual
- C. Bounded, Strong, Session, Eventual, Prefix
- D. Session, Strong, Bounded, Eventual, Prefix

Answer: B. Strong → Bounded Staleness → Session → Consistent Prefix → Eventual.

Q178. A Cosmos container shows one physical partition throttling while others idle. Root cause?

- A. Too many regions
- B. Hot partition from low-cardinality/skewed partition key
- C. Wrong consistency level
- D. TTL misconfigured

Answer: B. Skewed partition keys concentrate traffic; choose high-cardinality keys aligned to query patterns.

Q179. Globally distributed app requires writes accepted in every region with RPO≈0. Configure:

- A. Single write region + read replicas
- B. Cosmos DB multi-region writes
- C. GRS storage
- D. SQL geo-replication

Answer: B. Multi-region writes give active-active write availability with conflict resolution policies.

Q180. Cosmos DB workload is spiky and low volume overall; cheapest throughput model?

- A. Provisioned standard RU/s
- B. Autoscale RU/s
- C. Serverless (per-request RUs)
- D. Dedicated gateway

Answer: C. Serverless bills per consumed RU with no minimum — best for low/spiky traffic; autoscale suits variable but sustained loads.

Q181. Which storage redundancy survives both a zone outage AND a regional disaster?

- A. LRS
- B. ZRS
- C. GRS
- D. GZRS

Answer: D. GZRS = zone-redundant locally + geo-replicated to the pair region; RA-GZRS adds secondary reads.

Q182. Compliance requires blob data be immutable (WORM) for 7 years. Use:

- A. Soft delete
- B. Time-based retention immutability policy
- C. Cool tier
- D. Snapshots

Answer: B. Immutable storage time-based retention enforces write-once-read-many; legal holds cover open-ended cases.

Q183. Logs accessed rarely after 30 days must auto-move to cheaper tiers and delete after 2 years. Use:

- A. Manual scripts
- B. Lifecycle management policies
- C. AzCopy cron
- D. Archive on write

Answer: B. Lifecycle rules tier (hot→cool→cold→archive) and expire blobs by age/last-access automatically.

Q184. Retrieval from which blob tier can take hours and requires rehydration?

- A. Hot
- B. Cool
- C. Cold
- D. Archive

Answer: D. Archive is offline storage; rehydrate (standard/priority) before reading.

Q185. Big-data analytics needs a hierarchical filesystem with POSIX ACLs on blob storage. Enable:

- A. Versioning
- B. ADLS Gen2 hierarchical namespace
- C. NFS Azure Files
- D. Premium block blobs only

Answer: B. HNS turns Blob storage into ADLS Gen2 — directories, POSIX ACLs, analytics-optimized.

Q186. Branch offices need local fast access to a central file share with cloud as source of truth. Recommend:

- A. Blob + AzCopy
- B. Azure File Sync with cloud tiering on local servers
- C. OneDrive
- D. Data Box

Answer: B. File Sync caches hot files on local Windows servers, tiering cold data to Azure Files.

Q187. Which is the safest SAS type for delegating temporary blob access?

- A. Account SAS
- B. Service SAS with account key
- C. User-delegation SAS (Entra-signed)
- D. Never use SAS

Answer: C. User-delegation SAS is signed with Entra credentials, avoiding account keys and enabling identity-based revocation.

Q188. Session state and hot lookups need sub-millisecond reads shared across app instances. Recommend:

- A. SQL table
- B. Azure Cache for Redis
- C. Cosmos strong consistency
- D. Blob storage

Answer: B. Redis is the in-memory cache for session/hot-path data — classic cache-aside companion.

Q189. Which feature lets downstream services react to every insert/update in a Cosmos container?

- A. TTL
- B. Change feed
- C. Unique keys
- D. Synapse Link

Answer: B. Change feed streams document changes — powering event sourcing, materialized views, and integrations.

Business Continuity

Q190. RTO vs RPO: a 4-hour RTO and 15-minute RPO means...

- A. Restore within 15 min, lose up to 4 h data
- B. Restore within 4 h, lose at most 15 min of data
- C. Both are 4 h
- D. Backups every 4 h

Answer: B. RTO bounds downtime; RPO bounds data loss.

Q191. Which statement about backup vs replication is correct?

- A. Replication replaces backup
- B. Backup protects against corruption/deletion (point-in-time); replication protects against infrastructure loss — you need the right mix of both
- C. Backup is always cheaper DR
- D. GRS storage is a backup

Answer: B. Replication faithfully copies corruption/ransomware too; only point-in-time backups recover from it.

Q192. Which vault type backs up Azure VMs and SQL-in-VM workloads?

- A. Backup vault
- B. Recovery Services vault
- C. Key vault
- D. Storage account

Answer: B. Recovery Services vaults handle VMs, SQL/SAP in VMs, Files, and MARS/MABS; Backup vaults cover blobs/disks/PostgreSQL.

Q193. Ransomware protection for backups includes which features?

- A. Faster backups
- B. Soft delete, immutable vaults, multi-user authorization
- C. GRS only
- D. Larger retention only

Answer: B. Soft delete retains deleted backups, immutability blocks tampering, MUA requires a second identity for destructive ops.

Q194. A 3-tier app needs regional DR with RTO of minutes, RPO of seconds, and ordered multi-tier failover with scripts. Recommend:

- A. Azure Backup with cross-region restore
- B. Azure Site Recovery with recovery plans
- C. GRS storage
- D. Manual redeploy from IaC

Answer: B. ASR continuously replicates and orchestrates failover via recovery plans; backup restores are far slower.

Q195. How do you validate ASR readiness without disrupting production?

- A. Real failover at midnight
- B. Test failover into an isolated network
- C. Restore one VM
- D. Trust the dashboard

Answer: B. Test failovers spin up replicas in isolated networks — run them regularly.

Q196. Apps must reconnect automatically to the secondary SQL region after failover without connection string changes. Configure:

- A. Active geo-replication only
- B. Auto-failover groups (listener endpoints)
- C. Geo-restore
- D. Export/import

Answer: B. Failover groups provide stable read-write/read-only listener endpoints that follow the failover.

Q197. Azure SQL point-in-time restore covers up to 35 days. Requirement: restore capability for 7 years. Add:

- A. Bigger PITR window
- B. Long-term retention (LTR) policies
- C. Geo-replication
- D. Snapshots

Answer: B. LTR stores weekly/monthly/yearly full backups up to 10 years.

Q198. Cheapest availability improvement that survives datacenter failure within a region:

- A. Second region active-active
- B. Availability zone redundancy
- C. Bigger VMs
- D. ASR

Answer: B. Zone redundancy is the low-cost first step; multi-region only where RTO/RPO justify it.

Q199. Which routing setup implements active-passive regional failover for an HTTP app?

- A. Front Door priority routing to primary with health-probe failover to secondary
- B. Round-robin DNS
- C. Weighted 50/50
- D. NSG rules

Answer: A. Priority routing sends traffic to the healthy highest-priority origin, failing over automatically.

Q200. Which service injects controlled faults (zone down, CPU pressure) to validate resilience?

- A. Load Testing
- B. Chaos Studio
- C. Monitor
- D. Advisor

Answer: B. Chaos Studio runs fault-injection experiments against your workloads.

Q201. Storage account with RA-GRS: what does customer-initiated account failover do?

- A. Nothing without support ticket
- B. Promotes secondary region to primary (potential data loss up to Last Sync Time)
- C. Zero data loss failover
- D. Converts to ZRS

Answer: B. Failover promotes the async secondary — data after Last Sync Time is lost; plan around it.

Monitoring

Q202. Which store do KQL queries in Azure Monitor run against?

- A. Metrics database
- B. Log Analytics workspace
- C. Storage account
- D. Event Hubs

Answer: B. Logs land in Log Analytics; KQL queries, log alerts, and workbooks operate there.

Q203. You need distributed tracing across APIM → Container Apps → Service Bus consumers. Use:

- A. Activity log
- B. Application Insights (workspace-based) with correlation
- C. Metrics explorer
- D. Service Health

Answer: B. App Insights provides end-to-end transaction tracing and the Application Map across services.

Q204. Alert when p95 latency KQL query over 15 min exceeds a threshold. Which alert type?

- A. Metric alert
- B. Log (scheduled query) alert
- C. Activity log alert
- D. Smart detection only

Answer: B. Custom KQL aggregations need log alerts; simple platform metrics use faster metric alerts.

Q205. What executes notifications and automation when any alert fires?

- A. Runbooks only
- B. Action groups
- C. Event Grid mandatory
- D. Logic Apps only

Answer: B. Action groups define receivers (email/SMS/webhook/Function/Logic App/ITSM) reused across alert rules.

Q206. Platform logs must go to a SIEM outside Azure. Which diagnostic-settings destination fits?

- A. Log Analytics
- B. Storage account
- C. Event Hubs
- D. Workbooks

Answer: C. Event Hubs streams logs to external SIEM/analytics; Storage is archive; Log Analytics is Azure-native analysis.

Q207. Teams must query only their own resources' logs in a shared central workspace. Use:

- A. Separate workspaces per team always
- B. Resource-context RBAC / table-level RBAC
- C. Export copies per team

- D. No solution

Answer: B. Resource-context access scopes queries to resources the user can read — keeping central design with per-team isolation.

Q208. Which tool proactively probes your public endpoint from multiple regions and alerts on failures?

- A. Resource Health
- B. App Insights availability tests
- C. Network Watcher
- D. Advisor

Answer: B. Availability (web) tests run synthetic checks globally with alerting.

Q209. Azure region incident notifications (platform-side) come from...

- A. Resource Health
- B. Service Health alerts
- C. App Insights
- D. Defender

Answer: B. Service Health covers Azure incidents/maintenance; Resource Health reports your specific resource's state.

Q210. A SOC needs cross-source security correlation, hunting, and automated response playbooks. Recommend:

- A. Defender for Cloud alone
- B. Microsoft Sentinel (SIEM/SOAR) on Log Analytics
- C. Advisor
- D. Activity log

Answer: B. Sentinel is the cloud-native SIEM/SOAR; Defender for Cloud feeds it posture and workload alerts.

Q211. Which practices control Log Analytics cost?

- A. Log everything forever
- B. Retention tuning, archive/long-term tiers, daily caps, Basic table plans, App Insights sampling
- C. Delete the workspace monthly
- D. Use metrics only

Answer: B. Cost levers: retention, table plans, caps, and sampling — an exam-favorite operational question.

Q212. NSG/VNet flow logs and packet capture for network troubleshooting come from...

- A. Azure Firewall
- B. Network Watcher
- C. Front Door
- D. Traffic Manager

Answer: B. Network Watcher provides flow logs, connection monitor, IP flow verify, and packet capture.

Compute

Q213. Two VMs in an availability set vs spread across availability zones — which SLA applies to zones?

- A. 99.9%
- B. 99.95%
- C. 99.99%
- D. 100%

Answer: C. Zone-spread multi-VM deployments get 99.99%; availability sets 99.95%; single premium-SSD VM 99.9%.

Q214. Interruptible batch rendering should minimize compute cost. Recommend:

- A. Reserved instances
- B. Spot VMs (evictable, up to ~90% off)
- C. Premium VMs
- D. Dedicated hosts

Answer: B. Spot pricing suits fault-tolerant, interruptible workloads; never for production-critical services.

Q215. Which Functions plan eliminates cold starts with pre-warmed instances and adds VNet integration?

- A. Consumption
- B. Premium (Elastic)
- C. Free
- D. Shared

Answer: B. Premium keeps warm instances, supports VNet and unlimited duration — for latency-sensitive event apps.

Q216. Fan-out/fan-in of 500 parallel activity functions with aggregated results requires...

- A. Plain HTTP functions
- B. Durable Functions orchestration
- C. Logic Apps only
- D. Webjobs

Answer: B. Durable Functions patterns: chaining, fan-out/fan-in, async HTTP, monitors, human interaction.

Q217. An AKS API server must not be reachable from the internet. Configure:

- A. NSG on nodes
- B. Private cluster (private API endpoint)
- C. Internal load balancer only
- D. Deny policy

Answer: B. Private AKS clusters place the control-plane endpoint on a private IP (Private Link).

Q218. Pods on AKS need Entra tokens to access Key Vault without stored secrets. Use:

- A. Kubeconfig secrets
- B. Entra Workload Identity (federated service accounts)
- C. Node managed identity for all pods
- D. Client secrets in env vars

Answer: B. Workload identity federates K8s service accounts to Entra — per-pod identity, no secrets.

Q219. Massively parallel Monte-Carlo simulation across thousands of cores, job-based. Best service:

- A. App Service
- B. Azure Batch
- C. Container Apps
- D. Logic Apps

Answer: B. Batch schedules large-scale parallel/HPC jobs on managed pools (Spot-capable).

Q220. Which VM cost lever provides discounts while retaining flexibility to change VM families and regions?

- A. Reserved instances
- B. Savings plan for compute
- C. Spot
- D. Hybrid Benefit

Answer: B. Savings plans commit to \$/hr spend flexibly; reservations lock family/region for deeper discounts.

Migration & Integration

Q221. Which of the 5 Rs describes moving VMs as-is to Azure IaaS?

- A. Refactor
- B. Rehost
- C. Rearchitect
- D. Replace

Answer: B. Rehost = lift-and-shift; fastest migration, least cloud optimization.

Q222. Before migrating 300 on-prem VMs you need discovery, dependency mapping, right-sizing, and cost estimates. Use:

- A. ASR directly
- B. Azure Migrate assessment
- C. Advisor
- D. Data Box

Answer: B. Azure Migrate handles discovery/assessment including dependency visualization and sizing.

Q223. Minimal-downtime online migration of SQL Server databases to SQL MI uses...

- A. Backup/restore over weekend
- B. Database Migration Service (online)
- C. BACPAC export
- D. Copy database wizard

Answer: B. DMS online mode continuously syncs until cutover, minimizing downtime.

Q224. You must move 800 TB to Azure from a site with poor bandwidth. Recommend:

- A. AzCopy over VPN
- B. Azure Data Box offline transfer
- C. ExpressRoute first
- D. File Sync

Answer: B. Data Box ships encrypted appliances — offline bulk transfer when network is impractical.

Q225. Orchestrating nightly ETL copying data from on-prem SQL + SaaS APIs into a lake calls for...

- A. Functions timers
- B. Azure Data Factory with self-hosted integration runtime
- C. Event Grid
- D. Stream Analytics

Answer: B. ADF orchestrates hybrid ETL; the self-hosted IR reaches on-prem sources.

Q226. Real-time SQL queries over Event Hubs streams with windowing (tumbling/hopping) suggests...

- A. Databricks batch
- B. Azure Stream Analytics
- C. ADF mapping flows
- D. Synapse dedicated pool

Answer: B. Stream Analytics is SQL-on-streams with native windowing functions.

Q227. Sub-second interactive analytics over billions of telemetry rows (KQL) is the sweet spot of...

- A. Azure SQL
- B. Azure Data Explorer (ADX)
- C. Blob storage
- D. Redis

Answer: B. ADX/Data Explorer powers telemetry/time-series analytics — the same engine behind Log Analytics.

Q228. Business analysts need a unified SaaS platform: lakehouse, warehouse, real-time, and Power BI in one. Recommend:

- A. Separate services stitched manually
- B. Microsoft Fabric (OneLake)
- C. HDInsight
- D. ACI

Answer: B. Fabric unifies analytics workloads over OneLake — Microsoft's strategic analytics direction.

Q229. Low-code integration across Salesforce, SAP, and on-prem systems with connectors and approval steps:

- A. Durable Functions
- B. Logic Apps
- C. Service Bus alone
- D. Event Grid

Answer: B. Logic Apps' 1400+ connectors and designer suit workflow/B2B integration; code-first orchestration → Durable Functions.

Auth & Identity

Q230. Key Vault must guarantee deleted secrets/vaults can't be permanently removed during retention (CMK requirement). Enable:

- A. RBAC
- B. Soft delete + purge protection
- C. Firewall
- D. Premium SKU

Answer: B. Purge protection blocks permanent deletion until retention expires — mandatory for CMK encryption scenarios.

Q231. FIPS 140-2 Level 3 dedicated HSM control of keys requires...

- A. Key Vault Standard
- B. Key Vault Premium shared HSM
- C. Managed HSM
- D. Storage SSE

Answer: C. Managed HSM is the dedicated, single-tenant FIPS 140-2 L3 offering; Premium uses shared HSMs.

Q232. Which Key Vault authorization model does Microsoft recommend for new vaults?

- A. Access policies
- B. Azure RBAC (data-plane roles)
- C. SAS
- D. Shared keys

Answer: B. RBAC unifies management, supports PIM/conditions, and replaces legacy per-vault access policies.

RBAC & Governance

Q233. Governance must extend Azure Policy and RBAC to VMs running in AWS and on-premises. Use:

- A. Azure Migrate
- B. Azure Arc-enabled servers
- C. Site Recovery
- D. VPN

Answer: B. Arc projects external machines into ARM so Policy, RBAC, tags, Monitor, and Defender apply.

Design Patterns

Q234. Which IaC feature blocks out-of-band portal edits to deployed resources and cleans up removed ones?

- A. Template specs
- B. Deployment stacks with deny settings
- C. Blueprints (current)
- D. Locks only

Answer: B. Deployment stacks manage deployments as units with deny settings — the Blueprints successor.

Q235. Bicep vs Terraform — which statement is accurate?

- A. Terraform is Azure-only
- B. Bicep is Azure-native (day-0 ARM support, no state file); Terraform is multicloud with state management
- C. Bicep requires state files
- D. Terraform can't deploy Azure

Answer: B. Bicep compiles to ARM with no separate state; Terraform brings multicloud + ecosystem with state files.

Q236. CI/CD should deploy to Azure with zero stored cloud credentials. Configure:

- A. Service principal secret in repo settings
- B. OIDC workload identity federation from GitHub Actions/Azure DevOps
- C. PAT tokens
- D. Publish profiles

Answer: B. Federated OIDC exchanges pipeline tokens for Entra tokens — no long-lived secrets to leak or rotate.

Q237. Which managed service load-tests your API with JMeter/Locust scripts in CI/CD?

- A. Chaos Studio
- B. Azure Load Testing
- C. App Insights
- D. Batch

Answer: B. Azure Load Testing runs managed large-scale load tests and integrates with pipelines for regression gates.

AZ-305 Exam

Q238. Requirement: ‘store app feature flags and non-secret settings centrally with Key Vault references for secrets.’ Recommend:

- A. Key Vault for everything
- B. Azure App Configuration + Key Vault references
- C. Cosmos DB
- D. Env vars

Answer: B. App Configuration handles settings/feature flags; secrets stay in Key Vault via references.

Q239. A workload needs 99.99% availability. Single-region design maxes at 99.95% due to a component. Next step per WAF?

- A. Accept it silently
- B. Add zone/regional redundancy for that component or adjust the requirement with stakeholders
- C. Buy support plan
- D. Increase VM size

Answer: B. Architects surface the gap: redesign the weakest serial component or renegotiate the target.

Q240. Case study: ‘minimize administrative effort’ for SQL HA across regions with app transparency. Best answer:

- A. Log shipping to VM
- B. Auto-failover groups
- C. Manual geo-restore runbook
- D. Storage replication

Answer: B. ‘Minimize admin effort’ + transparency = failover groups with listener endpoints — no app changes.

Q241. Which combination satisfies ‘analyze last month’s platform logs cheaply, rarely queried, kept 5 years’?

- A. Keep all in interactive retention
- B. Log Analytics with archive/long-term retention tier (restore/search jobs when needed)
- C. Event Hubs forever
- D. Premium SSD storage

Answer: B. Long-term/archive tiers store cheaply; search/restore jobs rehydrate on demand.

Q242. An architect must pick between App Service, Container Apps, and AKS for 12 microservices; team has no K8s skills; needs Dapr, scale-to-zero. Recommend:

- A. AKS
- B. Container Apps
- C. App Service
- D. ACI

Answer: B. ACA delivers microservice capabilities (Dapr, KEDA, revisions) without K8s operational burden.

Q243. Which service estimates savings from moving on-prem servers to Azure including power/real-estate costs?

- A. Pricing Calculator
- B. TCO Calculator
- C. Cost analysis
- D. Advisor

Answer: B. TCO Calculator compares total on-prem cost vs Azure; Pricing Calculator prices planned Azure architectures.

API Management

Q244. Your rate-limit-by-key policy uses counter-key="constant". What behavior results?

- A. Per-IP limits
- B. All callers share one global bucket and starve each other
- C. Policy is ignored
- D. 429 for every call

Answer: B. The counter key defines the bucket; a constant means every caller shares it. Use expressions like `context.Request.IpAddress` or a JWT claim.

Q245. After importing an API, calls return 401 despite a valid Entra token. The validate-jwt audience is set to the CLIENT app ID. Fix?

- A. Extend token lifetime
- B. Set audience to the API's identifier (aud claim of the token)
- C. Remove issuer check
- D. Use subscription keys instead

Answer: B. aud must match the API/resource the token was issued FOR, not the client that requested it — a classic misconfiguration.

Q246. A daemon requests scope=api://orders/Orders.Read with client_credentials and fails. Why?

- A. Secret expired
- B. App-only flows must request the /.default scope; named scopes are for delegated flows
- C. Wrong tenant
- D. Roles missing in manifest

Answer: B. Client credentials uses `api://orders/.default`, which rolls up granted application permissions (app roles).

Service Bus

Q247. Processing takes 90 s but the queue lock duration is 30 s with no renewal. Result?

- A. Message dead-letters immediately
- B. Message unlocks mid-processing and is redelivered — duplicates
- C. Nothing, locks auto-extend
- D. Namespace throttles

Answer: B. Expired locks redeliver messages; renew locks or extend duration, and make handlers idempotent.

Q248. A queue has TTL 5 minutes and dead-lettering-on-expiration disabled. Backlogged messages...

- A. Wait forever
- B. Silently disappear when TTL expires
- C. Move to DLQ
- D. Throttle senders

Answer: B. Without dead-letter-on-expiration, expired messages are simply dropped — enable it so nothing vanishes unnoticed.

Q249. Sending to a session-enabled queue fails with ‘session identifier missing’. Why?

- A. Queue is full
- B. Session-enabled entities require every message to carry a SessionId
- C. Wrong tier
- D. AMQP disabled

Answer: B. Session queues enforce SessionId on send; receivers accept sessions rather than plain messages.

RBAC & Governance

Q250. A VM’s managed identity gets 403 reading blobs despite a ‘Reader’ role on the storage account. Root cause?

- A. Token expired
- B. Reader is control-plane only; data access needs Storage Blob Data Reader
- C. Wrong region
- D. Firewall

Answer: B. Control-plane Reader sees the resource’s configuration, not blob contents — data-plane roles (DataActions) are required.

Q251. Your IMDS token request uses resource=https://management.azure.com/ but you call blob storage. Result?

- A. Works fine
- B. 401/403 — audience mismatch; request https://storage.azure.com/
- C. Slower calls
- D. Token refresh loop

Answer: B. Tokens are audience-bound; the storage data plane rejects management-plane tokens.

Networking

Q252. After creating a private endpoint, a spoke VM still resolves the storage FQDN to a public IP. Least likely cause?

- A. Private DNS zone not linked to the VNet
- B. Missing DNS zone group (no A record)
- C. Custom DNS not forwarding to 168.63.129.16
- D. NSG blocking port 443

Answer: D. NSGs affect connectivity, not name resolution. The other three are the classic private-endpoint DNS failure modes.

Q253. Spoke A cannot reach spoke B though both peer with the hub. Peering is healthy. What is missing?

- A. Global peering
- B. UDRs sending spoke traffic to the hub firewall (peering is non-transitive)
- C. More address space
- D. Service endpoints

Answer: B. Transitivity must be built: route spoke↔spoke traffic through an NVA/firewall in the hub with UDRs (or use direct peering/VWAN).

Q254. In the hub-spoke SaaS design, why does Azure Policy deny public network access on PaaS resources when private endpoints already exist?

- A. Policy speeds up traffic
- B. Defense in depth: architecture provides the path, policy prevents drift back to public exposure
- C. Cheaper egress
- D. DNS requires it

Answer: B. Governance enforces the architecture over time — someone will eventually click ‘enable public access’ without it.

Container Apps

Q255. Canary traffic weights are configured but 100% of traffic hits the old revision. Likely cause?

- A. KEDA missing
- B. App is in single-revision mode — weights need multiple-revision mode
- C. Ingress internal
- D. Min replicas 0

Answer: B. Single revision mode always routes to the latest revision; traffic splitting requires multiple revision mode.

Q256. A KEDA Service Bus scale rule never scales the worker. Which is NOT a plausible cause?

- A. Wrong queue name in metadata
- B. Invalid connection secret for the scaler
- C. messageCount threshold never reached
- D. Dapr not enabled

Answer: D. Dapr is unrelated to KEDA scale rules; the others are the usual suspects.

Auth & Identity

Q257. jwt.ms shows your token has scp=User.Read but your API checks roles. What happened?

- A. Token corrupt
- B. You obtained a delegated token; the API expects app roles (application permissions)
- C. Wrong signing key
- D. Expired token

Answer: B. scp = delegated scopes (user flows); roles = app roles. Match the flow to what the API authorizes on.

Q258. Why are two cloud-only break-glass accounts excluded from Conditional Access?

- A. To save licenses
- B. So a CA misconfiguration or MFA outage cannot lock every admin out of the tenant
- C. They’re service accounts
- D. Compliance requires it

Answer: B. Emergency access accounts guarantee tenant access during CA/MFA failures — monitored and alerted on every use.

Q259. FinSecure chose password hash sync over ADFS mainly because...

- A. It’s more configurable
- B. Cloud auth keeps working during on-prem outages with near-zero infrastructure
- C. It avoids licenses
- D. ADFS is deprecated

Answer: B. PHS is the simplest, most resilient hybrid identity option — federation adds servers, certificates, and failure modes.

Data Storage

Q260. ShopSphere keeps EU order data on EU SQL servers and pairs failover groups within the same geography because...

- A. Latency only
- B. GDPR data residency must survive DR — failing over to a US region would violate it
- C. Cost
- D. SQL requires it

Answer: B. Residency by architecture: DR topology must respect the same data-boundary constraints as production.

Q261. Which Cosmos consistency did the e-commerce scenario choose for carts, and why?

- A. Strong — money involved
- B. Session — shoppers see their own cart instantly without global write latency
- C. Eventual — fastest
- D. Bounded — compliance

Answer: B. Session gives read-your-own-writes for each user at low latency — carts don't need global strong consistency.

Business Continuity

Q262. ForgeWorks runs Region B as a scaled-down warm standby instead of cold IaC rebuild because...

- A. IaC is unreliable
- B. 15-minute RTO can't wait for full provisioning; warm standby scales out in minutes
- C. Cheaper than cold
- D. Zones require it

Answer: B. Standby temperature follows RTO: cold = hours (cheapest), warm = minutes, hot/active = seconds (priciest).

Q263. Why does ransomware defeat replication-only DR strategies?

- A. It disables ASR
- B. Replication faithfully copies the encrypted/corrupted data; only point-in-time immutable back-ups recover
- C. It attacks backups first
- D. It doesn't

Answer: B. Replication has no memory; backup provides history. Immutability + MUA protect that history from tampering.

Q264. Azure SQL PITR restores go to...

- A. The same database in place
- B. A new database — the app must cut over or you swap names
- C. A file share
- D. The paired region only

Answer: B. PITR creates a new database on the same server; plan the rename/cutover step in your runbook and measure its time.

Q265. Soft delete was enabled the day AFTER a blob was deleted. Can it be recovered?

- A. Yes, always
- B. No — protection features only cover events after they're enabled
- C. Only with GRS
- D. Only via support

Answer: B. Enable soft delete/versioning/immutability BEFORE you need them; retention must also exceed detection time.

Monitoring

Q266. An SRE must trace one failed order across APIM → ACA API → Service Bus → worker. Fastest tool?

- A. Activity log
- B. App Insights end-to-end transaction view (operation ID correlation)
- C. Metrics explorer
- D. NSG flow logs

Answer: B. Distributed tracing correlates every hop under one operation ID — the Application Map/transaction view shows the failing link.

Design Patterns

Q267. In the GitOps AKS design, rollback of a bad deployment is performed by...

- A. kubectl apply of old YAML from a laptop
- B. Reverting the Git commit — Flux reconciles the cluster to the previous state
- C. Restoring etcd
- D. Restarting nodes

Answer: B. Git is the source of truth; the revert is audited, repeatable, and automatic.

Q268. Why does the deployment stack's denyWriteAndDelete matter in the IaC lab?

- A. Faster deploys
- B. Blocks out-of-band portal edits so prod can't drift from the template
- C. Saves cost
- D. Enables what-if

Answer: B. Deny settings enforce that changes flow through the pipeline — drift prevention, not just detection.

Q269. The retailer taints its Spot node pool so only batch workloads schedule there because...

- A. Spot is faster
- B. Spot VMs can be evicted anytime; latency-critical pods must not land on them
- C. Licensing
- D. GPU drivers

Answer: B. Taints/tolerations keep interruptible capacity for interruptible work — cost optimization without risking SLOs.

AZ-305 Exam

Q270. 'Minimize administrative effort' for hybrid sign-in with cloud resilience. Best fit:

- A. ADFS farm
- B. Password hash sync
- C. Pass-through auth + 3 agents
- D. Certificate federation

Answer: B. PHS = no extra servers and works during on-prem outages — the 'minimize effort' phrasing points straight to it.

Q271. A case study says spoke-per-customer with central egress inspection. Which two limits should the architect plan for first?

- A. VM quotas
- B. IP address plan (supernet) and firewall SNAT/peering scale limits
- C. Storage IOPS
- D. DNS TTLs

Answer: B. Hub-spoke at 50+ spokes hits addressing and hub-throughput/SNAT limits before anything else — plan the /16 on day one.

Q272. Which statement about test failovers is the exam-correct posture?

- A. Optional if SLAs are high
- B. Run regularly into isolated networks; an untested DR plan doesn't count as a plan
- C. Only after incidents
- D. Production failovers monthly

Answer: B. ASR test failovers validate RTO without touching production — auditors and the exam both expect scheduled drills.

Q273. An architect must justify APIM Premium over Standard v2 in a design review. Which requirement clinches it?

- A. JWT validation
- B. Multi-region active gateways with internal VNet injection
- C. Rate limiting
- D. OpenAPI import

Answer: B. Features common to all tiers never justify Premium; multi-region + VNet injection are Premium differentiators.

Q274. CASE STUDY — FinSecure bank requires: no standing admin access, MFA on activation, quarterly access certification, and full audit of privileged actions. Which THREE features do you combine?

- A. PIM eligible assignments with MFA on activation
- B. PIM access reviews
- C. Entra audit logs / Log Analytics export
- D. Resource locks
- E. Azure Policy deny effect

Answers: A, B, C. PIM eligibility+MFA kills standing access, access reviews give certification, audit logs provide the trail. Locks and Policy govern resources, not privileged identity.

Q275. CASE STUDY — ShopSphere must keep EU customer PII in the EU while serving a global app, with regional failover that stays EU-compliant. Which TWO design choices are correct?

- A. Azure SQL in West Europe with failover group to North Europe
- B. Cosmos DB with write regions in East US and West Europe for PII
- C. Store PII pseudonymized references only in non-EU regions
- D. Enable RA-GRS from West Europe to East US for the PII database

Answers: A, C. WEU→NEU keeps both replicas in the EU; non-EU regions hold only pseudonymized references. Replicating PII to US regions (options B/D) violates residency.

Networking

Q276. Which THREE are required for a working private endpoint to Azure SQL from on-premises via ExpressRoute?

- A. Private endpoint in a VNet subnet
- B. Private DNS zone privatelink.database.windows.net with correct records reachable from on-prem resolvers
- C. ER private peering routing to the VNet

- D. Service endpoint on the gateway subnet
- E. Public firewall allow rule for the on-prem NAT IP

Answers: A, B, C. PE + DNS resolution to the private IP + private-peering routing are the triad. Service endpoints don't work from on-prem; public firewall rules defeat the purpose.

Auth & Identity

Q277. An API validates incoming JWTs. Which FOUR checks are mandatory before trusting the token?

- A. Signature against the issuer's published keys (JWKS)
- B. Issuer (iss) claim
- C. Audience (aud) claim
- D. Expiry (exp/nbf)
- E. Token length > 500 characters

Answers: A, B, C, D. Signature, issuer, audience, and time validity are the core validation set; token length is meaningless.

Business Continuity

Q278. Which TWO protect backups against ransomware actors with stolen admin credentials?

- A. Immutable vault
- B. Multi-user authorization for destructive operations
- C. GRS redundancy
- D. Longer retention

Answers: A, B. Immutability blocks tampering; MUA requires a second identity to approve destructive changes. GRS/retention don't stop malicious deletion attempts.

Data Storage

Q279. Which TWO Cosmos DB choices best fit an unpredictable, low-average, spiky workload with occasional bursts?

- A. Serverless capacity mode
- B. Autoscale provisioned throughput
- C. Standard (manual) provisioned at peak
- D. Strong consistency to smooth load

Answers: A, B. Serverless (low/spiky) or autoscale (variable but sustained) both avoid paying peak 24/7. Manual-at-peak wastes money; consistency isn't a capacity lever.

API Management

Q280. Which THREE belong in an enterprise APIM security baseline for a public API?

- A. validate-jwt (or validate-azure-ad-token) at the gateway
- B. WAF (Front Door/App GW) in front of APIM
- C. Managed identity or mTLS from APIM to backends
- D. Rely on subscription keys as sole authentication
- E. Expose backends publicly for redundancy

Answers: A, B, C. Token validation + WAF edge + secured backend channel. Subscription keys alone aren't authentication; public backends bypass the gateway.

AZ-305 Exam

Q281. Arrange the CAF migration journey in the correct order:

- A. Strategy — define business justification
- B. Plan — inventory and skill planning
- C. Ready — build landing zones
- D. Adopt — migrate and innovate
- E. Govern & Manage — ongoing operations

Answer: the steps are listed above in the correct order (A→E). CAF: Strategy → Plan → Ready → Adopt, with Govern and Manage as continuing disciplines.

Business Continuity

Q282. Arrange the regional failover runbook for the ForgeWorks MES (SQL MI failover group + ASR + Traffic Manager):

- A. Confirm regional outage via Service Health and monitoring
- B. Fail over the SQL MI auto-failover group to the secondary
- C. Execute the ASR recovery plan to boot app VMs in order
- D. Verify app health probes in the secondary region
- E. Switch Traffic Manager/Front Door priority to the secondary

Answer: the steps are listed above in the correct order (A→E). Verify the disaster first, bring data tier up, then compute, validate, then shift traffic — shifting traffic before the stack is healthy causes a second outage.

Auth & Identity

Q283. Arrange the OAuth 2.0 authorization code + PKCE flow:

- A. Client generates code_verifier and code_challenge
- B. User authenticates at Entra ID authorize endpoint (challenge sent)
- C. Entra returns an authorization code to the redirect URI
- D. Client exchanges code + code_verifier at the token endpoint
- E. Client calls the API with the access token

Answer: the steps are listed above in the correct order (A→E). PKCE binds the code exchange to the verifier, preventing code interception attacks.

Networking

Q284. Arrange the packet path for an internet user reaching an internal-VNet APIM-fronted API:

- A. User resolves the app domain to Front Door anycast edge
- B. Front Door applies WAF rules and routes to origin via Private Link
- C. Internal APIM gateway validates JWT and applies policies
- D. APIM forwards to the Container Apps internal ingress
- E. Container app calls the data tier over private endpoints

Answer: the steps are listed above in the correct order (A→E). Edge → WAF → gateway policies → private backend → private data — every hop private after the edge.

Migration & Integration

Q285. Arrange an online (minimal-downtime) SQL Server → SQL MI migration:

- A. Run Data Migration Assistant compatibility assessment
- B. Provision SQL Managed Instance and networking

- C. Start DMS online migration (continuous sync)
- D. Validate data and performance on the target
- E. Cutover: stop writes, complete sync, repoint connection strings

Answer: the steps are listed above in the correct order (A→E). Assess → provision → sync → validate → brief cutover window. Offline mode would move the downtime to the full copy phase.

Design Patterns

Q286. Arrange the GitOps deployment flow for RetailRocket's AKS platform:

- A. Developer merges code PR; GitHub Actions builds and scans the image
- B. Pipeline pushes the signed image to ACR
- C. Pipeline opens a PR updating the image tag in the GitOps config repo
- D. Flux detects the merged config change and reconciles the cluster
- E. Canary weights shift traffic; burn-rate alerts gate promotion or rollback

Answer: the steps are listed above in the correct order (A→E). CI produces artifacts; CD is a Git merge that Flux pulls — pipelines never hold cluster credentials.

AZ-305 Exam

Q287. CASE STUDY — 'Minimize administrative effort' + 'apps must not change connection strings after regional DB failover' + 'automatic failover'. Recommend:

- A. Active geo-replication with manual DNS updates
- B. Auto-failover groups with grace period
- C. Geo-restore runbook
- D. Cross-region read replica promotion

Answer: B. Failover groups: automatic policy + listener endpoints = zero app changes, minimal admin effort — every qualifier satisfied.

Q288. CASE STUDY — A workload needs 99.99% compute SLA in ONE region at the LOWEST cost. Recommend:

- A. Two VMs in an availability set
- B. VMs spread across availability zones behind a Standard LB
- C. Single VM with premium SSD
- D. Multi-region active-active

Answer: B. Zone spread hits 99.99%; availability sets top out at 99.95%; multi-region exceeds requirement (and budget).

Monitoring

Q289. CASE STUDY — CloudNest tenants must view only their own spoke's logs; the SOC needs estate-wide correlation. Cheapest compliant design:

- A. 50 workspaces + 1 SOC workspace with cross-workspace queries
- B. One central workspace: resource-context RBAC for tenants, Sentinel for SOC
- C. Export each tenant's logs to their own storage account
- D. Per-tenant Sentinel instances

Answer: B. Central workspace + resource-context access satisfies isolation and correlation with one ingestion pipeline; 50 workspaces multiplies cost and ops.

Data Storage

Q290. CALCULATE — A read-heavy Cosmos container needs 8,000 RU/s at peak (4 h/day) and 800 RU/s otherwise. Which throughput model is cheapest?

- A. Standard provisioned 8,000 RU/s
- B. Autoscale with max 8,000 RU/s (scales 800–8,000)
- C. Serverless
- D. Two containers

Answer: B. Autoscale bills per-hour on actual scale between 10%–100% of max — ideal for daily peaks. Standard pays 8,000 all day; serverless gets expensive at sustained high RU consumption.

AZ-305 Exam

Q291. CALCULATE — App Gateway (99.95%) → 2 zone-redundant web VMs (99.99%) → SQL zone-redundant (99.99%). Approximate composite SLA:

- A. 99.99%
- B. 99.93%
- C. 99.95%
- D. 99.97%

Answer: B. Serial chain: $0.9995 \times 0.9999 \times 0.9999 \approx 0.9993 \rightarrow 99.93\%$. The weakest serial link dominates.

Container Apps

Q292. TROUBLESHOOT — An ACA app with minReplicas=0 shows 8-second first-request latency each morning. Cheapest fix preserving off-hours savings:

- A. minReplicas=1 always
- B. Scheduled scale rule (cron) warming one replica during business hours
- C. Move to AKS
- D. Bigger CPU allocation

Answer: B. A cron scale rule warms replicas only when users are active — cold starts eliminated during hours, scale-to-zero kept overnight.

Service Bus

Q293. TROUBLESHOOT — Orders are processed twice occasionally under load. Consumers take ~90 s per message; lock duration is 60 s. Best fix:

- A. Switch to receive-and-delete
- B. Renew locks (or raise lock duration) AND make handlers idempotent
- C. Add more consumers
- D. Disable dead-lettering

Answer: B. Lock expiry mid-processing causes redelivery. Fix the lock lifetime and keep idempotency — at-least-once delivery makes duplicates always possible.

19. Flashcards

273 flashcards — cover the right column and recall.

API Management

Prompt	Answer
Three components of Azure API Management APIM tier for multi-region deployment, VNet injection, availability zones APIM Consumption tier traits	API gateway (data plane), management plane, developer portal. Premium (Premium v2 for injection in v2 family). Serverless, per-call billing, scale to zero, no VNet injection — great for spiky/lightweight APIs.
Versions vs revisions	Versions = breaking changes, consumer-visible. Revisions = non-breaking iterations; test with ;rev=N, then make current.
Policy pipeline sections Policy scope order	inbound → backend → outbound, plus on-error. Global → Workspace → Product → API → Operation; composed with .
rate-limit vs quota	rate-limit = short-term burst throttling (429). quota = long-term call/bandwidth cap per renewal period.
validate-jwt policy checks	Signature via OpenID config JWKS, issuer, audience, expiry, required claims.
Products in APIM	Bundles of APIs with visibility, terms, subscription requirement/approval, and product-scope policies.
Subscription keys — proper role	Consumer identification/analytics; pair with OAuth 2.0 JWT validation for real authentication.
Named values	Key-value store for policies; can reference Key Vault secrets.
Internal VNet mode pattern	Front Door/App GW (WAF) public edge → internal APIM → private backends.
Self-hosted gateway	Containerized APIM gateway run on-prem/other clouds; Azure keeps the management plane — hybrid/multicloud APIs.
Workspaces (Premium)	Federated API management: teams own APIs/products with isolated workspace gateways in one instance.
authentication-managed-identity policy	Gateway acquires an Entra token with its managed identity to call the backend — no stored secrets.
Backends entity extras	Reusable backend config: credentials, mTLS certs, circuit breaker rules, load-balanced pools.
APIOps	GitOps for APIM: extractor + publisher pipelines promote configuration across environments via Git PRs.
Standard v2 networking capability	Outbound VNet integration to reach private backends (injection requires Premium/Premium v2).
AI gateway policies	llm-token-limit (token quotas), llm-semantic-cache (semantic response reuse) for LLM backends.
mock-response policy use	API-first development: return stub responses before the backend exists.

Prompt	Answer
Developer tier warning	Full features but NO SLA — never for production.

Service Bus

Prompt	Answer
Queue vs topic	Queue = point-to-point, competing consumers. Topic = pub/sub via independent filtered subscriptions.
Peek-lock semantics	At-least-once: message locked; complete/abandon/dead-letter/defer; lock expiry → redelivery → idempotent handlers required.
Receive-and-delete semantics	At-most-once: fastest, but messages lost on consumer crash.
Dead-letter queue triggers	MaxDeliveryCount exceeded, TTL expired, filter evaluation errors, explicit dead-lettering.
Sessions	FIFO ordering + stateful processing for messages sharing SessionId; required for strict ordering.
Duplicate detection	Dedupes by MessageId within a time window — idempotent enqueue for producer retries.
Premium tier exclusives	Dedicated MUs, private endpoints, zone redundancy, Geo-DR, CMK, 100 MB messages, JMS 2.0.
Geo-DR gotcha	Classic Geo-DR replicates METADATA ONLY — messages are not copied to the secondary.
Claim-check pattern	Store large payload in Blob Storage; send a small reference message instead.
Service Bus vs Event Grid vs Event Hubs	SB = enterprise messages/commands. Event Grid = discrete reactive events (push). Event Hubs = high-throughput event streaming.
When Storage Queues win	>80 GB queue storage, simple semantics, lowest cost — no sessions/dedup/ordering needed.
Data-plane RBAC roles	Azure Service Bus Data Owner / Data Sender / Data Receiver — scope to namespace or entity; use with managed identity.
Scheduled delivery vs deferral	Scheduled = visible at future time (producer-set). Deferral = consumer sets aside; retrieve by sequence number.
Correlation vs SQL filters	Correlation = fast exact-match on properties. SQL = flexible expressions, higher cost.
Queue-based load leveling	Queue buffers producer bursts so consumers process at a sustainable steady rate.
Competing consumers	Multiple workers receive from one queue for horizontal scale; each message to one worker.
Saga pattern	Distributed transaction = sequence of local transactions + compensating actions, coordinated via messages.
Auto-forwarding	Automatically chain queue/subscription to another entity — fan-in and decoupling topologies.

Prompt	Answer
Protocol	AMQP 1.0 (HTTPS fallback for some operations).

RBAC & Governance

Prompt	Answer
Role assignment formula	Security principal + role definition + scope.
Scope hierarchy	Management group → subscription → resource group → resource; assignments inherit downward.
Actions vs DataActions	Actions = control plane (ARM). DataActions = data plane (blob contents, queue messages...).
NotActions meaning	Subtracts from Actions — NOT a deny; other assignments can still grant those ops.
Deny assignments	Explicit deny that wins over grants; created by Azure (deployment stacks, managed apps), not directly by users.
Owner vs Contributor	Both manage everything; only Owner can assign roles.
Effective permissions	Union of all role assignments minus deny assignments.
Entra ID roles vs Azure RBAC	Separate planes: directory administration vs Azure resources. Global Admin can elevate to UAA at root '/' if needed.
System-assigned vs user-assigned MI	System: 1:1, dies with resource. User: standalone, shareable, survives deletion, pre-provisionable.
PIM	Just-in-time role activation: eligibility, approval, MFA, time-bound, access reviews, full audit.
ABAC conditions	Attribute-based filters on role assignments (e.g., blob tags) — finer than role+scope alone.
Custom role key property	AssignableScopes controls where the role can be assigned.
RBAC vs Azure Policy	RBAC = who can act. Policy = what configurations are allowed regardless of who.
Policy effects	Deny, Audit, Append, Modify, DeployIfNotExists, AuditIfNotExists, Disabled.
Resource locks	CanNotDelete (edit ok, no delete) and ReadOnly (no changes); apply to critical prod resources.
Management groups purpose	Hierarchy above subscriptions to inherit Policy + RBAC — backbone of landing zones.
Blueprints status	Deprecated — use template specs + deployment stacks.
Least-privilege checklist	Narrow data roles > Contributor; lowest scope; assign to groups; PIM for privileged; audit regularly.
Workload identity federation	External workloads (GitHub Actions, other clouds, K8s) exchange their OIDC tokens for Entra tokens — no stored secrets.

Networking

Prompt	Answer
NSG evaluation	Stateful; priority 100–4096, lowest first, first match wins; service tags and ASGs simplify rules.
ASG	Application security group: label NICs by workload so NSG rules reference groups, not IPs.
VNet peering transitivity	Non-transitive. Spoke↔spoke traffic must route via hub NVA/firewall (UDRs) or direct peering.
Hub-spoke hub contents	Azure Firewall/NVA, VPN/ER gateways, Bastion, DNS/Private Resolver — shared services.
Force egress through firewall	UDR 0.0.0.0/0 → firewall private IP on workload subnets.
VPN GW vs ExpressRoute	VPN: IPsec over internet, cheap, fast setup. ER: private circuit, SLA, up to 100 Gbps; pair with VPN failover.
Private endpoint	Private IP NIC mapped to a specific PaaS instance; on-prem reachable; lets you disable public access; needs private DNS zone.
Service endpoint	Subnet-to-service backbone path; service stays public; no on-prem reach; simpler/cheaper, weaker isolation.
Private endpoint DNS zone example	privatelink.blob.core.windows.net linked to VNets so the FQDN resolves to the private IP.
Azure DNS Private Resolver	Managed inbound/outbound endpoints for hybrid DNS — replaces DNS forwarder VMs.
Load balancer decision tree	L4 regional: Azure LB. L7 regional: App Gateway (WAF). L7 global HTTP: Front Door.
Front Door	DNS/global any-protocol: Traffic Manager. Global anycast L7: CDN caching, TLS offload, WAF, path routing, multi-region origin failover; Premium adds Private Link origins.
Traffic Manager routing methods	Priority, weighted, performance, geographic, subnet, multivalue — DNS-based, protocol-agnostic.
Azure Firewall Premium	Adds TLS inspection, IDPS, URL filtering, web categories over Standard.
Bastion	Browser-based RDP/SSH to VMs without public IPs.
NAT Gateway	Scalable predictable outbound SNAT; prevents port exhaustion.
Virtual WAN	Managed global transit: branches (SD-WAN/VPN), ER, P2S, hub firewalls, any-to-any routing.
DDoS protection options	Infrastructure (free, always-on), Network Protection (VNet-level plan), IP Protection (per-IP).

Container Apps

Prompt	Answer
ACA under the hood	AKS + KEDA (scaling) + Dapr (building blocks) + Envoy (ingress), fully abstracted — no K8s API access.
ACA vs AKS	ACA: serverless, no cluster ops, no CRDs/operators. AKS: full K8s control, custom operators, meshes, node access.
ACA vs App Service vs ACI vs Functions	ACA: microservices/containers. App Service: web apps + slots. ACI: single simple containers. Functions: event-driven code with bindings.
Environment	Isolation boundary sharing VNet + Log Analytics; internal (private ingress) or external.
Revisions	Immutable app snapshots; traffic splitting across revisions = canary/blue-green.
Scale rules	KEDA-based: HTTP concurrency, CPU/memory, Service Bus/Event Hubs/Kafka/cron scalers; minReplicas 0 = scale to zero.
ACA Jobs	Run-to-completion workloads: manual, scheduled (cron), or event-driven (KEDA).
Workload profiles	Mix Consumption (serverless) and Dedicated (fixed vCPU/mem, GPU) in one environment; enables UDR/firewall egress control.
Dapr in ACA	Sidecar per app: service invocation with mTLS/retries, pub/sub, state stores, bindings, secrets — opt-in per app.
ACR image pulls	Managed identity + AcrPull role — never admin credentials.
ACA secrets best practice	Key Vault references resolved via managed identity.
Cold start mitigation	Set minReplicas ≥ 1 for latency-sensitive APIs; scale to zero for batch/event workloads.

Auth & Identity

Prompt	Answer
OAuth 2.0 vs OIDC	OAuth = authorization (access tokens). OIDC = authentication layer on top (ID tokens).
Three token types	ID token (client authenticates user), access token (sent to APIs), refresh token (silent renewal).
Flow: web/SPA/mobile user sign-in	Authorization code + PKCE (implicit flow is deprecated).
Flow: daemon / service-to-service	Client credentials — prefer certificate, managed identity, or workload identity federation over client secrets.
Flow: API calling downstream API as the user	On-Behalf-Of (OBO) token exchange.
Flow: browserless/input-constrained device	Device code flow.
ROPC	Legacy username/password flow — breaks MFA/Conditional Access; avoid.

Prompt	Answer
App registration vs service principal	Registration = global app definition (home tenant). Service principal/enterprise app = its identity instance per tenant.
scp vs roles claims	scp = delegated scopes (user context). roles = app roles (app-only tokens or user role assignments).
Admin consent rule	Application (app-only) permissions ALWAYS require admin consent.
JWT validation checklist	Signature (JWKS), issuer, audience, expiry/nbf → then authorize on scp/roles. Every hop validates (zero trust).
APIM token enforcement	validate-jwt / validate-azure-ad-token with required-claims; reject at the edge, backend re-validates.
Entra External ID	CIAM (successor to Azure AD B2C): consumer sign-up/sign-in, social IdPs, custom branding.
B2B vs CIAM	B2B = partner guests in YOUR tenant. External ID/CIAM = customers in a separate consumer directory.
Conditional Access	If-then policy engine on sign-in signals (user, device, location, risk) → require MFA/compliant device/block.
Managed identity value	Azure-managed Entra identity for resources — zero secrets for Azure-to-Azure auth.
SAS appropriate use	Time-boxed, scoped delegation (e.g., blob download links); prefer user-delegation SAS; otherwise use Entra RBAC.
mTLS pattern	Mutual certificate auth for high-assurance service-to-service (APIM client certs to backend), layered with JWT + private network.
Zero Trust principles	Verify explicitly. Least privilege (JIT/JEA). Assume breach (segment, encrypt, monitor).
Multi-tenant token validation	Issuer varies per tenant — validate iss/tid against an allowlist or pattern.

Design Patterns

Prompt	Answer
API Gateway pattern	Single entry point handling auth, throttling, routing, transformation — APIM's role.
Backends for Frontends	Separate tailored backend per client type (mobile/web/partner).
Strangler Fig	Route slices of a monolith to new services via a facade until the monolith is gone.
Gateway offloading / routing / aggregation	Offload cross-cutting concerns / direct requests / compose multiple backend calls into one response.
Circuit breaker	Trip after repeated failures, fail fast, probe recovery — prevents cascading failure. Pair with retry+backoff.
Bulkhead	Partition resource pools so one consumer/workload can't exhaust capacity for others.

Prompt	Answer
Retry guidance	Exponential backoff + jitter, idempotent operations only, cap attempts, honor Retry-After.
CQRS	Separate command (write) and query (read) models — scale and optimize independently.
Event sourcing	Append-only event log as source of truth; state rebuilt by replay; projections serve reads.
Cache-aside	Check cache → miss → read store → populate cache; invalidate on writes.
Saga	Distributed transaction via local transactions + compensating actions (choreography or orchestration).
Claim-check	Reference message + payload in blob storage — beats message size limits.
Sidecar / Ambassador	Co-located helper container for connectivity, security, observability (e.g., Dapr).
Anti-corruption layer	Translation layer isolating a new system from legacy models/protocols.
Deployment stamps / Geode	Stamps: repeatable isolated stack units per tenant/region. Geode: active-active regional nodes behind global routing.
Health endpoint monitoring	Expose functional health probes for LBs/orchestrators to gate traffic.
Queue-based load leveling	Broker buffer between spiky producers and steady consumers.
Throttling pattern	Reject/degrade excess load to protect the service (APIM rate-limit/quota).

AZ-305 Exam

Prompt	Answer
AZ-305 domains + weights (Apr 2026)	Identity/governance/monitoring 25–30% • Data storage 25–30% • Business continuity 10–15% • Infrastructure 30–35%.
AZ-305 logistics	~50 questions incl. case studies, 700/1000 to pass, ~\$165; Expert cert requires AZ-104 prerequisite.
Composite SLA — serial	Multiply: $99.95\% \times 99.9\% \approx 99.85\%$. Every serial dependency lowers availability.
Composite SLA — parallel	$1 - (1 - A)^2$: two 99.9% regions $\approx 99.9999\%$ theoretical.
RTO vs RPO	RTO = max downtime tolerated. RPO = max data loss window tolerated.
Well-Architected pillars	Reliability, Security, Cost Optimization, Operational Excellence, Performance Efficiency.
CAF phases	Strategy → Plan → Ready (landing zones) → Adopt (Migrate/Innovate) + ongoing Govern & Manage.
Answer strategy	Map stated requirements to the CHEAPEST compliant service; never gold-plate; watch ‘minimize cost/effort’ qualifiers.

Prompt	Answer
Zone vs region redundancy	Zones survive datacenter failure (same region, low latency). Region pairs/multi-region survive regional disasters.
Hybrid identity options	Password hash sync (simplest, recommended) → pass-through auth → ADFS federation (most infra, most control).
Cost tools	Pricing Calculator (pre-deploy), TCO Calculator (on-prem compare), Cost Management (actuals), Advisor (savings recs), Reservations/Savings Plans (commit discounts).
Monitoring stack	Azure Monitor metrics + Log Analytics (KQL) + App Insights (APM) + alerts/action groups + workbooks.
Free official AZ-305 practice	Microsoft Learn Practice Assessment: learn.microsoft.com/credentials/certifications/exams/az-305/practice/assessment
Top prep resources	MS Learn AZ-305 paths, John Savill's AZ-305 playlist (YouTube), MeasureUp official practice tests, Tutorials Dojo/Whizlabs mocks, Architecture Center case studies.
Architecture Center + WAF role in exam	AZ-305 explicitly aligns to Well-Architected Framework and Cloud Adoption Framework — justify every recommendation with them.
Common exam traps	Developer APIM tier in prod (no SLA) • Geo-DR replicates metadata only • service endpoint ≠ private endpoint • NotActions ≠ deny • scp ≠ roles.
Reference architecture (memorize)	Front Door+WAF → internal APIM → ACA/AKS (managed identities) → Service Bus → workers → private-endpoint data tier; hub VNet with firewall/ER; Key Vault + Log Analytics throughout.

API Management

Prompt	Answer
Multi-region APIM behavior	Premium replicates gateways to added regions; primary region hosts the management plane; route callers to nearest gateway.
CORS in APIM	cors policy at the appropriate scope — required before browser-based consumers can call your APIs.

Service Bus

Prompt	Answer
Max delivery count	Deliveries before automatic dead-lettering (default 10) — tune with retry expectations.

Prompt	Answer
Prefetch	Client pre-fetches messages to boost throughput; balance against lock expiry risk.

RBAC & Governance

Prompt	Answer
Reader vs Monitoring Reader	Use purpose-built limited roles where possible; Reader sees configuration of everything in scope.
Elevate access (Global Admin)	Break-glass toggle granting UAA at root '/'; audited; remove after use.

Networking

Prompt	Answer
Global VNet peering	Cross-region, cross-subscription private connectivity over Microsoft backbone; non-transitive.
Forced tunneling	Send internet-bound traffic to on-prem/NVA via UDR/BGP for inspection/compliance.

Container Apps

Prompt	Answer
Ingress options	External or internal-only; HTTPS by default (Envoy), TCP supported, IP restrictions, custom domains, session affinity.
Traffic splitting example	Revision A 90% / Revision B 10% → observe metrics → shift to 100% B or roll back instantly.

Auth & Identity

Prompt	Answer
OBO flow mechanics	API exchanges the incoming user token (assertion) at Entra for a downstream-API token preserving user identity.
Token lifetime defaults	Access tokens ~60-90 min; refresh via MSAL silently; design APIs to tolerate token refresh mid-session.

Design Patterns

Prompt	Answer
Choreography vs orchestration (saga)	Choreography: services react to events (loose coupling). Orchestration: central coordinator (Durable Functions/Logic Apps) directs steps.
Materialized view	Precompute query-optimized views from the write model — common CQRS read side.

AZ-305 Exam

Prompt	Answer
Case study time strategy	Read requirements/constraints first, skim background, answer requirement-by-requirement; ~2 case studies typical.
Renewal	Microsoft certs renew annually via a FREE online, open-book renewal assessment on Microsoft Learn.

Data Storage

Prompt	Answer
SQL Database vs Managed Instance vs SQL on VM	SQL DB: cloud-first single DBs. MI: near-full SQL Server compat (Agent, cross-DB), VNet-native — lift-and-shift. VM: full OS/engine control.
DTU vs vCore	DTU: bundled blend, simple/small. vCore: independent compute/storage, Hybrid Benefit, reservations, serverless.
SQL service tiers	General Purpose: default, remote storage. Business Critical: local SSD, readable secondary, lowest latency. Hyperscale: 100+ TB, instant restores, fast replicas.
SQL serverless	Auto-scales vCores and auto-pauses; pay storage only while paused — intermittent dev/test.
Elastic pools	Many variable-usage databases share provisioned resources — multi-tenant SaaS cost optimization.
Cosmos DB consistency levels	Strong → Bounded Staleness → Session (default; read-your-own-writes) → Consistent Prefix → Eventual.
Cosmos RU/s models	Provisioned standard, autoscale (variable sustained), serverless (per-request; low/spiky traffic).
Cosmos partition key rules	High cardinality, even distribution, matches dominant filters; avoid hot partitions; 20 GB logical partition cap.
Cosmos multi-region writes	Active-active writes in all regions, RPO≈0, conflict resolution policies; 99.999% SLA potential.

Prompt	Answer
Cosmos change feed	Ordered stream of container changes — event sourcing, materialized views, downstream triggers.
Storage redundancy ladder	LRS (1 DC) → ZRS (3 zones) → GRS/RA-GRS (async pair region) → GZRS/RA-GZRS (zones + region; highest durability).
Blob access tiers	Hot → Cool (30d) → Cold (90d) → Archive (180d, offline, rehydration hours). Lifecycle policies auto-tier/expire.
Immutable storage	WORM: time-based retention or legal hold — compliance (SEC 17a-4) and ransomware defense.
ADLS Gen2	Blob + hierarchical namespace: directories, POSIX ACLs — the data-lake foundation.
Azure File Sync	On-prem Windows servers become hot caches of Azure Files with cloud tiering.
Data store quick picks	OLTP relational: SQL. Global NoSQL: Cosmos. Objects/lake: Blob/ADLS. SMB shares: Files. Cache: Redis. Telemetry analytics: ADX.

Business Continuity

Prompt	Answer
RTO / RPO	RTO = max downtime. RPO = max data loss. Drive every BC design decision from these two numbers per workload tier.
Backup vs replication	Backup = point-in-time recovery (corruption/ransomware/deletion). Replication = infrastructure loss. Replication copies corruption too — need both.
Vault types	Recovery Services vault: VMs, SQL-in-VM, Files, MARS. Backup vault: blobs, disks, PostgreSQL.
Backup ransomware defenses	Soft delete (14+ days), immutable vaults, multi-user authorization for destructive changes.
Azure Site Recovery	Continuous VM replication + recovery plans (ordered tiers, scripts) + test failover in isolated networks. RPO seconds-minutes.
SQL auto-failover groups	Cross-region group failover with stable RW/RO listener endpoints — apps reconnect without config changes.
SQL backup retention	PITR 1-35 days; long-term retention (LTR) weekly/monthly/yearly up to 10 years; geo-restore from geo-backups.
Availability ladder	Single VM 99.9% → availability set 99.95% → zones 99.99% → multi-region (paired) for regional disasters.
Active-passive vs active-active	Passive: priority routing, cheaper, higher RTO. Active: both serve traffic; needs multi-write or partitioned data.
Storage account failover	Customer-initiated with (RA-)GRS promotes secondary; data after Last Sync Time is lost (async replication).

Prompt	Answer
Chaos Studio	Managed fault injection (zone down, CPU stress, network latency) to validate resilience assumptions.

Monitoring

Prompt	Answer
Azure Monitor data split	Metrics: numeric time-series, near-real-time alerts. Logs: Log Analytics workspace, KQL, richer/slower.
Application Insights	APM: distributed tracing, Application Map, live metrics, availability tests, smart detection; workspace-based.
Alert types	Metric (fast, platform metrics), log/scheduled-query (KQL), activity log, smart detection → all fire action groups.
Action groups	Reusable receiver sets: email/SMS/push/webhook/Function/Logic App/ITSM. Alert processing rules suppress/route at scale.
Diagnostic settings destinations	Log Analytics (analyze), Storage (cheap archive), Event Hubs (external SIEM/stream).
Workspace design default	Centralize per environment/region for correlation; resource-context/table RBAC isolates teams; split only for sovereignty/chargeback.
Log cost levers	Retention tuning, long-term/archive tiers, daily caps, Basic/Auxiliary table plans, App Insights sampling.
Service Health vs Resource Health	Service Health: Azure platform incidents/maintenance (alertable). Resource Health: your resource's availability state.
Defender for Cloud vs Sentinel	Defender: CSPM (secure score) + workload protection. Sentinel: SIEM/SOAR — cross-source correlation, hunting, playbooks.
Network Watcher	NSG/VNet flow logs, connection monitor, IP flow verify, packet capture — network diagnostics toolkit.

Compute

Prompt	Answer
Compute decision tree	OS control: VM/VMSS. Web PaaS: App Service. Event code: Functions. Microservices sans K8s: ACA. Full K8s: AKS. One container: ACI. HPC: Batch.
VM availability SLAs	99.9% single (premium SSD) / 99.95% availability set / 99.99% across zones.
Spot VMs	Up to ~90% off, evictable with 30s notice — batch, CI, stateless only.

Prompt	Answer
Functions hosting plans	Consumption (scale-to-zero, cold starts) / Flex Consumption / Premium (pre-warmed, VNet, unlimited) / Dedicated plan.
Durable Functions patterns	Chaining, fan-out/fan-in, async HTTP APIs, monitor, human interaction, saga orchestration.
AKS security baseline	Private cluster, Entra workload identity for pods, Entra+K8s RBAC, Defender for Containers, GitOps (Flux).
Cost commitment options	Reservations (deepest, locked family/region), savings plans (flexible \$/hr commit), Hybrid Benefit (licenses), Spot (interruptible).

Migration & Integration

Prompt	Answer
The 5 Rs	Rehost (lift-shift) · Refactor (minor PaaS changes) · Rearchitect (microservices) · Rebuild (cloud-native rewrite) · Replace (SaaS).
Azure Migrate	Discovery, dependency mapping, right-size + cost assessment, then migration of servers/DBs/web apps.
Database Migration Service	Online (minimal downtime, continuous sync) or offline migrations to Azure SQL family; DMA assesses compatibility first.
Azure Data Box	Encrypted appliance for offline TB-PB transfers when bandwidth is impractical.
ADF self-hosted IR	Integration runtime installed on-prem to reach private/on-prem data sources from Data Factory pipelines.
Analytics quick picks	ETL orchestration: ADF. Unified SaaS analytics: Fabric. Spark/ML: Databricks. Streams SQL: Stream Analytics. Telemetry KQL: ADX. Low-code workflows: Logic Apps.

Auth & Identity

Prompt	Answer
Key Vault protection duo	Soft delete (default, recoverable) + purge protection (no permanent delete until retention ends — required for CMK).
Key Vault vs Managed HSM	Standard: software keys. Premium: shared HSM-backed. Managed HSM: dedicated single-tenant FIPS 140-2 Level 3 pool.
Key Vault access model	Prefer Azure RBAC data roles (Secrets User, Crypto User...) over legacy access policies; private endpoints + disabled public access.
App Configuration	Central non-secret settings + feature flags, with Key Vault references for secrets — pairs with safe deployment/feature rings.

RBAC & Governance

Prompt	Answer
Azure Arc	Projects on-prem/multicloud servers, K8s, and data services into ARM — Policy, RBAC, tags, Monitor, Defender everywhere.

Design Patterns

Prompt	Answer
Bicep vs Terraform	Bicep: Azure-native, day-0 ARM, no state file. Terraform: multicloud, state-managed, huge ecosystem. Both exam-valid IaC.
Deployment stacks	Manage deployments as units: deny settings block out-of-band edits; deleted-from-template resources cleaned up. Blueprints successor.
Pipeline auth best practice	OIDC workload identity federation from GitHub Actions/Azure DevOps — zero stored cloud secrets.
Pre-production validation services	Azure Load Testing (JMeter/Locust at scale) + Chaos Studio (fault injection) wired into CI/CD gates.

API Management

Prompt	Answer
[Simply] APIM in plain words	A hotel front desk: guests never enter the kitchen (backends). The desk checks keys (auth), enforces house rules (rate limits), answers from memory (cache), and forwards requests to the right department (routing).
[Simply] Rate limit vs quota, simply	Speed limit (right now) vs monthly fuel allowance (long term). You can obey one and still violate the other.
[Simply] Versions vs revisions, simply	New book edition (readers choose it; story changes) vs a reprint fixing typos (silently replaces stock; nothing breaks).
Config: per-caller throttling key	rate-limit-by-key with counter-key=@(context.Request.IpAddress) or a JWT claim — a constant key makes every caller share one bucket.

Service Bus

Prompt	Answer
[Simply] Queue in plain words	A deli ticket machine: any free clerk serves the next ticket; nobody is served twice; a rush just makes the line longer (load leveling).

Prompt	Answer
[Simply] Topic in plain words	A magazine subscription list: print once, every subscriber gets a copy, some only take the sports edition (filters).
[Simply] DLQ in plain words	The post office's undeliverable-mail shelf: after failed attempts, the letter waits with a note saying why, for a human to investigate.
[Simply] Claim-check in plain words	A coat check: pass the small ticket (blob reference) around the party, not the winter coat (huge payload).
[Simply] Event vs message, simply	Doorbell (fact happened, react if you care) vs courier package (payload the sender expects processed, signature required).
Config: lock duration vs processing time	If processing exceeds the lock, the message redelivers mid-work → duplicates. Renew locks or extend duration; always code idempotent handlers.
Config: TTL + dead-letter-on-expiration	TTL expiry with the flag off silently DELETES messages; with it on, they land in the DLQ where monitoring can catch them.

RBAC & Governance

Prompt	Answer
[Simply] RBAC in plain words	Office keycards: who you are (principal) + which doors open (role) + which floors (scope). Cards go to teams (groups), not individuals.
[Simply] PIM in plain words	The master-key sign-out sheet: borrow with a reason, manager approves, auto-return in 4 hours, every borrow logged.
[Simply] Managed identity in plain words	A badge issued by the building itself — the app never carries a password, so there's nothing to steal or leak to GitHub.
Config: control vs data plane 403	'Reader' on a storage account ≠ reading blobs. Data access needs DataActions roles like Storage Blob Data Reader, scoped as low as the container.
Config: IMDS token request	curl 169.254.169.254/meta-data/identity/oauth2/token?resource=https://storage.azure.com — the resource/audience must match the service you call.

Networking

Prompt	Answer
[Simply] Private endpoint in plain words	A private elevator installed straight into your office: the public lobby is bricked over, and the building directory (private DNS) routes visitors to your floor.

Prompt	Answer
[Simply] Hub-spoke in plain words	An airport hub: every regional flight connects through it, where security (firewall), customs (gateways), and the control tower (DNS/monitoring) live once.
[Simply] ER vs VPN, simply	A private rail line vs an armored truck on the public highway — both safe, but the rail line never touches public roads and carries far more freight.
[Simply] LB family in plain words	Front Door = global concierge. App Gateway = building receptionist with metal detector (WAF). Load Balancer = elevator dispatcher. Traffic Manager = phone directory (DNS).
Config: private endpoint DNS trio	(1) privatelink zone exists, (2) zone linked to the consuming VNet, (3) DNS zone group creates the A record. Custom DNS must forward to 168.63.129.16.
Config: spoke↔spoke traffic	Peering is non-transitive — add UDRs (0.0.0.0/0 or spoke CIDRs → hub firewall) to route between spokes through the hub.

Container Apps

Prompt	Answer
[Simply] Scale to zero in plain words	Motion-sensor lights: dark (free) until someone walks in, with a half-second flicker (cold start) as they switch on.
[Simply] Canary revisions in plain words	Test the new recipe on 10% of tables; complaints → old recipe returns instantly; praise → roll it out to every table.
[Simply] KEDA in plain words	Open supermarket checkouts because six people are queuing (queue depth), not because the clock says 5 pm.
Config: traffic split prerequisites	revisions-mode multiple + ingress traffic set -revision-weight A=90 B=10; single-revision mode ignores weights.
Config: Service Bus scale rule	scale-rule-type azure-servicebus with queueName/namespace/messageCount metadata + connection auth secret; min-replicas 0 enables scale-to-zero.

Auth & Identity

Prompt	Answer
[Simply] AuthN vs AuthZ, simply	Airport ID check (who you are) vs boarding pass (what you may do: this flight, this seat). Both checks, in that order.
[Simply] Access token in plain words	A valet key: starts the car but not the trunk (limited scopes), expires soon, and you never hand over your real key (password).

Prompt	Answer
[Simply] Zero trust in plain words	A nightclub where every door has a bouncer: front-door entry opens nothing else; wristbands match what you paid for; cameras assume someone snuck in.
[Simply] Key Vault in plain words	Hotel safe with a logbook: valuables out from under mattresses (config files), every opening logged, can't discard the safe while full (purge protection).
Config: app-only token request	client_credentials + scope=api://your-api/.default (rolls up app roles). Named scopes in that flow fail — they're for delegated auth.
Config: break-glass accounts	Two cloud-only accounts excluded from Conditional Access, long random passwords, alert on every sign-in — insurance against CA/MFA lockout.

Data Storage

Prompt	Answer
[Simply] Storage redundancy in plain words	Photo backups: one shoebox (LRS), three rooms (ZRS), a copy mailed to grandma — slightly delayed (GRS), rooms AND grandma (GZRS).
[Simply] Blob tiers in plain words	Closet (Hot) → attic (Cool) → storage unit (Cold) → bank vault: cheap but retrieval needs an appointment (Archive rehydration).
[Simply] Cosmos consistency in plain words	Group chat: Strong = nobody sees a message until everyone can. Session = you always see your own instantly (default). Eventual = all arrive, maybe out of order.
[Simply] Partition key in plain words	Filing cabinets: file orders by customer and lookups fly; file everything under '2026' and one drawer jams (hot partition) while others sit empty.

Business Continuity

Prompt	Answer
[Simply] RTO/RPO in plain words	'How long until the shop reopens?' (RTO) vs 'how many receipts did we lose?' (RPO). Price both fears separately.
[Simply] Backup vs replication, simply	Mirror vs photo album: the mirror instantly shows the ketchup stain too (corruption); the album lets you return to before the spill.
[Simply] Composite SLA in plain words	Christmas lights: serial = every extra bulb can kill the string (multiply, availability drops). Parallel = two strings side by side (availability jumps).

Prompt	Answer
Standby temperature ladder	Cold (rebuild from IaC, hours, cheapest) → warm (scaled-down, minutes) → hot/active-active (seconds, priciest). RTO picks the rung.
Config: SQL PITR restore	az sql db restore creates a NEW database at the chosen time — plan the cutover/rename step and rehearse it; that time is your real RTO.
Config: protection before disaster	Soft delete, versioning, immutability only cover events AFTER enablement, and retention must exceed detection time (found on day 9 > 7-day retention = gone).

Monitoring

Prompt	Answer
[Simply] Metrics vs logs, simply	Car dashboard (instant, cheap, alarms) vs trip journal (rich history for later questions — KQL).
[Simply] Distributed tracing in plain words	A barcode that follows one package end-to-end: scan one order ID and see every warehouse, truck, and doorstep it touched.
[Simply] Policy vs RBAC, simply	Door keys (who may build) vs building codes (what anyone may build). Codes stop wooden shacks even for landowners.

Design Patterns

Prompt	Answer
[Simply] GitOps in plain words	The cluster reads its recipe book (Git) and cooks accordingly; fix the book, not the kitchen — rollback = revert the page, and the book's margin notes are your audit log.
Scenario: FinSecure identity stack	PHS (resilient, zero infra) + Conditional Access baseline + PIM JIT + break-glass + policy-carrying management groups. License P2 only for admins.
Scenario: ShopSphere data residency	EU data on EU SQL servers, failover groups paired within the geography — DR that respects GDPR by architecture, not paperwork.
Scenario: ForgeWorks recovery stack	Zones for common failures; warm standby region for disasters; ASR for legacy VMs; failover groups for SQL; immutable GRS backups for ransomware; quarterly drills.
Scenario: CloudGate isolation	Spoke-per-customer Bicep stamps, all egress via hub Firewall Premium, private endpoints + policy deny on public access, central logs with resource-context RBAC.

Prompt	Answer
Scenario: RetailRun supply chain	OIDC federation (CI) + workload identity (pods) + private ACR with scanning + approved-registry policy + Flux GitOps = zero stored secrets, full provenance.

AZ-305 Exam

Prompt	Answer
Lab habit that wins the exam	After every hands-on step ask: which requirement (cost, RTO, isolation, least privilege) did this flag serve? AZ-305 tests the why, not the syntax.
Wording traps decoder	'Minimize cost' → cheapest compliant. 'Minimize admin effort' → most managed/PaaS. 'Most secure' → private + MI + JIT. 'Minimize changes' → lift-and-shift options.